

System V Application Binary Interface
MD64 Architecture Processor Supplement
(With LP64 and ILP32 Programming Models)
Version 1.0

Edited by
H.J. Lu¹, Michael Matz², Milind Girkar³, Jan Hubička⁴,
Andreas Jaeger⁵, Mark Mitchell⁶

September 28, 2021

¹hongjiu.lu@intel.com

²matz@suse.de

³milind.girkar@intel.com

⁴jh@suse.cz

⁵aj@suse.de

⁶mark@codesourcery.com

Contents

1	Introduction	11
2	Software Installation	12
3	Low Level System Information	13
3.1	Machine Interface	13
3.1.1	Processor rchitecture	13
3.1.2	Data Representation	15
3.2	Function Calling Sequence	19
3.2.1	Registers	19
3.2.2	The Stack Frame	20
3.2.3	Parameter Passing	21
3.3	Operating System Interface	28
3.3.1	Exception Interface	28
3.3.2	Virtual ddress Space	28
3.3.3	Page Size	28
3.3.4	Virtual ddress ssignments	28
3.4	Process Initialization	31
3.4.1	Initial Stack and Register State	31
3.4.2	Thread State	33
3.4.3	uxiliary Vector	34
3.5	Coding Examples	37
3.5.1	rchitectural Constraints	37
3.5.2	Conventions	40
3.5.3	Position-Independent Function Prologue	40
3.5.4	Data Objects	41
3.5.5	Function Calls	47
3.5.6	Branching	50

3.5.7	Variable Argument Lists	53
3.6	DW RF Definition	57
3.6.1	DW RF Release Number	58
3.6.2	DW RF Register Number Mapping	58
3.7	Stack Unwind Algorithm	58
4	Object Files	61
4.1	ELF Header	61
4.1.1	Machine Information	61
4.1.2	Number of Program Headers	62
4.2	Sections	62
4.2.1	Section Flags	62
4.2.2	Section types	63
4.2.3	Special Sections	63
4.2.4	EH_FRT ME sections	64
4.3	Symbol Table	68
4.4	Relocation	69
4.4.1	Relocation Types	69
4.4.2	Large Models	74
5	Program Loading and Dynamic Linking	75
5.1	Program Loading	75
5.1.1	Program header	76
5.2	Dynamic Linking	76
5.2.1	Program Interpreter	81
5.2.2	Initialization and Termination Functions	82
6	Libraries	83
6.1	C Library	83
6.1.1	Global Data Symbols	83
6.1.2	Floating Point Environment Functions	83
6.2	Unwind Library Interface	83
6.2.1	Exception Handler Framework	84
6.2.2	Data Structures	86
6.2.3	Throwing an Exception	88
6.2.4	Exception Object Management	91
6.2.5	Context Management	91
6.2.6	Personality Routine	94
6.3	Unwinding Through Assembly Code	97

7 Development Environment	100
8 Execution Environment	101
9 Conventions	102
9.1 C++	103
9.2 Fortran	104
9.2.1 Names	104
9.2.2 Representation of Fortran Types	105
9.2.3 Argument Passing	106
9.2.4 Functions	107
9.2.5 COMMON blocks	108
9.2.6 Intrinsic s	109
10 ILP32 Programming Model	119
10.1 Parameter Passing	119
10.2 Address Space	119
10.3 Thread-Local Storage Support	119
10.3.1 Global Thread-Local Variable	119
10.3.2 Static Thread-Local Variable	121
10.3.3 TLS Linker Optimization	123
10.4 Kernel Support	126
10.5 Coding Examples	126
10.5.1 Indirect Branch	126
11 Alternate Code Sequences For Security	127
11.1 Code Sequences without PLT	127
11.1.1 Indirect Call via the GOT Slot	127
11.1.2 Thread-Local Storage without PLT	128
12 Intel MPX Extension	133
12.1 Parameter Passing and Returning of Values	133
12.1.1 Classification	133
12.1.2 Passing	134
12.1.3 Returning of Values	135
12.1.4 Variable Argument Lists	137
12.2 Program Loading and Dynamic Linking	137

Linux Conventions	140
.1 Execution of 32-bit Programs	140
.2 MD64 Linux Kernel Conventions	140
.2.1 Calling Conventions	140
.2.2 Stack Layout	141
.2.3 Miscellaneous Remarks	141
B Linker Optimization	142
B.1 Combine GOTPLT and GOT Slots	142
B.2 Optimize GOTPCRELX Relocations	143

List of Tables

3.1	Micro- rchitecture Levels	14
3.2	Hardware Exceptions and Signals	29
3.3	Floating-Point Exceptions	29
3.4	x87 Floating-Point Control Word	31
3.5	MXCSR Status Bits	31
3.6	rFL GS Bits	32
4.1	MD64 Specific Section Header Flag, <code>sh_flags</code>	62
4.2	Section Header Types	63
4.3	Special sections	63
4.4	dditional Special Sections for the Large Code Model	64
4.5	Common Information Entry (CIE)	65
4.6	CIE ugmentation Section Content	66
4.7	Frame Descriptor Entry (FDE)	67
4.8	FDE ugmentation Section Content	68
4.9	Relocation Types	71
4.10	Large Model Relocation Types	74
5.1	Program Header Types	76
7.1	Predefined Pre-Processor Symbols	100
9.1	Mil intrinsics	109
9.2	F77 intrinsics	111
9.3	F90 intrinsics	112
9.4	Math intrinsics	112
9.5	Unix intrinsics	114
10.1	General Dynamic Model Code Sequence	120
10.2	General Dynamic Model Code Sequence with TLSDESC	120

10.3	Initial Exec Model Code Sequence	120
10.4	Initial Exec Model Code Sequence, II	121
10.5	Local Dynamic Model Code Sequence With Lea	121
10.6	Local Dynamic Model Code Sequence With dd	121
10.7	General Dynamic Model Code Sequence with TLSDESC	122
10.8	Local Dynamic Model Code Sequence, II	122
10.9	Local Exec Model Code Sequence With Lea	122
10.10	Local Exec Model Code Sequence With dd	122
10.11	Local Exec Model Code Sequence, II	123
10.12	Local Exec Model Code Sequence, III	123
10.13	GD -> IE Code Transition	123
10.14	GDesc -> IE Code Transition	123
10.15	GD -> LE Code Transition	124
10.16	GDesc -> LE Code Transition	124
10.17	IE -> LE Code Transition With Lea	124
10.18	IE -> LE Code Transition With dd	124
10.19	IE -> LE Code Transition, II	125
10.20	LD -> LE Code Transition With Lea	125
10.21	LD -> LE Code Transition With dd	125
10.22	LD -> LE Code Transition, II	125
10.23	Indirect Branch	126
11.1	General Dynamic Model Code Sequence (LP64)	129
11.2	General Dynamic Model Code Sequence (ILP32)	129
11.3	Local Dynamic Model Code Sequence (LP64)	130
11.4	Local Dynamic Model Code Sequence (ILP32)	130
11.5	GD -> IE Code Transition (LP64)	130
11.6	GD -> IE Code Transition (ILP32)	131
11.7	GD -> LE Code Transition (LP64)	131
11.8	GD -> LE Code Transition (ILP32)	131
11.9	LD -> LE Code Transition (LP64)	132
11.10	LD -> LE Code Transition (ILP32)	132
B.1	Call and Jmp Conversion	143
B.2	Mov Conversion	144
B.3	Test and Binop Conversion	144

List of Figures

3.1	Scalar Types	16
3.2	Bit-Field Ranges	18
3.3	Stack Frame with Base Pointer	20
3.4	Register Usage	25
3.5	Parameter Passing Example	27
3.6	Register Allocation Example	27
3.7	Virtual Address Configuration	30
3.8	Conventional Segment Arrangements	30
3.9	Initial Process Stack	33
3.10	<code>auxv_t</code> Type Definition	34
3.11	Auxiliary Vector Types	35
3.12	Position-Independent Function Prolog Code	41
3.13	Absolute Load and Store (Small Model)	42
3.14	Position-Independent Load and Store (Small PIC Model)	43
3.15	Absolute Load and Store (Medium Model)	44
3.16	Position-Independent Load and Store (Medium PIC Model)	45
3.17	Position-Independent Load and Store (Medium PIC Model), continued	45
3.18	Absolute Global Data Load and Store	46
3.19	Faster Absolute Global Data Load and Store	46
3.20	Position-Independent Global Data Load and Store	47
3.21	Faster Position-Independent Global Data Load and Store	47
3.22	Position-Independent Direct Function Call (Small and Medium Model)	47
3.23	Position-Independent Indirect Function Call	48
3.24	Absolute Direct and Indirect Function Call	48
3.25	Position-Independent Direct and Indirect Function Call	49
3.26	Absolute Branching Code	50
3.27	Implicit Calculation of Target Address	51
3.28	Position-Independent Branching Code	51
3.29	Absolute Switch Code	52

3.30	Position-Independent Switch Code	53
3.31	Parameter Passing Example with Variable- argument List	54
3.32	Register llocation Example for Variable- rgument List	54
3.33	Register Save rea	55
3.34	<code>va_list</code> Type Declaration	55
3.35	Sample Implementation of <code>va_arg(l, int)</code>	57
3.36	DW RF Register Number Mapping	59
3.37	Pointer Encoding Specification Byte	60
4.1	Relocatable Fields	69
5.1	Global Offset Table	77
5.2	Procedure Linkage Table (small and medium models)	78
5.3	Final Large Code Model PLT	81
5.4	MD64 Program Interpreter	82
6.1	Examples for Unwinding in ssembler	99
9.1	Example mapping of names	105
9.2	Mapping of Fortran to C types	105
11.1	Function Call without PLT (Small and Medium Models)	127
11.2	Function ddress without PLT (Small and Medium Models)	128
11.3	<code>__tls_get_addr</code> Call	129
12.1	Bound Register Usage	134
12.2	Bounds Passing Example	136
12.3	Bounds llocation Example	136
12.4	BND Procedure Linkage Table (small and medium models)	138
12.5	dditional Procedure Linkage Table (small and medium models)	138
12.6	Indirect branch	139
B.1	Procedure Linkage Table Entry Via GOTPLT Slot	142
B.2	Procedure Linkage Table Entry Via GOT Slot	143

Revision History

Intel Changes to this document made by Intel Corporation are copyright Intel Corporation and licensed under CC BY 4.0: <https://creativecommons.org/licenses/by/4.0/legalcode>.

- 1.0** Document `_Unwind_GetIPInfo`. Update to match C++ BI. dd chapter for alternate code sequences For security.
- 0.99** dd description of TLS relocations (thanks to Alexandre Oliva) and mention the decimal floating point and VX types (thanks to H.J. Lu).
- 0.98** Various clarifications and fixes according to feedback from Sun, thanks to Terrence Miller. DW RF register numbers for some system registers, thanks to Jan Beulich.
dd `R_X86_64_SIZE32` and `R_X86_64_SIZE64` relocations; extend meaning of `e_phnum` to handle more than 0xffff program headers, thanks to Rod Evans. dd footnote about passing of `decimal` datatypes. Specify that `_Bool` is booleanized at the caller.
- 0.97** Integrate Fortran BI.
- 0.96** Use `SHF_X86_64_L_RGE` instead `SHF_M64_L_RGE` (thanks to Evandro Menezes). Correct various grammatical errors noted by Mark F. Haigh, who also noted that there are no global VL s in C99. Thanks also to Robert R. Henry.
- 0.95** Include description of the medium PIC memory model (thanks to Jan Hubička) and large model (thanks to Evandro Menezes).
- 0.94** dd sections in Development Environment, Program Loading, a description of EH_FR ME sections and general cleanups to make text in this BI self-contained. Thanks to Michael Walker and Terrence Miller.
- 0.93** dd sections about program headers, new section types and special sections for unwinding information. Thanks to Michael Walker.
- 0.92** Fix some typos (thanks to Bryan Ford), add section about stack layout in the Linux kernel. Fix example in figure 3.5 (thanks to Tom Horsley). dd section on unwinding through assembler (written by Michal Ludvig). Remove `mmxext` feature (thanks to Evandro Menezes). dd section on Fortran (by Steven Bosscher) and stack unwinding (by Jan Hubička).
- 0.91** Clarify that x87 is default mode, not MMX (by Hans Peter nvin).
- 0.90** Change DW RF register numbers again; mention that `_m128` needs alignment; fix typo in figure 3.3; add some comments on kernel expectations; mention TLS extensions; add example for passing of variable-argument lists; change semantics of `%rax` in variable-argument lists; improve formatting; mention that X87 class is not used for passing; make `/lib64` a Linux specific section; rename x86-64 to MD64;

describe passing of complex types. Special thanks to Endi Kleen, Michal Ludvig, Michael Matz, David O'Brien and Eric Young for their comments.

- 0.21** Define `__int128` as class INTEGER in register passing. Mention that `%a1` is used for variadic argument lists. Fix some textual problems. Thanks to H. Peter Nivin, Bo Thorsen, and Michael Matz.
- 0.20 — 2002-07-11** Change DW RF register number values of `%rbx`, `%rsi`, `%rwi` (thanks to Michal Ludvig). Fix footnotes for fundamental types (thanks to H. Peter Nivin). Specify `size_t` (thanks to Bo Thorsen and Andreas Schwab). Add new section on floating point environment functions.
- 0.19 — 2002-03-27** Set name of Linux dynamic linker, mention `%fs`. Incorporate changes from H. Peter Nivin <hpa@zytor.com> for booleans and define handling of sub-64-bit integer types in registers.

Chapter 1

Introduction

The MD64¹ architecture² is an extension of the x86 architecture. Any processor implementing the MD64 architecture specification will also provide compatibility modes for previous descendants of the Intel 8086 architecture, including 32-bit processors such as the Intel 386, Intel Pentium, and MD K6-2 processor. Operating systems conforming to the MD64 BI may provide support for executing programs that are designed to execute in these compatibility modes. The MD64 BI does not apply to such programs; this document applies only to programs running in the “long” mode provided by the MD64 architecture.

Binaries using the MD64 instruction set may program to either a 32-bit model, in which the C data types `int`, `long` and all pointer types are 32-bit objects (ILP32); or to a 64-bit model, in which the C `int` type is 32-bits but the C `long` type and all pointer types are 64-bit objects (LP64). This specification covers both LP64 and ILP32 programming models.

Except where otherwise noted, the MD64 architecture BI follows the conventions described in the Intel386 BI. Rather than replicate the entire contents of the Intel386 BI, the MD64 BI indicates only those places where changes have been made to the Intel386 BI.

No attempt has been made to specify an BI for languages other than C. However, it is assumed that many programming languages will wish to link with code written in C, so that the BI specifications documented here apply there too.³

¹ MD64 has been previously called x86-64. The latter name is used in a number of places out of historical reasons instead of MD64.

²The architecture specification is available on the web at <http://www.x86-64.org/documentation>.

³See section 9.1 for details on C++ BI.

Chapter 2

Software Installation

This document does not specify how software must be installed on an MD64 architecture machine.

Chapter 3

Low Level System Information

3.1 Machine Interface

3.1.1 Processor Architecture

Any program can expect that an MD64 processor implements the baseline features mentioned in table 3.1. Most feature names correspond to CPUID bits, as described in the processor manual. Exceptions are OSFXSR and SCE, which are controlled by bits in the %cr4 register and the I 32_EFER MSR.

In addition to the MD64 baseline architecture, several micro-architecture levels implemented by later CPU modules have been defined, starting at level x86-64-v2. These levels are intended to support loading of optimized implementations on those systems that are compatible with them (see below). The levels are cumulative in the sense that features from previous levels are implicitly included in later levels.

Levels x86-64-v3 and x86-64-v4 are only available if the corresponding features have been fully enabled. This means that the system must pass the full sequence of checks in the processor manual for these features, including verification of the XCR0 feature flags obtained using `xgetbv`.

Recommended Uses of Micro-architecture Levels

The names for the micro-architecture levels of table 3.1 are expected to be used as directory names (to be searched by the dynamic linker, based on the levels supported by the current CPU), and by compilers, to select groups of CPU features. Distributions may also specify that they require CPU support for a certain level.

For example, to select the second level, x86-64-v3, a programmer would build a

Table 3.1: Micro- rchitecture Levels

Level Name	CPU Feature	Example instruction
(baseline)	CMOV CX8 FPU FXSR MMX OSFXSR SCE SSE SSE2	cmov cmpxchg8b fld fxsave emms fxsave syscall cvtsi2ss cvtpi2pd
x86-64-v2	CMPXCHG16B L HF-S HF POPCNT SSE3 SSE4_1 SSE4_2 SSSE3	cmpxchg16b lahf popcnt addsubpd blendpd pcmpestrd phaddd
x86-64-v3	VX VX2 BMI1 BMI2 F16C FM LZCNT MOVBE OSXSV	vzeroall vpermd andn bzhi vcvtph2ps vfmadd132pd lzcnt movbe xgetbv
x86-64-v4	VX512F VX512BW VX512CD VX512DQ VX512VL	kmovw vdbpsadbw vplzcntd vpmullq n/a

shared object with the `-march=x86-64-v3` GCC flag. The resulting shared object needs to be installed into the directory `/usr/lib64/glibc-hwcaps/x86-64-v3` or `/usr/lib/x86_64-linux-gnu/glibc-hwcaps/x86-64-v3` (in case of distributions with a multi-arch file system layout). In order to support systems that only implement the K8 baseline, a fallback implementation must be installed into the default locations, `/usr/lib64` or `/usr/lib/x66_64-linux-gnu`. It has to be built with `-march=x86-64` (the upstream GCC default). If this guideline is not followed, loading the library will fail on systems that do not support the level for which the optimized shared object was built.

Shared objects that are installed under the matching `glibc-hwcaps` subdirectory can use the CPU features for this level and earlier levels without further detection logic. Run-time detection for other CPU features not listed in this section, or listed only under later levels, is still required (even if all current CPUs implement certain CPU features together).

If a distribution requires support for a certain level, they build everything with the appropriate `-march=` option and install the built binaries in the default file system locations. When targeting such distributions, programmers can build their binaries with the same `-march=` option and install them into the default locations. Optimized shared objects for later levels can still be installed into subdirectories with the appropriate name.

3.1.2 Data Representation

Within this specification, the term *byte* refers to a 8-bit object, the term *two byte* refers to a 16-bit object, the term *four byte* refers to a 32-bit object, the term *eight byte* refers to a 64-bit object, and the term *sixteen byte* refers to a 128-bit object.¹

Fundamental Types

Figure 3.1 shows the correspondence between ISO C’s scalar types and the processor’s. `__int128`, `_Float16`, `__float80`, `__float128`, `__m64`, `__m128`, `__m256` and `__m512` types are optional.

The `__float128` type uses a 15-bit exponent, a 113-bit mantissa (the high order significant bit is implicit) and an exponent bias of 16383.²

¹The Intel386 BI uses the term *halfword* for a 16-bit object, the term *word* for a 32-bit object, the term *doubleword* for a 64-bit object. But most I -32 processor specific documentation define a *word* as a 16-bit object, a *doubleword* as a 32-bit object, a *quadword* as a 64-bit object and a *double quadword* as a 128-bit object.

²Initial implementations of the MD64 architecture are expected to support operations on the `__float128` type only via software emulation.

Figure 3.1: Scalar Types

Type	C	<code>sizeof</code>	Alignment (bytes)	MD64 Architecture
Integral	<code>_Bool</code>	1	1	boolean
	<code>char</code>	1	1	signed byte
	<code>signed char</code>			
	<code>unsigned char</code>	1	1	unsigned byte
	<code>signed short</code>	2	2	signed twobyte
	<code>unsigned short</code>	2	2	unsigned twobyte
	<code>signed int</code>	4	4	signed fourbyte
	<code>enum</code>			
	<code>unsigned int</code>	4	4	unsigned fourbyte
	<code>signed long (LP64)</code>	8	8	signed eightbyte
	<code>unsigned long (LP64)</code>	8	8	unsigned eightbyte
	<code>signed long (ILP32)</code>	4	4	signed fourbyte
	<code>unsigned long (ILP32)</code>	4	4	unsigned fourbyte
	<code>signed long long</code>	8	8	signed eightbyte
	<code>unsigned long long</code>	8	8	unsigned eightbyte
Pointer	<code>ny-type * (LP64)</code>	8	8	unsigned eightbyte
	<code>ny-type (*) () (LP64)</code>			
	<code>ny-type * (ILP32)</code>	4	4	unsigned fourbyte
	<code>ny-type (*) () (ILP32)</code>			
Floating-point	<code>_Float16</code>	2	2	16-bit (IEEE-754)
	<code>float</code>	4	4	single (IEEE-754)
	<code>double</code>	8	8	double (IEEE-754)
	<code>__float80</code>	16	16	80-bit extended (IEEE-754)
	<code>long double</code>	16	16	80-bit extended (IEEE-754)
	<code>__float128</code>	16	16	128-bit extended (IEEE-754)
Decimal-floating-point	<code>long double</code>	16	16	128-bit extended (IEEE-754)
	<code>_Decimal32</code>	4	4	32bit BID (IEEE-754R)
	<code>_Decimal64</code>	8	8	64bit BID (IEEE-754R)
Packed	<code>_Decimal128</code>	16	16	128bit BID (IEEE-754R)
	<code>__m64</code>	8	8	MMX and 3DNow!
	<code>__m128</code>	16	16	SSE and SSE-2
	<code>__m256</code>	32	32	VX
	<code>__m512</code>	64	64	VX-512

This type is called `bool` in C++.

These types are optional.

C++ and some implementations of C permit enums larger than an `int`. The underlying type is bumped to an `unsigned int`, `long int` or `unsigned long int`, in that order.

The `long long`, `signed long long`, `unsigned long long` and `double` types have 4-byte alignment in the Intel386 BI.

The `long double` type is 128-bit, the same as the `__float128` type, on the `ndroid™` platform. More information on the `ndroid™` platform is available from <http://www.android.com/>.

The `_Float16` type, from ISO/IEC TS 18661-3:2015, is optional.

The `long double` type uses a 15 bit exponent, a 64-bit mantissa with an explicit high order significant bit and an exponent bias of 16383.³ Although a `long double` requires 16 bytes of storage, only the first 10 bytes are significant. The remaining six bytes are tail padding, and the contents of these bytes are undefined.

The `_int128` type is stored in little-endian order in memory, i.e., the 64 low-order bits are stored at a lower address than the 64 high-order bits.

The value of `_alignof(max_align_t)` is 16.

`null` pointer (for all types) has the value zero.

The type `size_t` is defined as `unsigned long` for LP64 and `unsigned int` for ILP32.

Booleans, when stored in a memory object, are stored as single byte objects the value of which is always 0 (`false`) or 1 (`true`). When stored in integer registers (except for passing as arguments), all 8 bytes of the register are significant; any nonzero value is considered `true`.

Like the Intel386 architecture, the MD64 architecture in general does not require all data accesses to be properly aligned. Misaligned data accesses are slower than aligned accesses but otherwise behave identically. The only exceptions are that `_m128`, `_m256` and `_m512` must always be aligned properly.

Aggregates and Unions

Structures and unions assume the alignment of their most strictly aligned component. Each member is assigned to the lowest available offset with the appropriate alignment. The size of any object is always a multiple of the object's alignment.

An array uses the same alignment as its elements, except that a local or global array variable of length at least 16 bytes or a C99 variable-length array variable always has alignment of at least 16 bytes.⁴

Structure and union objects can require padding to meet size and alignment constraints. The contents of any padding is undefined.

Bit-Fields

C struct and union definitions may include bit-fields that define integral values of a specified size.

The BI does not permit bit-fields having the type `_m64`, `_m128`, `_m256` or `_m512`. Programs using bit-fields of these types are not portable.

³This type is the x87 double extended precision data type.

⁴The alignment requirement allows the use of SSE instructions when operating on the array. The compiler cannot in general calculate the size of a variable-length array (`VL`), but it is expected that most `VL`'s will require at least 16 bytes, so it is logical to mandate that `VL`'s have at least a 16-byte alignment.

Figure 3.2: Bit-Field Ranges

Bit-field Type	Width w	Range		
signed char	1 to 8	2 to 2	1	
char		0 to 2	1	
unsigned char		0 to 2	1	
signed short	1 to 16	2 to 2	1	
short		0 to 2	1	
unsigned short		0 to 2	1	
signed int	1 to 32	2 to 2	1	
int		0 to 2	1	
unsigned int		0 to 2	1	
signed long (LP64)	1 to 64	2 to 2	1	
long (LP64)		0 to 2	1	
unsigned long (LP64)		0 to 2	1	
long (ILP32)	1 to 32	0 to 2	1	
unsigned long (ILP32)		0 to 2	1	
signed long long	1 to 64	2 to 2	1	
long long		0 to 2	1	
unsigned long long		0 to 2	1	

Bit-fields that are neither signed nor unsigned always have non-negative values. Although they may have type char, short, int, or long (which can have negative values), these bit-fields have the same range as a bit-field of the same size with the corresponding unsigned type. Bit-fields obey the same size and alignment rules as other structure and union members.

Iso:

- bit-fields are allocated from right to left
- bit-fields must be contained in a storage unit appropriate for its declared type
- bit-fields may share a storage unit with other struct / union members

Unnamed bit-fields' types do not affect the alignment of a structure or union.

3.2 Function Calling Sequence

This section describes the standard function calling sequence, including stack frame layout, register usage, parameter passing and so on.

The standard calling sequence requirements apply only to global functions. Local functions that are not reachable from other compilation units may use different conventions. Nevertheless, it is recommended that all functions use the standard calling sequence when possible.

3.2.1 Registers

The MD64 architecture provides 16 general purpose 64-bit registers. In addition the architecture provides 16 SSE registers, each 128 bits wide and 8 x87 floating point registers, each 80 bits wide. Each of the x87 floating point registers may be referred to in MMX/3DNow! mode as a 64-bit register. All of these registers are global to all procedures active for a given thread.

Intel VX (Advanced Vector Extensions) provides 16 256-bit wide VX registers ($\%ymm_0 - \%ymm_{15}$). The lower 128-bits of $\%ymm_0 - \%ymm_{15}$ are aliased to the respective 128b-bit SSE registers ($\%xmm_0 - \%xmm_{15}$). Intel VX-512 provides 32 512-bit wide SIMD registers ($\%zmm_0 - \%zmm_{31}$). The lower 128-bits of $\%zmm_0 - \%zmm_{31}$ are aliased to the respective 128b-bit SSE registers ($\%xmm_0 - \%xmm_{31}$ ⁵). The lower 256-bits of $\%zmm_0 - \%zmm_{31}$ are aliased to the respective 256-bit VX registers ($\%ymm_0 - \%ymm_{31}$ ⁶). For purposes of parameter passing and function return, $\%xmm_N$, $\%ymm_N$ and $\%zmm_N$ refer to the same register. Only one of them can be used at the same time. We use vector register to refer to either SSE, VX or VX-512 register. In addition, Intel VX-512 also provides 8 vector mask registers ($\%k_0 - \%k_7$), each 64-bit wide.

Intel Advanced Matrix Extensions (Intel MX) is a programming paradigm consisting of two components: a set of 2-dimensional registers (tiles) representing sub-arrays from a larger 2-dimensional memory image, and accelerators able to operate on tiles. Capability of Intel MX implementation is enumerated by palettes. Two palettes are supported: palette 0 represents the initialized state and palette 1 consists of 8 tile registers ($\%tmm_0 - \%tmm_7$) of up to 1 KB size, which is controlled by a tile control register.

This subsection discusses usage of each register. Registers $\%rbp$, $\%rbx$ and $\%r12$ through $\%r15$ “belong” to the calling function and the called function is required to preserve their values. In other words, a called function must preserve these registers’ values for its

⁵ $\%xmm_{15} - \%xmm_{31}$ are only available with Intel VX-512.

⁶ $\%ymm_{15} - \%ymm_{31}$ are only available with Intel VX-512.

Figure 3.3: Stack Frame with Base Pointer

Position	Contents	Frame
$8n+16$ (%rbp)	memory argument eightbyte n ...	Previous
16 (%rbp)	memory argument eightbyte 0	
8 (%rbp)	return address	
0 (%rbp)	previous %rbp value	
-8 (%rbp)	unspecified ...	Current
0 (%rsp)	variable size	
-128 (%rsp)	red zone	

caller. Remaining registers “belong” to the called function.⁷ If a calling function wants to preserve such a register value across a function call, it must save the value in its local stack frame.

The CPU shall be in x87 mode upon entry to a function. Therefore, every function that uses the *MMX* registers is required to issue an `emms` or `femms` instruction after using *MMX* registers, before returning or calling another function.⁸ The direction flag `DF` in the `%rFL GS` register must be clear (set to “forward” direction) on function entry and return. Other user flags have no specified role in the standard calling sequence and are *not* preserved across calls.

The control bits of the `MXCSR` register are callee-saved (preserved across calls), while the status bits are caller-saved (not preserved). The x87 status word register is caller-saved, whereas the x87 control word is callee-saved.

3.2.2 The Stack Frame

In addition to registers, each function has a frame on the run-time stack. This stack grows downwards from high addresses. Figure 3.3 shows the stack organization.

The end of the input argument area shall be aligned on a 16 (32 or 64, if `__m256` or `__m512` is passed on stack) byte boundary.⁹ In other words, the stack needs to be 16 (32

⁷Note that in contrast to the Intel386 BI, `%rdi`, and `%rsi` belong to the called function, not the caller.

⁸If x87 registers are caller-saved, so callees that make use of the *MMX* registers may use the faster `femms` instruction.

⁹The maximum aligned boundary is the maximum alignment of all variables passed on stack. In C11,

or 64) byte aligned immediately before the call instruction is executed. Once control has been transferred to the function entry point, i.e. immediately after the return address has been pushed, `%rsp` points to the return address, and the value of `%rsp + 8` is a multiple of 16 (32 or 64).¹⁰

The 128-byte area beyond the location pointed to by `%rsp` is considered to be reserved and shall not be modified by signal or interrupt handlers.¹¹ Therefore, functions may use this area for temporary data that is not needed across function calls. In particular, leaf functions may use this area for their entire stack frame, rather than adjusting the stack pointer in the prologue and epilogue. This area is known as the red zone.

3.2.3 Parameter Passing

After the argument values have been computed, they are placed either in registers or pushed on the stack. The way how values are passed is described in the following sections.

Definitions We first define a number of classes to classify arguments. The classes are corresponding to MD64 register classes and defined as:

INTEGER This class consists of integral types that fit into one of the general purpose registers.

SSE The class consists of types that fit into a vector register.

SSEUP The class consists of types that fit into a vector register and can be passed and returned in the upper bytes of it.

X87, X87UP These classes consists of types that will be returned via the x87 FPU.

COMPLEX_X87 This class consists of types that will be returned via the x87 FPU.

NO_CL SS This class is used as initializer in the algorithms. It will be used for padding and empty structures and unions.

MEMORY This class consists of types that will be passed and returned in memory via the stack.

variable of type `typedef struct { __lignas (512) int i; } var_t;` is aligned to 512 bytes.

¹⁰The conventional use of `%rbp` as a frame pointer for the stack frame may be avoided by using `%rsp` (the stack pointer) to index into the stack frame. This technique saves two instructions in the prologue and epilogue and makes one additional general-purpose register (`%rbp`) available.

¹¹Locations within 128 bytes can be addressed using one-byte displacements.

Classification The size of each argument gets rounded up to eightbytes.¹²

The basic types are assigned their natural classes:

- Arguments of types (signed and unsigned) `_Bool`, `char`, `short`, `int`, `long`, `long long`, and pointers are in the INTEGER class.
- Arguments of types `_Float16`, `float`, `double`, `_Decimal32`, `_Decimal64` and `__m64` are in class SSE.
- Arguments of types `__float128`, `_Decimal128` and `__m128` are split into two halves. The least significant ones belong to class SSE, the most significant one to class SSEUP.
- Arguments of type `__m256` are split into four eightbyte chunks. The least significant one belongs to class SSE and all the others to class SSEUP.
- Arguments of type `__m512` are split into eight eightbyte chunks. The least significant one belongs to class SSE and all the others to class SSEUP.
- The 64-bit mantissa of arguments of type `long double` belongs to class X87, the 16-bit exponent plus 6 bytes of padding belongs to class X87UP.
- Arguments of type `__int128` offer the same operations as INTEGERS, yet they do not fit into one general purpose register but require two registers. For classification purposes `__int128` is treated as if it were implemented as:

```
typedef struct {
    long low, high;
} __int128;
```

with the exception that arguments of type `__int128` that are stored in memory must be aligned on a 16-byte boundary.

- Arguments of `complex T` where `T` is one of the types `_Float16`, `float`, `double` or `__float128` are treated as if they are implemented as:

```
struct complexT {
    T real;
    T imag;
};
```

¹²Therefore the stack will always be eightbyte aligned.

- variable of type `complex long double` is classified as type `COMPLEX_X87`.

The classification of aggregate (structures and arrays) and union types works as follows:

1. If the size of an object is larger than eight eightbytes, or it contains unaligned fields, it has class `MEMORY`¹³.
2. If a C++ object is non-trivial for the purpose of calls, as specified in the C++ BI¹⁴, it is passed by invisible reference (the object is replaced in the parameter list by a pointer that has class `INTEGER`)¹⁵.
3. If the size of the aggregate exceeds a single eightbyte, each is classified separately. Each eightbyte gets initialized to class `NO_CL_SS`.
4. Each field of an object is classified recursively so that always two fields are considered. The resulting class is calculated according to the classes of the fields in the eightbyte:
 - (a) If both classes are equal, this is the resulting class.
 - (b) If one of the classes is `NO_CL_SS`, the resulting class is the other class.
 - (c) If one of the classes is `MEMORY`, the result is the `MEMORY` class.
 - (d) If one of the classes is `INTEGER`, the result is the `INTEGER`.
 - (e) If one of the classes is `X87`, `X87UP`, `COMPLEX_X87` class, `MEMORY` is used as class.
 - (f) Otherwise class `SSE` is used.
5. Then a post merger cleanup is done:

¹³The post merger clean up described later ensures that, for the processors that do not support the `_m256` type, if the size of an object is larger than two eightbytes and the first eightbyte is not SSE or any other eightbyte is not SSEUP, it still has class `MEMORY`. This in turn ensures that for processors that do support the `_m256` type, if the size of an object is four eightbytes and the first eightbyte is SSE and all other eightbytes are SSEUP, it can be passed in a register. This also applies to the `_m512` type. That is for processors that support the `_m512` type, if the size of an object is eight eightbytes and the first eightbyte is SSE and all other eightbytes are SSEUP, it can be passed in a register, otherwise, it will be passed in memory.

¹⁴See section 9.1 for details on C++ BI.

¹⁵An object whose type is non-trivial for the purpose of calls cannot be passed by value because such objects must have the same address in the caller and the callee. Similar issues apply when returning an object from a function. See C++17 [class.temporary] paragraph 3.

- (a) If one of the classes is MEMORY, the whole argument is passed in memory.
- (b) If X87UP is not preceded by X87, the whole argument is passed in memory.
- (c) If the size of the aggregate exceeds two eightbytes and the first eightbyte isn't SSE or any other eightbyte isn't SSEUP, the whole argument is passed in memory.
- (d) If SSEUP is not preceded by SSE or SSEUP, it is converted to SSE.

Passing Once arguments are classified, the registers get assigned (in left-to-right order) for passing as follows:

1. If the class is MEMORY, pass the argument on the stack at an address respecting the arguments alignment (which might be more than its natural alignment).
2. If the class is INTEGER, the next available register of the sequence `%rdi, %rsi, %rdx, %rcx, %r8` and `%r9` is used¹⁶.
3. If the class is SSE, the next available vector register is used, the registers are taken in the order from `%xmm0` to `%xmm7`.
4. If the class is SSEUP, the eightbyte is passed in the next available eightbyte chunk of the last used vector register.
5. If the class is X87, X87UP or COMPLEX_X87, it is passed in memory.

When a value of type `_Bool` is returned or passed in a register or on the stack, bit 0 contains the truth value and bits 1 to 7 shall be zero¹⁷.

If there are no registers available for any eightbyte of an argument, the whole argument is passed on the stack. If registers have already been assigned for some eightbytes of such an argument, the assignments get reverted.

¹⁶Note that `%r11` is neither required to be preserved, nor is it used to pass arguments. Making this register available as scratch register means that code in the PLT need not spill any registers when computing the address to which control needs to be transferred. `%al` is used to indicate the number of vector arguments passed to a function requiring a variable number of arguments. `%r10` is used for passing a function's static chain pointer.

¹⁷Other bits are left unspecified, hence the consumer side of those values can rely on it being 0 or 1 when truncated to 8 bit.

Figure 3.4: Register Usage

Register	Usage	callee saved
%rax	temporary register; with variable arguments passes information about the number of vector registers used; 1 st return register	No
%rbx	callee-saved register	Yes
%rcx	used to pass 4 th integer argument to functions	No
%rdx	used to pass 3 ^r argument to functions; 2 ⁿ return register	No
%rsp	stack pointer	Yes
%rbp	callee-saved register; optionally used as frame pointer	Yes
%rsi	used to pass 2 ⁿ argument to functions	No
%rdi	used to pass 1 st argument to functions	No
%r8	used to pass 5 th argument to functions	No
%r9	used to pass 6 th argument to functions	No
%r10	temporary register, used for passing a function's static chain pointer	No
%r11	temporary register	No
%r12–r14	callee-saved registers	Yes
%r15	callee-saved register; optionally used as GOT base pointer	Yes
%xmm0–%xmm1	used to pass and return floating point arguments	No
%xmm2–%xmm7	used to pass floating point arguments	No
%xmm8–%xmm15	temporary registers	No
%tmm0–%tmm7	temporary registers	No
%mm0–%mm7	temporary registers	No
%k0–%k7	temporary registers	No
%st0,%st1	temporary registers, used to return long double arguments	No
%st2–%st7	temporary registers	No
%fs	thread pointer	Yes
mxcsr	SSE2 control and status word	partial
x87 SW	x87 status word	No
x87 CW	x87 control word	Yes
tilecfg	Tile control register	No

Once registers are assigned, the arguments passed in memory are pushed on the stack in reversed (right-to-left¹⁸) order.

For calls that may call functions that use varargs or stdargs (prototype-less calls or calls to functions containing ellipsis (...) in the declaration) `%_al`¹⁹ is used as hidden argument to specify the number of vector registers used. The contents of `%_al` do not need to match exactly the number of registers, but must be an upper bound on the number of vector registers used and is in the range 0–8 inclusive.

When passing `__m256` or `__m512` arguments to functions that use varargs or stdarg, function prototypes must be provided. Otherwise, the run-time behavior is undefined.

Returning of Values The returning of values is done according to the following algorithm:

1. Classify the return type with the classification algorithm.
2. If the type has class MEMORY, then the caller provides space for the return value and passes the address of this storage in `%rdi` as if it were the first argument to the function. In effect, this address becomes a “hidden” first argument. This storage must not overlap any data visible to the callee through other names than this argument.

On return `%rax` will contain the address that has been passed in by the caller in `%rdi`.

3. If the class is INTEGER, the next available register of the sequence `%rax`, `%rdx` is used.
4. If the class is SSE, the next available vector register of the sequence `%xmm0`, `%xmm1` is used.
5. If the class is SSEUP, the eightbyte is returned in the next available eightbyte chunk of the last used vector register.
6. If the class is X87, the value is returned on the X87 stack in `%st0` as 80-bit x87 number.
7. If the class is X87UP, the value is returned together with the previous X87 value in `%st0`.

¹⁸Right-to-left order on the stack makes the handling of functions that take a variable number of arguments simpler. The location of the first argument can always be computed statically, based on the type of that argument. It would be difficult to compute the address of the first argument if the arguments were pushed in left-to-right order.

¹⁹Note that the rest of `%rax` is undefined, only the contents of `%_al` is defined.

8. If the class is COMPLEX_X87, the real part of the value is returned in %st0 and the imaginary part in %st1.

As an example of the register passing conventions, consider the declarations and the function call shown in Figure 3.5. The corresponding register allocation is given in Figure 3.6, the stack frame offset given shows the frame before calling the function.

Figure 3.5: Parameter Passing Example

```

typedef struct {
    int a, b;
    double d;
} structparm;
structparm s;
int e, f, g, h, i, j, k;
long double ld;
double m, n;
__m256 y;
__m512 z;

extern void func (int e, int f,
                  structparm s, int g, int h,
                  long double ld, double m,
                  __m256 y,
                  __m512 z,
                  double n, int i, int j, int k);

func (e, f, s, g, h, ld, m, y, z, n, i, j, k);

```

Figure 3.6: Register Allocation Example

General Purpose Registers	Floating Point Registers	Stack Frame Offset
%rdi: e	%xmm0: s.d	0: ld
%rsi: f	%xmm1: m	16: j
%rdx: s.a,s.b	%ymm2: y	24: k
%rcx: g	%zmm3: z	
%r8: h	%xmm4: n	
%r9: i		

3.3 Operating System Interface

3.3.1 Exception Interface

As the MD64 manuals describe, the processor changes mode to handle *exceptions*, which may be synchronous, floating-point/coprocessor or asynchronous. Synchronous and floating-point/coprocessor exceptions, being caused by instruction execution, can be explicitly generated by a process. This section, therefore, specifies those exception types with defined behavior. The MD64 architecture classifies exceptions as *faults*, *traps*, and *aborts*. See the Intel386 BI for more information about their differences.

Hardware Exception Types

The operating system defines the correspondence between hardware exceptions and the signals specified by `signal(B _OS)` as shown in table 3.2. Contrary to the i386 architecture, the MD64 does not define any instructions that generate a bounds check fault in long mode.

3.3.2 Virtual Address Space

Although the MD64 architecture uses 64-bit pointers, implementations are only required to handle 48-bit addresses. Therefore, conforming processes may only use addresses from `0x00000000 00000000` to `0x00007ffff ffffffff`²⁰.

Processes begin with three logical segments, commonly called text, data, and stack. Use of shared libraries add other segments and a process may dynamically create segments.

3.3.3 Page Size

Systems are permitted to use any power-of-two page size between 4KB and 64KB, inclusive.

3.3.4 Virtual Address Assignments

Conceptually processes have the full address space available. In practice, however, several factors limit the size of a process.

- The system reserves a configuration dependent amount of virtual space.

²⁰0x0000ffff ffffffff is not a canonical address and cannot be used.

Table 3.2: Hardware Exceptions and Signals

Number	Exception name	Signal
0	divide error fault	SIGFPE
1	single step trap/fault	SIGTR P
2	non-maskable interrupt	none
3	breakpoint trap	SIGTR P
4	overflow trap	SIGSEGV
5	(reserved)	
6	invalid opcode fault	SIGILL
7	no coprocessor fault	SIGFPE
8	double fault abort	none
9	coprocessor overrun abort	SIGSEGV
10	invalid TSS fault	none
11	segment no present fault	none
12	stack exception fault	SIGSEGV
13	general protection fault/abort	SIGSEGV
14	page fault	SIGSEGV
15	(reserved)	
16	coprocessor error fault	SIGFPE
other	(unspecified)	SIGILL

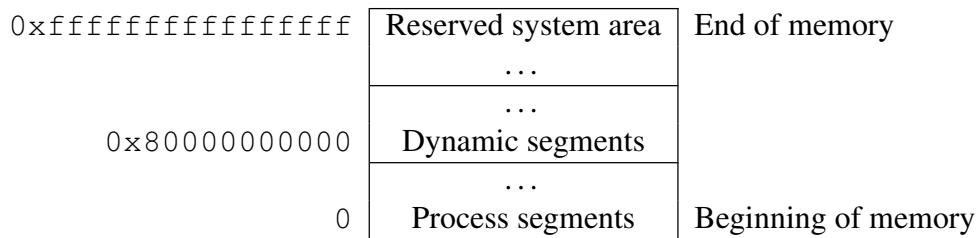
Table 3.3: Floating-Point Exceptions

Code	Reason
FPE_FLTDIV	floating-point divide by zero
FPE_FLTOVF	floating-point overflow
FPE_FLTUND	floating-point underflow
FPE_FLTRES	floating-point inexact result
FPE_FLTINV	invalid floating-point operation

- The system reserves a configuration dependent amount of space per process.
- A process whose size exceeds the system's available combined physical memory and secondary storage cannot run. Although some physical memory must be present to run any process, the system can execute processes that are bigger than physical memory, paging them to and from secondary storage. Nonetheless, both physical memory and secondary storage are shared resources. System load, which can vary from one program execution to the next, affects the available amount.

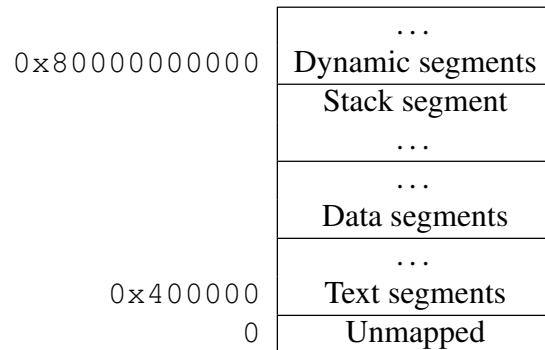
Programs that dereference null pointers are erroneous and a process should not expect 0x0 to be a valid address.

Figure 3.7: Virtual Address Configuration



Although applications may control their memory assignments, the typical arrangement appears in figure 3.8.

Figure 3.8: Conventional Segment Arrangements



3.4 Process Initialization

3.4.1 Initial Stack and Register State

Special Registers

The MD64 architecture defines floating point instructions. At process startup the two floating point units, SSE2 and x87, both have all floating-point exception status flags cleared. The status of the control words is as defined in tables 3.4 and 3.5.

Table 3.4: x87 Floating-Point Control Word

Field	Value	Note
RC	0	Round to nearest
PC	11	Double extended precision
PM	1	Precision masked
UM	1	Underflow masked
OM	1	Overflow masked
ZM	1	Zero divide masked
DM	1	De-normal operand masked
IM	1	Invalid operation masked

Table 3.5: MXCSR Status Bits

Field	Value	Note
FZ	0	Do not flush to zero
RC	0	Round to nearest
PM	1	Precision masked
UM	1	Underflow masked
OM	1	Overflow masked
ZM	1	Zero divide masked
DM	1	De-normal operand masked
IM	1	Invalid operation masked
D Z	0	De-normals are not zero

The `rFL GS` register contains the system flags, such as the direction flag and the carry flag. The low 16 bits (`FL GS` portion) of `rFL GS` are accessible by application software. The state of them at process initialization is shown in table 3.6.

Table 3.6: `rFL GS` Bits

Field	Value	Note
DF	0	Direction forward
CF	0	No carry
PF	0	Even parity
F	0	No auxiliary carry
ZF	0	No zero result
SF	0	Unsigned result
OF	0	No overflow occurred

Stack State

This section describes the machine state that `exec` (B _OS) creates for new processes. Various language implementations transform this initial program state to the state required by the language standard.

For example, a C program begins executing at a function named `main` declared as:

```
extern int main ( int argc , char *argv[ ] , char* envp[ ] );
```

where

argc is a non-negative argument count

argv is an array of argument strings, with `argv[argc] == 0`

envp is an array of environment strings, terminated by a null pointer.

When `main()` returns its value is passed to `exit()` and if that has been over-ridden and returns, `_exit()` (which must be immune to user interposition).

The initial state of the process stack, i.e. when `_start` is called is shown in figure 3.9.

Figure 3.9: Initial Process Stack

Purpose	Start address	Length
Unspecified	High addresses	
Information block, including argument strings, environment strings, auxiliary information ...		varies
Unspecified		
Null auxiliary vector entry		1 eightbyte
auxiliary vector entries ...		2 eightbytes each
0		eightbyte
Environment pointers ...		1 eightbyte each
0	8+8*argc+%rsp	eightbyte
rgument pointers	8+%rsp	argc eightbytes
rgument count	%rsp	eightbyte
Undefined	Low addresses	

rgument strings, environment strings, and the auxiliary information appear in no specific order within the information block and they need not be compactly allocated.

Only the registers listed below have specified values at process entry:

- %rbp** The content of this register is unspecified at process initialization time, but the user code should mark the deepest stack frame by setting the frame pointer to zero.
- %rsp** The stack pointer holds the address of the byte with lowest address which is part of the stack. It is guaranteed to be 16-byte aligned at process entry.
- %rdx** a function pointer that the application should register with `atexit(B _OS)`.

It is unspecified whether the data and stack segments are initially mapped with execute permissions or not. Applications which need to execute code on the stack or data segments should take proper precautions, e.g., by calling `mprotect()`.

3.4.2 Thread State

New threads inherit the floating-point state of the parent thread and the state is private to the thread thereafter.

3.4.3 Auxiliary Vector

The auxiliary vector is an array of the following structures (ref. figure 3.10), interpreted according to the `a_type` member.

Figure 3.10: `auxv_t` Type Definition

```
typedef struct
{
    int a_type;
    union {
        long a_val;
        void *a_ptr;
        void (*a_fnc)();
    } a_un;
} auxv_t;
```

The MD64 BI uses the auxiliary vector types defined in figure 3.11.

Figure 3.11: Auxiliary Vector Types

Name	Value	a_un
T_NULL	0	ignored
T_IGNORE	1	ignored
T_EXECFD	2	a_val
T_PHDR	3	a_ptr
T_PHENT	4	a_val
T_PHNUM	5	a_val
T_P_GESZ	6	a_val
T_B_SE	7	a_ptr
T_FL_GS	8	a_val
T_ENTRY	9	a_ptr
T_NOTEFL	10	a_val
T_UID	11	a_val
T_EUID	12	a_val
T_GID	13	a_val
T_EGID	14	a_val
T_PL_TFORM	15	a_ptr
T_HWC_P	16	a_val
T_CLKTCK	17	a_val
T_SECURE	23	a_val
T_B_SE_PL_TFORM	24	a_ptr
T_R_NDOM	25	a_ptr
T_HWC_P2	26	a_val
T_EXECFN	31	a_ptr

T_NULL The auxiliary vector has no fixed length; instead its last entry's `a_type` member has this value.

T_IGNORE This type indicates the entry has no meaning. The corresponding value of `a_un` is undefined.

T_EXECFD At process creation the system may pass control to an interpreter program. When this happens, the system places either an entry of type `T_EXECFD` or one of type `T_PHDR` in the auxiliary vector. The entry for type `T_EXECFD` uses the `a_val`

member to contain a file descriptor open to read the application program’s object file.

T_PHDR The system may create the memory image of the application program before passing control to the interpreter program. When this happens, the `a_ptr` member of the `T_PHDR` entry tells the interpreter where to find the program header table in the memory image.

T_PHENT The `a_val` member of this entry holds the size, in bytes, of one entry in the program header table to which the `T_PHDR` entry points.

T_PHNUM The `a_val` member of this entry holds the number of entries in the program header table to which the `T_PHDR` entry points.

T_P GESZ If present, this entry’s `a_val` member gives the system page size, in bytes.

T_B SE The `a_ptr` member of this entry holds the base address at which the interpreter program was loaded into memory. See “Program Header” in the System V BI for more information about the base address.

T_FL GS If present, the `a_val` member of this entry holds one-bit flags. Bits with undefined semantics are set to zero.

T_ENTRY The `a_ptr` member of this entry holds the entry point of the application program to which the interpreter program should transfer control.

T_NOTEELF The `a_val` member of this entry is non-zero if the program is in another format than ELF.

T_UID The `a_val` member of this entry holds the real user id of the process.

T_EUID The `a_val` member of this entry holds the effective user id of the process.

T_GID The `a_val` member of this entry holds the real group id of the process.

T_EGID The `a_val` member of this entry holds the effective group id of the process.

T_PL TFORM The `a_ptr` member of this entry points to a string containing the platform name.

T_HWC P The `a_val` member of this entry contains an bitmask of CPU features. It mask to the value returned by CPUID 1.EDX.

T_CLKTCK The `a_val` member of this entry contains the frequency at which times() increments.

T_SECURE The `a_val` member of this entry contains one if the program is in secure mode (for example started with suid). Otherwise zero.

T_B SE_PL TFORM The `a_ptr` member of this entry points to a string identifying the base architecture platform (which may be different from the platform).

T_R NDOM The `a_ptr` member of this entry points to 16 securely generated random bytes.

T_HWC P2 The `a_val` member of this entry contains the extended hardware feature mask. Currently it is 0, but may contain additional feature bits in the future.

T_EXECFN The `a_ptr` member of this entry is a pointer to the file name of the executed program.

3.5 Coding Examples

This section discusses example code sequences for fundamental operations such as calling functions, accessing static objects, and transferring control from one part of a program to another. Unlike previous material, this material is not normative.

3.5.1 Architectural Constraints

The MD64 architecture usually does not allow an instruction to encode arbitrary 64-bit constants as immediate operand. Most instructions accept 32-bit immediates that are sign extended to the 64-bit ones. Additionally the 32-bit operations with register destinations implicitly perform zero extension making loads of 64-bit immediates with upper half set to 0 even cheaper.

Additionally the branch instructions accept 32-bit immediate operands that are sign extended and used to adjust the instruction pointer. Similarly an instruction pointer relative addressing mode exists for data accesses with equivalent limitations.

In order to improve performance and reduce code size, it is desirable to use different code models depending on the requirements.

Code models define constraints for symbolic values that allow the compiler to generate better code. Basically code models differ in addressing (absolute versus position independent), code size, data size and address range. We define only a small number of code models that are of general interest:

Small code model The virtual address of code executed is known at link time. Additionally all symbols are known to be located in the virtual addresses in the range from 0 to $2^3 \cdot 2^{24} - 1$ or from 0x00000000 to 0x7effffff²¹.

This allows the compiler to encode symbolic references with offsets in the range from $-2^3 \cdot 2^{24}$ to 2^{24} or from 0x80000000 to 0x01000000 directly in the sign extended immediate operands, with offsets in the range from 0 to $2^3 \cdot 2^{24}$ or from 0x00000000 to 0x7f000000 in the zero extended immediate operands and use instruction pointer relative addressing for the symbols with offsets in the range -2^{24} to 2^{24} or 0xff000000 to 0x01000000.

This is the fastest code model and we expect it to be suitable for the vast majority of programs.

Kernel code model The kernel of an operating system is usually rather small but runs in the negative half of the address space. So we define all symbols to be in the range from $2^{64} - 2^3$ to $2^{64} - 2^{24}$ or from 0xfffffff80000000 to 0xfffffff000000.

This code model has advantages similar to those of the small model, but allows encoding of zero extended symbolic references only for offsets from 2^3 to $2^3 + 2^{24}$ or from 0x80000000 to 0x81000000. The range offsets for sign extended reference changes to 0 to $2^3 + 2^{24}$ or 0x00000000 to 0x81000000.

Medium code model In the medium model, the data section is split into two parts — the data section still limited in the same way as in the small code model and the large data section having no limits except for available addressing space. The program layout must be set in a way so that large data sections (.ldata, .lrodata, .lbss) come after the text and data sections.

This model requires the compiler to use `movabs` instructions to access large static data and to load addresses into registers, but keeps the advantages of the small code model for manipulation of addresses in the small data and text sections (specially needed for branches).

By default only data larger than 65535 bytes will be placed in the large data section.

Large code model The large code model makes no assumptions about addresses and sizes of sections.

²¹The number 24 is chosen arbitrarily. It allows for all memory of objects of size up to 2^{24} or 16M bytes to be addressed directly because the base address of such objects is constrained to be less than $2^3 \cdot 2^{24}$ or $x7f \cdot 2^{24}$. Without such constraint only the base address would be accessible directly, but not any offset variant of it.

The compiler is required to use the `movabs` instruction, as in the medium code model, even for dealing with addresses inside the text section. Additionally, indirect branches are needed when branching to addresses whose offset from the current instruction pointer is unknown.

It is possible to avoid the limitation on the text section in the small and medium models by breaking up the program into multiple shared libraries, so this model is strictly only required if the text of a single function becomes larger than what the medium model allows.

Small position independent code model (PIC) Unlike the previous models, the virtual addresses of instructions and data are not known until dynamic link time. So all addresses have to be relative to the instruction pointer.

Additionally the maximum distance between a symbol and the end of an instruction is limited to $2^3 - 2^{24} - 1$ or `0x7effffff`, allowing the compiler to use instruction pointer relative branches and addressing modes supported by the hardware for every symbol with an offset in the range -2^{24} to 2^{24} or `0xffff000000` to `0x01000000`.

Medium position independent code model (PIC) This model is like the previous model, but similarly to the medium static model adds large data sections at the end of object files.

In the medium PIC model, the instruction pointer relative addressing can not be used directly for accessing large static data, since the offset can exceed the limitations on the size of the displacement field in the instruction. Instead an unwind sequence consisting of `movabs`, `lea` and `add` needs to be used.

Large position independent code model (PIC) This model is like the previous model, but makes no assumptions about the distance of symbols.

The large PIC model implies the same limitation as the medium PIC model regarding addressing of static data. Additionally, references to the global offset table and to the procedure linkage table and branch destinations need to be calculated in a similar way. Further the size of the text segment is allowed to be up to 16EB in size, hence similar restrictions apply to all address references into the text segments, including branches.

Only small code model and small position independent code model (PIC) are used in ILP32 binaries.

3.5.2 Conventions

In this document some special assembler symbols are used in the coding examples and discussion. They are:

- `name@GOT`: specifies the offset to the GOT entry for the symbol `name` from the base of the GOT.
- `name@GOTOFF`: specifies the offset to the location of the symbol `name` from the base of the GOT.
- `name@GOTPCREL`: specifies the offset to the GOT entry for the symbol `name` from the current code location.
- `name@PLT`: specifies the offset to the PLT entry of symbol `name` from the current code location.
- `name@PLTOFF`: specifies the offset to the PLT entry of symbol `name` from the base of the GOT.
- `_GLOB_L_OFFSET_T BLE_`: specifies the offset to the base of the GOT from the current code location.

3.5.3 Position-Independent Function Prologue

In the small code model all addresses (including GOT entries) are accessible via the IP-relative addressing provided by the MD64 architecture. Hence there is no need for an explicit GOT pointer and therefore no function prologue for setting it up is necessary.

In the medium and large code models a register has to be allocated to hold the address of the GOT in position-independent objects, because the MD64 IS does not support an immediate displacement larger than 32 bits.

`s %r15` is preserved across function calls, it is initialized in the function prolog to hold the GOT address²² for non-leaf functions which call other functions through the PLT. Other functions are free to use any other register. Throughout this document, `%r15` will be used in examples.

²²If, at code generation-time, it is determined that either no other functions are called (leaf functions), the called functions addresses can be resolved and are within 2GB, or no global data objects are referred to, it is not necessary to store the GOT address in `%r15` and the prolog code that initializes it may be omitted.

Figure 3.12: Position-Independent Function Prolog Code
medium model:

```
leaq    _GLOB_L_OFFSET_T BLE_(%rip),%r15 # GOTPC32 reloc
```

large model:

```
pushq   %r15          # save %r15
leaq    1f(%rip),%r11      # absolute %rip
1: movabs $_GLOB_L_OFFSET_T BLE_,%r15  # offset to the GOT (R_X86_64_GOTPC64)
   leaq    (%r11,%r15),%r15      # absolute address of the GOT
```

For the medium model the GOT pointer is directly loaded, for the large model the absolute value of `%rip` is added to the relative offset to the base of the GOT in order to obtain its absolute address (see figure 3.12).

3.5.4 Data Objects

This section describes only objects with static storage. Stack-resident objects are excluded since programs always compute their virtual address relative to the stack or frame pointers.

Because only the `movabs` instruction uses 64-bit addresses directly, depending on the code model either `%rip`-relative addressing or building addresses in registers and accessing the memory through the register has to be used.

For absolute addresses `%rip`-relative encoding can be used in the small model. In the medium model the `movabs` instruction has to be used for accessing addresses.

Position-independent code cannot contain absolute address. To access a global symbol the address of the symbol has to be loaded from the Global Offset Table. The address of the entry in the GOT can be obtained with a `%rip`-relative instruction in the small model.

Small models

Figure 3.13: Absolute Load and Store (Small Model)

<pre>extern int src[65536]; extern int dst[65536]; extern int *ptr; static int lsrc[65536]; static int ldst[65536]; static int *lptr; dst[0] = src[0]; ptr = dst[0]; *ptr = src[0]; ldst[0] = lsrc[0]; lptr = ldst; *lptr = lsrc[0];</pre>	<pre>.extern src .extern dst .extern ptr .local lsrc .comm lsrc,262144,4 .local ldst .comm ldst,262144,4 .local lptr .comm lptr,8,8 .text movl src(%rip), %eax movl %eax, dst(%rip) movq \$dst, ptr(%rip) movq ptr(%rip),%rax movl src(%rip),%edx movl %edx, (%rax) movl lsrc(%rip), %eax movl %eax, ldst(%rip) movq \$dst, lptr(%rip) movq lptr(%rip),%rax movl lsrc(%rip),%edx movl %edx, (%rax)</pre>
--	---

Figure 3.14: Position-Independent Load and Store (Small PIC Model)

extern int src[65536];	.extern src
extern int dst[65536];	.extern dst
extern int *ptr;	.extern ptr
static int lsrc[65536];	.local lsrc
static int ldst[65536];	.comm lsrc,262144,4
static int *lptr;	.local ldst
dst[0] = src[0];	.comm ldst,262144,4
ptr = dst;	.local lptr
*ptr = src[0];	.comm lptr,8,8
ldst[0] = lsrc[0];	.text
lptr = ldst;	movq src@GOTPCREL(%rip), %rax
*lptr = lsrc[0];	movl (%rax), %edx
	movq dst@GOTPCREL(%rip), %rax
	movl (%rax), %edx
	movq ptr@GOTPCREL(%rip), %rax
	movq dst@GOTPCREL(%rip), %rdx
	movq (%rdx), (%rax)
	movq ptr@GOTPCREL(%rip), %rax
	movq dst@GOTPCREL(%rip), %rdx
	movq (%rdx), (%rax)
	movq lsrc(%rip), %eax
	movl %eax, ldst(%rip)
	lea ldst(%rip), %rdx
	movq %rdx, lptr(%rip)
	movq lptr(%rip), %rax
	movl lsrc(%rip), %edx
	movl %edx, (%rax)

Medium models

Figure 3.15: Absolute Load and Store (Medium Model)

extern int src[65536]; extern int dst[65536]; extern int *ptr; static int lsrc[65536]; static int ldst[65536]; static int *lptr; dst[0] = src[0]; ptr = dst; *ptr = src[0]; ldst[0] = lsrc[0]; lptr = ldst; *lptr = lsrc[0];	.extern src .extern dst .extern ptr .local lsrc .comm lsrc,262144,4 ²³ .local ldst .comm ldst,262144,4 .local lptr .comm lptr,8,8 .text movabsl src, %eax movabsl %eax, dst movabsq \$dst,%rdx movq %rdx, ptr movq ptr(%rip),%rdx movabsl src,%eax movl %eax, (%rdx) movabsl lsrc, %eax movabsl %eax, ldst movabsq \$ldst,%rdx movabsq %rdx, lptr movq lptr(%rip),%rdx movabsl lsrc,%eax movl %eax, (%rdx)
---	--

Figure 3.16: Position-Independent Load and Store (Medium PIC Model)

extern int src[65536]; extern int dst[65536]; extern int *ptr; static int lsrc[65536]; static int ldst[65536]; static int *lptr; dst[0] = src[0]; ptr = dst; *ptr = src[0];	.extern src .extern dst .extern ptr .local lsrc .comm lsrc,262144,4 .local ldst .comm ldst,262144,4 .local lptr .comm lptr,8,8 .text movq src@GOTPCREL(%rip), %rax movl (%rax), %edx movq dst@GOTPCREL(%rip), %rax movl %edx, (%rax) movq ptr@GOTPCREL(%rip), %rax movq dst@GOTPCREL(%rip), %rdx movq %rdx, (%rax) movq ptr@GOTPCREL(%rip), %rax movq (%rax), %rdx movq src@GOTPCREL(%rip), %rax movl (%rax), %eax movl %eax, (%rdx)
---	---

Figure 3.17: Position-Independent Load and Store (Medium PIC Model), continued

ldst[0] = lsrc[0]; lptr = ldst; *lptr = lsrc[0];	movabsq lsrc@GOTOFF, %rax movl (%rax,%r15), %eax movabsq ldst@GOTOFF, %rdx movl %eax, (%rdx,%r15) movabsq ldst@GOTOFF, %rax addq %r15, %rax movq %rax, lptr(%rip) movabsq lsrc@GOTOFF, %rax movl (%rax,%r15), %eax movq lptr(%rip),%rdx movl %eax, (%rdx)
--	---

Large Models

gain, in order to access data at any position in the 64-bit addressing space, it is necessary to calculate the address explicitly²⁴, not unlike the medium code model.

Figure 3.18: Absolute Global Data Load and Store

static int src; static int dst; extern int *ptr;	Lsrc: .long Ldst: .long .extern ptr
dst = src;	movabs \$Lsrc,%rax ;R_X86_64_64 movabs \$Ldst,%rdx ;R_X86_64_64 movl (%rax),%ecx movl %ecx,(%rdx)
ptr = &dst;	movabs \$ptr,%rax ;R_X86_64_64 movabs \$Ldst,%rdx ;R_X86_64_64 movq %rdx,(%rax)
*ptr = src;	movabs \$Lsrc,%rax ;R_X86_64_64 movabs \$ptr,%rdx ;R_X86_64_64 movl (%rax),%ecx movq (%rdx),%rdx movl %ecx,(%rdx)

Figure 3.19: Faster Absolute Global Data Load and Store

*ptr = src;	movabs \$ptr,%rdx ;R_X86_64_64 movl Lsrc(%rip),%ecx movq (%rdx),%rdx movl %ecx,(%rdx)
-------------	--

For position-independent code access to both static and external global data assumes that the GOT address is stored in a dedicated register. In these examples we assume it is in %r15²⁵ (see Function Prologue):

²⁴If, at code generation-time, it is determined that a referred to global data object address is resolved within 2GB, the %rip-relative addressing mode can be used instead. See example in figure 3.19.

²⁵If, at code generation-time, it is determined that a referred to global data object address is resolved within 2GB, the %rip-relative addressing mode can be used instead. See example in figure 3.21.

Figure 3.20: Position-Independent Global Data Load and Store

static int src; static int dst; extern int *ptr;	Lsrc: .long Ldst: .long .extern ptr
dst = src;	movabs \$Lsrc@GOTOFF,%rax ;R_X86_64_GOTOFF64 movabs \$Ldst@GOTOFF,%rdx ;R_X86_64_GOTOFF64 movl (%rax,%r15),%ecx movl %ecx,(%rdx,%r15)
ptr = &dst;	movabs \$ptr@GOT,%rax ;R_X86_64_GOT64 movabs \$Ldst@GOTOFF,%rdx ;R_X86_64_GOTOFF64 movq (%rax,%r15),%rax leaq (%rdx,%r15),%rcx movq %rcx,(%rax)
*ptr = src;	movabs \$Lsrc@GOTOFF,%rax ;R_X86_64_GOTOFF64 movabs \$ptr@GOT,%rdx ;R_X86_64_GOT64 movl (%rax,%r15),%ecx movq (%rdx,%r15),%rdx movl %ecx,(%rdx)

Figure 3.21: Faster Position-Independent Global Data Load and Store

*ptr = src;	movabs \$ptr@GOT,%rdx ;R_X86_64_GOT64 movl Lsrc(%rip),%ecx movq (%rdx,%r15),%rdx movl %ecx,(%rdx)
-------------	--

3.5.5 Function Calls

Small and Medium Models

Figure 3.22: Position-Independent Direct Function Call (Small and Medium Model)

extern void function (); function ();	.globl function call function@PLT
--	--------------------------------------

Figure 3.23: Position-Independent Indirect Function Call

<pre>extern void (*ptr) (); extern void name (); ptr = name; (*ptr)();</pre>	<pre>.globl ptr, name movq ptr@GOTPCREL(%rip), %rax movq name@GOTPCREL(%rip), %rdx movq %rdx, (%rax) movq ptr@GOTPCREL(%rip), %rax call *(%rax)</pre>
---	---

Large models

It cannot be assumed that a function is within 2GB in general. Therefore, it is necessary to explicitly calculate the desired address reaching the whole 64-bit address space.

Figure 3.24: Absolute Direct and Indirect Function Call

<pre>static (*ptr) (void); extern foo (void); static bar (void); foo (); bar ();</pre>	<pre>Lptr: .quad .globl foo Lbar: ... movabs \$foo,%r11 ;R_X86_64_64 call *%r11 movabs \$Lbar,%r11 ;R_X86_64_64 call *%r11</pre>
<pre>ptr = foo; ptr = bar;</pre>	<pre>movabs \$Lptr,%rax ;R_X86_64_64 movabs \$foo,%r11 ;R_X86_64_64 movq %r11,(%rax) movabs \$Lbar,%r11 ;R_X86_64_64 movq %r11,(%rax)</pre>
<pre>(*ptr)();</pre>	<pre>movabs \$Lptr,%r11 ;R_X86_64_64 call *(%r11)</pre>

nd in the case of position-independent objects ²⁶:

²⁶See subsection “Implementation advice” for some optimizations.

Figure 3.25: Position-Independent Direct and Indirect Function Call

static (*ptr) (void); extern foo (void); static bar (void);	Lptr: .quad .globl foo Lbar: ...
foo (); bar ();	movabs \$foo@PLTOFF,%r11 ;R_X86_64_PLTOFF64 call *(%r11,%r15) movabs \$Lbar@GOTOFF,%r11 ;R_X86_64_GOTOFF64 leaq (%r11,%r15),%r11 call *%r11
ptr = foo; ptr = bar;	movabs \$Lptr@GOTOFF,%rax ;R_X86_64_GOTOFF64 movabs \$foo@PLTOFF,%r11 ;R_X86_64_PLTOFF64 leaq (%r11,%r15),%r11 movq %r11,(%rax,%r15) movabs \$Lbar@GOTOFF,%r11 ;R_X86_64_GOTOFF64 leaq (%r11,%r15),%r11 movq %r11,(%rax,%r15)
(*ptr) ();	movabs \$Lptr@GOTOFF,%r11 ;R_X86_64_GOTOFF64 call *(%r11,%r15)

Implementation advice

If, at code generation-time, certain conditions are determined, it's possible to generate faster or smaller code sequences as the large model normally requires. When:

(absolute) target of function call is within 2GB, a direct call or `%rip`-relative addressing might be used:

bar ();	call Lbar
ptr = bar;	movabs \$Lptr,%rax ;R_X86_64_64 leaq \$Lbar(%rip),%r11 movq %r11,(%rax)

(PIC) the base of GOT is within 2GB an indirect call to the GOT entry might be implemented like so:

foo ();	call * (foo@GOT) ;R_X86_64_GOTPCREL
---------	-------------------------------------

(PIC) the base of PLT is within 2GB, the PLT entry may be referred to relatively to `%rip`:

ptr = foo;	movabs \$Lptr@GOTOFF,%rax ;R_X86_64_GOTOFF64 leaq \$foo@PLT(%rip),%r11 ;R_X86_64_PLT32 movq %r11,(%rax,%r15)
------------	--

(PIC) target of function call is within 2GB and is either not global or bound locally, a direct call to the symbol may be used or it may be referred to relatively to `%rip`:

bar ();	call Lbar
ptr = bar;	movabs \$Lptr@GOTOFF,%rax ;R_X86_64_GOTOFF64 leaq \$Lbar(%rip),%r11 movq %r11,(%rax,%r15)

3.5.6 Branching

Small and Medium Models

As all labels are within 2GB no special care has to be taken when implementing branches. The full MD64 IS is usable.

Large Models

Because functions can be theoretically up to 16EB long, the maximum 32-bit displacement of conditional and unconditional branches in the MD64 IS are not enough to address the branch target. Therefore, a branch target address is calculated explicitly ²⁷. For absolute objects:

Figure 3.26: Absolute Branching Code

if (!a) { ... }	testl %eax,%eax jnz 1f movabs \$2f,%r11 ;R_X86_64_64 jmpq *%r11 1: ... 2:
goto Label; ... Label:	movabs \$Label,%r11 ;R_X86_64_64 jmpq *%r11 ... Label:

²⁷If, at code generation-time, it is determined that the target addresses are within 2GB, alternatively, branch target addresses may be calculated implicitly (see figure 3.27)

Figure 3.27: Implicit Calculation of Target Address

if (!a) { ... }	testl %eax, %eax jz 2f 1: 2:
goto Label; ... Label:	jmp Label ... Label:

For position-independent objects:

Figure 3.28: Position-Independent Branching Code

if (!a) { ... }	testl %eax, %eax jnz 1f movabs \$2f@GOTOFF,%r11 ;R_X86_64_GOTOFF64 leaq (%r11,%r15),%r11 jmpq *%r11 1: 2:
goto Label; Label:	movabs \$Label@GOTOFF,%r11 ;R_X86_64_GOTOFF64 leaq (%r11,%r15),%r11 jmpq *%r11 Label:

For absolute objects, the implementation of the `switch` statement is:

Figure 3.29: Absolute Switch Code

switch (a) { default: ... case 0: ... case 2: ... }	cmpl \$0,%eax jl .Ldefault cmpl \$2,%eax jg .Ldefault movabs \$.Ltable,%r11 ;R_X86_64_64 jmpq *(%r11,%eax,8) .section .lrodata,"aLM",@progbits,8 .align 8 .Ltable: .quad .Lcase0 ;R_X86_64_64 .quad .Ldefault ;R_X86_64_64 .quad .Lcase2 ;R_X86_64_64 .previous .Ldefault:Lcase0:Lcase2: ...
---	--

When building position-independent objects, the `switch` statement implementation changes to:

Figure 3.30: Position-Independent Switch Code

<pre> switch (a) { cmpl \$0,%eax jl .Ldefault cmpl \$2,%eax jg .Ldefault movabs \$.Ltable@GOTOFF,%r11 ;R_X86_64_GOTOFF64 leaq (%r11,%r15),%r11 movq *(%r11,%eax,8),%r11 leaq (%r11,%r15),%r11 jmpq *%r11 .section .lrodata,"aLM",@progbits,8 .align 8 .Ltable: .quad .Lcase0@GOTOFF ;R_X86_64_GOTOFF64 .quad .Ldefault@GOTOFF ;R_X86_64_GOTOFF64 .quad .Lcase2@GOTOFF ;R_X86_64_GOTOFF64 .previous default: .Ldefault: ... case 0: .Lcase0: ... case 2: .Lcase2: ... } </pre>	
---	--

28

3.5.7 Variable Argument Lists

Some otherwise portable C programs depend on the argument passing scheme, implicitly assuming that all arguments are passed on the stack, and arguments appear in increasing order on the stack. Programs that make these assumptions never have been portable, but they have worked on many implementations. However, they do not work on the MD64 architecture because some arguments are passed in registers. Portable C programs must use the header file `<stdarg.h>` in order to handle variable argument lists.

When a function taking variable-arguments is called, `%al` must be set to the total number of floating point parameters passed to the function in vector registers.²⁹

²⁸The jump-table is emitted in a different section so as to occupy cache lines without instruction bytes, thus avoiding exclusive cache subsystems to thrash.

²⁹This implies that the only legal values for `%al` when calling a function with variable-argument lists are 0 to 8 (inclusive).

When `__m256` or `__m512` is passed as variable-argument, it should always be passed on stack. Only named `__m256` and `__m512` arguments may be passed in register as specified in section 3.2.3.

Figure 3.31: Parameter Passing Example with Variable- argument List

```
int a, b;
long double ld;
double m, n;
__m256 u, y;
__m512 v, z;

extern void func (int a, double m, __m256 u, __m512 v, ...);

func (a, m, u, v, b, ld, y, z, n);
```

Figure 3.32: Register Allocation Example for Variable- argument List

General Purpose Registers	Floating Point Registers	Stack Frame Offset
%rdi: a	%xmm0: m	0: ld
%rsi: b	%ymm1: u	32: y
%rax: 3	%zmm2: v %xmm3: n	

The Register Save Area

The prologue of a function taking a variable argument list and known to call the macro `va_start` is expected to save the argument registers to the *register save area*. Each argument register has a fixed offset in the register save area as defined in the figure 3.33.

Only registers that might be used to pass arguments need to be saved. Other registers are not accessed and can be used for other purposes. If a function is known to never accept arguments passed in registers³⁰, the register save area may be omitted entirely.

The prologue should use `%a1` to avoid unnecessarily saving XMM registers. This is especially important for integer only programs to prevent the initialization of the XMM unit.

³⁰This fact may be determined either by exploring types used by the `va_arg` macro, or by the fact that the named arguments already are exhausted the argument registers entirely.

Figure 3.33: Register Save rea

Register	Offset
%rdi	0
%rsi	8
%rdx	16
%rcx	24
%r8	32
%r9	40
%xmm0	48
%xmm1	64
...	
%xmm15	288

The va_list Type

The `va_list` type is an array containing a single element of one structure containing the necessary information to implement the `va_arg` macro. The C definition of `va_list` type is given in figure 3.34.

Figure 3.34: `va_list` Type Declaration

```
typedef struct {
    unsigned int gp_offset;
    unsigned int fp_offset;
    void *overflow_arg_area;
    void *reg_save_area;
} va_list[1];
```

The `va_start` Macro

The `va_start` macro initializes the structure as follows:

reg_save_area The element points to the start of the register save area.

overflow_arg_area This pointer is used to fetch arguments passed on the stack. It is initialized with the address of the first argument passed on the stack, if any, and then always updated to point to the start of the next argument on the stack.

gp_offset The element holds the offset in bytes from `reg_save_area` to the place where the next available general purpose argument register is saved. In case all argument registers have been exhausted, it is set to the value 48 ($6 * 8$).

fp_offset The element holds the offset in bytes from `reg_save_area` to the place where the next available floating point argument register is saved. In case all argument registers have been exhausted, it is set to the value 304 ($6 * 8 + 16 * 16$).

The `va_arg` Macro

The algorithm for the generic `va_arg(l, type)` implementation is defined as follows:

1. Determine whether `type` may be passed in the registers. If not go to step 7.
2. Compute `num_gp` to hold the number of general purpose registers needed to pass `type` and `num_fp` to hold the number of floating point registers needed.
3. Verify whether arguments fit into registers. In the case:

$$l->gp_offset > 48 \quad num_gp * 8$$

or

$$l->fp_offset > 304 \quad num_fp * 16$$

go to step 7.

4. Fetch `type` from `l->reg_save_area` with an offset of `l->gp_offset` and/or `l->fp_offset`. This may require copying to a temporary location in case the parameter is passed in different register classes or requires an alignment greater than 8 for general purpose registers and 16 for XMM registers.

5. Set:

$$l->gp_offset = l->gp_offset + num_gp * 8$$

$$l->fp_offset = l->fp_offset + num_fp * 16$$

6. Return the fetched `type`.

7. Align `l->overflow_arg_area` upwards to a 16 byte boundary if alignment needed by `type` exceeds 8 byte boundary.

8. Fetch `type` from `l->overflow_arg_area`.

9. Set `l->overflow_arg_area` to:

```
l->overflow_arg_area + sizeof type)
```

10. Align `l->overflow_arg_area` upwards to an 8 byte boundary.

11. Return the fetched `type`.

The `va_arg` macro is usually implemented as a compiler builtin and expanded in simplified forms for each particular type. Figure 3.35 is a sample implementation of the `va_arg` macro.

Figure 3.35: Sample Implementation of `va_arg(l, int)`

movl	<code>l->gp_offset, %eax</code>	
cmpl	<code>\$48, %eax</code>	Is register available?
jae	<code>stack</code>	If not, use stack
leal	<code>\$8(%rax), %edx</code>	Next available register
addq	<code>l->reg_save_area, %rax</code>	address of saved register
movl	<code>%edx, l->gp_offset</code>	Update <code>gp_offset</code>
jmp	<code>fetch</code>	
stack:	<code>movq l->overflow_arg_area, %rax</code>	address of stack slot
	<code>leaq 8(%rax), %rdx</code>	Next available stack slot
	<code>movq %rdx,l->overflow_arg_area</code>	Update
fetch:	<code>movl (%rax), %eax</code>	Load argument

3.6 DW RF Definition

This section³¹ defines the Debug With Arbitrary Record Format (DW RF) debugging format for the MD64 processor family. The MD64 BI does not define a debug format. However, all systems that do implement DW RF on MD64 shall use the following definitions.

DW RF is a specification developed for symbolic, source-level debugging. The debugging information format does not favor the design of any compiler or debugger. For more information on DW RF, see *DW RF Debugging Format Standard*, available at: <http://www.dwarfstd.org/>.

³¹This section is structured in a way similar to the PowerPC ps BI

3.6.1 DW RF Release Number

The DW RF definition requires some machine-specific definitions. The register number mapping needs to be specified for the MD64 registers. In addition, starting with version 3 the DW RF specification requires processor-specific address class codes to be defined.

3.6.2 DW RF Register Number Mapping

Table 3.36³² outlines the register number mapping for the MD64 processor family.³³

3.7 Stack Unwind lgorithm

The stack frames are not self descriptive and where stack unwinding is desirable (such as for exception handling) additional unwind information needs to be generated. The information is stored in an allocatable section `.eh_frame` whose format is identical to `.debug_frame` defined by the DW RF debug information standard, see *DW RF Debugging Information Format*, with the following extensions:

Position independence In order to avoid load time relocations for position independent code, the FDE CIE offset pointer should be stored relative to the start of CIE table entry. Frames using this extension of the DW RF standard must set the CIE identifier tag to 1.

Outgoing arguments area delta To maintain the size of the temporarily allocated outgoing arguments area present on the end of the stack (when using `push` instructions), operation `GNU_RGS_SIZE (0x2e)` can be used. This operation takes a single `uleb128` argument specifying the current size. This information is used to adjust the stack frame when jumping into the exception handler of the function after unwinding the stack frame. Additionally the CIE augmentation shall contain an exact specification of the encoding used. It is recommended to use a PC relative encoding whenever possible and adjust the size according to the code model used.

CIE augmentations: The augmentation field is formed according to the augmentation field formating string stored in the CIE header.

The string may contain the following characters:

³²The table defines Return address to have a register number, even though the address is stored in `0(%rsp)` and not in a physical register.

³³This document does not define mappings for privileged registers.

Figure 3.36: DW RF Register Number Mapping

Register Name	Number	bbreviation
General Purpose Register R X	0	%rax
General Purpose Register RDX	1	%rdx
General Purpose Register RCX	2	%rcx
General Purpose Register RBX	3	%rbx
General Purpose Register RSI	4	%rsi
General Purpose Register RDI	5	%rdi
Frame Pointer Register RBP	6	%rbp
Stack Pointer Register RSP	7	%rsp
Extended Integer Registers 8-15	8-15	%r8-%r15
Return ddress R	16	
Vector Registers 0-7	17-24	%xmm0-%xmm7
Extended Vector Registers 8-15	25-32	%xmm8-%xmm15
Floating Point Registers 0-7	33-40	%st0-%st7
MMX Registers 0-7	41-48	%mm0-%mm7
Flag Register	49	%rFL GS
Segment Register ES	50	%es
Segment Register CS	51	%cs
Segment Register SS	52	%ss
Segment Register DS	53	%ds
Segment Register FS	54	%fs
Segment Register GS	55	%gs
Reserved	56-57	
FS Base address	58	%fs.base
GS Base address	59	%gs.base
Reserved	60-61	
Task Register	62	%tr
LDT Register	63	%ldtr
128-bit Media Control and Status	64	%mxcsr
x87 Control Word	65	%fcw
x87 Status Word	66	%fsw
Upper Vector Registers 16-31	67-82	%xmm16-%xmm31
Reserved	83-117	
Vector Mask Registers 0-7	118-125	%k0-%k7
Reserved	126-129	

Figure 3.37: Pointer Encoding Specification Byte

Mask	Meaning
0x1	Values are stored as <code>uleb128</code> or <code>sleb128</code> type (according to flag 0x8)
0x2	Values are stored as 2 bytes wide integers (<code>udata2</code> or <code>sdata2</code>)
0x3	Values are stored as 4 bytes wide integers (<code>udata4</code> or <code>sdata4</code>)
0x4	Values are stored as 8 bytes wide integers (<code>udata8</code> or <code>sdata8</code>)
0x8	Values are signed
0x10	Values are PC relative
0x20	Values are text section relative
0x30	Values are data section relative
0x40	Values are relative to the start of function

- z** Indicates that a `uleb128` is present determining the size of the augmentation section.
- L** Indicates the encoding (and thus presence) of an LSD pointer in the FDE augmentation.
The data field consist of single byte specifying the way pointers are encoded. It is a mask of the values specified by the table 3.37.
The default DW RF pointer encoding (direct 4-byte absolute pointers) is represented by value 0.
- R** Indicates a non-default pointer encoding for FDE code pointers. The formating is represented by a single byte in the same way as in the 'L' command.
- P** Indicates the presence and an encoding of a language personality routine in the CIE augmentation. The encoding is represented by a single byte in the same way as in the 'L' command followed by a pointer to the personality function encoded by the specified encoding.

When the augmentation is present, the first command must always be 'z' to allow easy skipping of the information.

In order to simplify manipulation of the unwind tables, the runtime library provide higher level PI to stack unwinding mechanism, for details see section 6.2.

Chapter 4

Object Files

4.1 ELF Header

4.1.1 Machine Information

Programming Model

As described in Section 1, binaries using the MD64 instruction set may program to either a 32-bit model (ILP32) or to a 64-bit model (LP64). This specification describes both binaries that use the ILP32 and the LP64 model.

File Class

For MD64 ILP32 objects, the file class value in `e_ident[EI_CL_SS]` must be `ELFCL_SS32`. For MD64 LP64 objects, the file class value must be `ELFCL_SS64`.

Data Encoding

For the data encoding in `e_ident[EI_D_T]`, MD64 objects use `ELFD_T_2LSB`.

Processor identification

Processor identification resides in the ELF headers `e_machine` member and must have the value `EM_X86_64`.¹

¹The value of this identifier is 62.

4.1.2 Number of Program Headers

The `e_phnum` member contains the number of entries in the program header table. The product of `e_phentsize` and `e_phnum` gives the table's size in bytes. If a file has no program header table, `e_phnum` holds the value zero.

If the number of program headers is greater than or equal to `PN_XNUM` (0xffff), this member has the value `PN_XNUM` (0xffff). The actual number of program header table entries is contained in the `sh_info` field of the section header at index 0. Otherwise, the `sh_info` member of the initial entry contains the value zero.

4.2 Sections

4.2.1 Section Flags

In order to allow linking object files of different code models, it is necessary to provide for a way to differentiate those sections which may hold more than 2GB from those which may not. This is accomplished by defining a processor-specific section attribute flag for `sh_flag` (see table 4.1).

Table 4.1: MD64 Specific Section Header Flag, `sh_flags`

Name	Value
<code>SHF_X86_64_L_RGE</code>	0x10000000

`SHF_X86_64_L_RGE` If an object file section does *not* have this flag set, then it may not hold more than 2GB and can be freely referred to in objects using smaller code models. Otherwise, only objects using larger code models can refer to them. For example, a medium code model object can refer to data in a section that sets this flag besides being able to refer to data in a section that does not set it; likewise, a small code model object can refer only to code in a section that does not set this flag.

4.2.2 Section types

Table 4.2: Section Header Types

sh_type name	Value
SHT_X86_64_UNWIND	0x70000001

SHT_X86_64_UNWIND This section contains unwind function table entries for stack unwinding. The contents are described in Section 4.2.4 of this document.

4.2.3 Special Sections

Table 4.3: Special sections

Name	Type	ttributes
.got	SHT_PROGBITS	SHF_LLOC+SHF_WRITE
.plt	SHT_PROGBITS	SHF_LLOC+SHF_EXECINSTR
.eh_frame	SHT_X86_64_UNWIND	SHF_LLOC

.got This section holds the global offset table.

.plt This section holds the procedure linkage table.

.eh_frame This section holds the unwind function table. The contents are described in Section 4.2.4 of this document.

The additional sections defined in table 4.4 are used by a system supporting the large code model.

Table 4.4: Additional Special Sections for the Large Code Model

Name	Type	Attributes
.lbss	SHT_NOBITS	SHF_LLOC+SHF_WRITE+SHF_X86_64_L RWE
.ldata	SHT_PROGBITS	SHF_LLOC+SHF_WRITE+SHF_X86_64_L RWE
.ldata1	SHT_PROGBITS	SHF_LLOC+SHF_WRITE+SHF_X86_64_L RWE
.lgot	SHT_PROGBITS	SHF_LLOC+SHF_WRITE+SHF_X86_64_L RWE
.plt	SHT_PROGBITS	SHF_LLOC+SHF_EXECINSTR+SHF_X86_64_L RWE
.lrodata	SHT_PROGBITS	SHF_LLOC+SHF_X86_64_L RWE
.lrodata1	SHT_PROGBITS	SHF_LLOC+SHF_X86_64_L RWE
.ltext	SHT_PROGBITS	SHF_LLOC+SHF_EXECINSTR+SHF_X86_64_L RWE

In order to enable static linking of objects using different code models, the following section ordering is suggested:

`.plt .init .fini .text .got .rodata .rodata1 .data .data1 .bss` These sections can have a combined size of up to 2GB.

`.plt .ltext .lgot .lrodata .lrodata1 .ldata .ldata1 .bss` These sections plus the above can have a combined size of up to 16EB.

4.2.4 EH_FRAME sections

The call frame information needed for unwinding the stack is output into one or more ELF sections of type SHT_X86_64_UNWIND. In the simplest case there will be one such section per object file and it will be named `.eh_frame`. A `.eh_frame` section consists of one or more subsections. Each subsection contains a CIE (Common Information Entry) followed by varying number of FDEs (Frame Descriptor Entry). A FDE corresponds to an explicit or compiler generated function in a compilation unit, all FDEs can access the CIE that begins their subsection for data. If the code for a function is not one contiguous block, there will be a separate FDE for each contiguous sub-piece.

If an object file contains C++ template instantiations there shall be a separate CIE immediately preceding each FDE corresponding to an instantiation.

Using the preferred encoding specified below, the `.eh_frame` section can be entirely resolved at link time and thus can become part of the text segment.

`EH_PE` encoding below refers to the pointer encoding as specified in the enhanced LSB Chapter 7 for `Eh_Frame_Hdr`.

Table 4.5: Common Information Entry (CIE)

Field	Length (byte)	Description
Length	4	Length of the CIE (not including this 4-byte field)
CIE id	4	Value 0 for <code>.eh_frame</code> (used to distinguish CIEs and FDEs when scanning the section)
Version	1	Value One (1)
CIE ugmentation String	string	Null-terminated string with legal values being "" or 'z' optionally followed by single occurrences of 'P', 'L', or 'R' in any order. The presence of character(s) in the string dictates the content of field 8, the ugmentation Section. Each character has one or two associated operands in the S (see table 4.6 for which ones). Operand order depends on position in the string ('z' must be first).
Code align Factor	uleb128	To be multiplied with the " advance Location" instructions in the Call Frame Instructions
Data align Factor	sleb128	To be multiplied with all offsets in the Call Frame Instructions
Ret address Reg	1/uleb128	"virtual" register representation of the return address. In Dwarf V2, this is a byte, otherwise it is uleb128. It is a byte in gcc 3.3.x
Optional CIE ugmentation Section	varying	Present if ugmentation String in ugmentation Section field 4 is not 0. See table 4.6 for the content.
Optional Call Frame Instructions	varying	

Table 4.6: CIE Augmentation Section Content

Char	Operands	Length (byte)	Description
z	size	uleb128	Length of the remainder of the Augmentation Section
P	personality_enc	1	Encoding specifier - preferred value is a pc-relative, signed 4-byte
	personality routine	(encoded)	Encoded pointer to personality routine (actually to the PLT entry for the personality routine)
R	code_enc	1	Non-default encoding for the code-pointers (FDE members initial_location and address_range and the operand for DW_CF_set_loc) - preferred value is pc-relative, signed 4-byte
L	lsda_enc	1	FDE augmentation bodies may contain LSD pointers. If so they are encoded as specified here - preferred value is pc-relative, signed 4-byte possibly indirect thru a GOT entry

Table 4.7: Frame Descriptor Entry (FDE)

Field	Length (byte)	Description
Length	4	Length of the FDE (not including this 4-byte field)
CIE pointer	4	Distance from this field to the nearest preceding CIE (the value is subtracted from the current address). This value can never be zero and thus can be used to distinguish CIE's and FDE's when scanning the .eh_frame section
Initial Location	var	Reference to the function code corresponding to this FDE. If 'R' is missing from the CIE augmentation String, the field is an 8-byte absolute pointer. Otherwise, the corresponding EH_PE encoding in the CIE augmentation Section is used to interpret the reference
Address Range	var	Size of the function code corresponding to this FDE. If 'R' is missing from the CIE augmentation String, the field is an 8-byte unsigned number. Otherwise, the size is determined by the corresponding EH_PE encoding in the CIE augmentation Section (the value is always absolute)
Optional FDE augmentation Section	var	Present if CIE augmentation String is non-empty. See table 4.8 for the content.
Optional Call Frame Instructions	var	

Table 4.8: FDE Augmentation Section Content

Char	Operands	Length (byte)	Description
z	length	uleb128	Length of the remainder of the augmentation Section
L	LSD	var	LSD pointer, encoded in the format specified by the corresponding operand in the CIE's augmentation body. (only present if length > 0).

The existence and size of the optional call frame instruction area must be computed based on the overall size and the offset reached while scanning the preceding fields of the CIE or FDE.

The overall size of a `.eh_frame` section is given in the ELF section header. The only way to determine the number of entries is to scan the section until the end, counting entries as they are encountered.

4.3 Symbol Table

The discussion of "Function Addresses" in Section 5.2 defines some special values for symbol table fields.

The STT_GNU_IFUNC² symbol type is optional. It is the same as STT_FUNC except that it always points to a function or piece of executable code which takes no arguments and returns a function pointer. If an STT_GNU_IFUNC symbol is referred to by a relocation, then evaluation of that relocation is delayed until load-time. The value used in the relocation is the function pointer returned by an invocation of the STT_GNU_IFUNC symbol.

The purpose of the STT_GNU_IFUNC symbol type is to allow the run-time to select between multiple versions of the implementation of a specific function. The selection made in general will take the currently available hardware into account and select the most appropriate version.

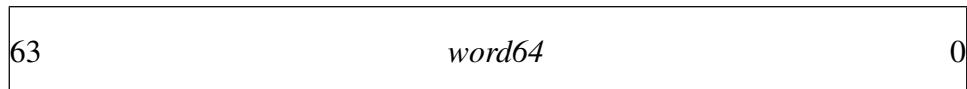
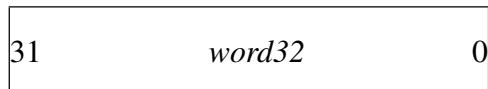
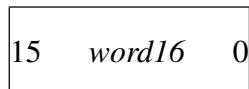
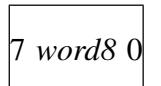
²It is specified in **Linux Extensions to g BI** at <https://github.com/hjl-tools/linux-abi>

4.4 Relocation

4.4.1 Relocation Types

Figure 4.4.1 shows the allowed relocatable fields.

Figure 4.1: Relocatable Fields



word8

This specifies a 8-bit field occupying 1 byte.

word16

This specifies a 16-bit field occupying 2 bytes with arbitrary byte alignment. These values use the same byte order as other word values in the MD64 architecture.

word32

This specifies a 32-bit field occupying 4 bytes with arbitrary byte alignment. These values use the same byte order as other word values in the MD64 architecture.

word64

This specifies a 64-bit field occupying 8 bytes with arbitrary byte alignment. These values use the same byte order as other word values in the MD64 architecture.

wordclass

This specifies *word64* for LP64 and specifies *word32* for ILP32.

The following notations are used for specifying relocations in table 4.9:

Represents the addend used to compute the value of the relocatable field.

B Represents the base address at which a shared object has been loaded into memory during execution. Generally, a shared object is built with a 0 base virtual address, but the execution address will be different.

G Represents the offset into the global offset table at which the relocation entry's symbol will reside during execution.

GOT Represents the address of the global offset table.

L Represents the place (section offset or address) of the Procedure Linkage Table entry for a symbol.

P Represents the place (section offset or address) of the storage unit being relocated (computed using `r_offset`).

S Represents the value of the symbol whose index resides in the relocation entry.

Z Represents the size of the symbol whose index resides in the relocation entry.

The MD64 LP64 BI architecture uses only `Elf64_Rela` relocation entries with explicit addends. The `r_addend` member serves as the relocation addend.

The MD64 ILP32 BI architecture uses only `Elf32_Rela` relocation entries in relocatable files. Executable files or shared objects may use either `Elf32_Rela` or `Elf32_Rel` relocation entries.

Table 4.9: Relocation Types

Name	Value	Field	Calculation
R_X86_64_NONE	0	none	none
R_X86_64_64	1	word64	S +
R_X86_64_PC32	2	word32	S + - P
R_X86_64_GOT32	3	word32	G +
R_X86_64_PLT32	4	word32	L + - P
R_X86_64_COPY	5	none	none
R_X86_64_GLOB_D_T	6	wordclass	S
R_X86_64_JUMP_SLOT	7	wordclass	S
R_X86_64_RELATIVE	8	wordclass	B +
R_X86_64_GOTPCREL	9	word32	G + GOT + - P
R_X86_64_32	10	word32	S +
R_X86_64_32S	11	word32	S +
R_X86_64_16	12	word16	S +
R_X86_64_PC16	13	word16	S + - P
R_X86_64_8	14	word8	S +
R_X86_64_PC8	15	word8	S + - P
R_X86_64_DTPMOD64	16	word64	
R_X86_64_DTPOFF64	17	word64	
R_X86_64_TPOFF64	18	word64	
R_X86_64_TLSGD	19	word32	
R_X86_64_TLSLD	20	word32	
R_X86_64_DTPOFF32	21	word32	
R_X86_64_GOTTPOFF	22	word32	
R_X86_64_TPOFF32	23	word32	
R_X86_64_PC64	24	word64	S + - P
R_X86_64_GOTOFF64	25	word64	S + - GOT
R_X86_64_GOTPC32	26	word32	GOT + - P
R_X86_64_SIZE32	32	word32	Z +
R_X86_64_SIZE64	33	word64	Z +
R_X86_64_GOTPC32_TLSDESC	34	word32	
R_X86_64_TLSDESC_C_LL	35	none	
R_X86_64_TLSDESC	36	word64	
R_X86_64_IRELATIVE	37	wordclass	indirect (B +)
R_X86_64_RELATIVE64	38	word64	B +
Deprecated	39		
Deprecated	40		
R_X86_64_GOTPCRELX	41	word32	G + GOT + - P
R_X86_64_RDX_GOTPCRELX	42	word32	G + GOT + - P

† This relocation is used only for LP64.

†† This relocation only appears in ILP32 executable files or shared objects.

The special semantics for most of these relocation types are identical to those used for

the Intel386 BI.³⁴

The R_X86_64_GOTPCREL relocation has different semantics from the R_X86_64_GOT32 or equivalent i386 R_386_GOTPC relocation. In particular, because the MD64 architecture has an addressing mode relative to the instruction pointer, it is possible to load an address from the GOT using a single instruction. The calculation done by the R_X86_64_GOTPCREL relocation gives the difference between the location in the GOT where the symbol's address is given and the location where the relocation is applied.

For the occurrence of name@GOTPCREL in the following assembler instructions:

```
call      *name@GOTPCREL(%rip)
jmp      *name@GOTPCREL(%rip)
mov      name@GOTPCREL(%rip), %reg
test     %reg, name@GOTPCREL(%rip)
binop    name@GOTPCREL(%rip), %reg
```

where binop is one of adc, add, and, cmp, or, sbb, sub, xor instructions, the R_X86_64_GOTPCRElx relocation, or the R_X86_64_Rex_GOTPCRElx relocation if the REX prefix is present, should be generated, instead of the R_X86_64_GOTPCREL relocation. See also section B.2.

The R_X86_64_32 and R_X86_64_32S relocations truncate the computed value to 32-bits. The linker must verify that the generated value for the R_X86_64_32 (R_X86_64_32S) relocation zero-extends (sign-extends) to the original 64-bit value.

program or object file using R_X86_64_8, R_X86_64_16, R_X86_64_PC16 or R_X86_64_PC8 relocations is not conformant to this BI, these relocations are only added for documentation purposes. The R_X86_64_16, and R_X86_64_8 relocations truncate the computed value to 16-bits resp. 8-bits.

The relocations R_X86_64_DTPMOD64, R_X86_64_DTPOFF64, R_X86_64_TPOFF64, R_X86_64_TLSGD, R_X86_64_TLSLD, R_X86_64_DTPOFF32, R_X86_64_GOTTPOFF and R_X86_64_TPOFF32 are listed for completeness. They are part of the Thread-Local Storage BI extensions and are documented in the document called “ELF Handling for Thread-Local Storage”⁵.

³Even though the MD64 architecture supports IP-relative addressing modes, a GOT is still required since the offset from a particular instruction to a particular data item cannot be known by the static linker.

⁴Note that the MD64 architecture assumes that offsets into GOT are 32-bit values, not 64-bit values. This choice means that a maximum of $2^{32} \cdot 8 = 2^{29}$ entries can be placed in the GOT. However, that should be more than enough for most programs. In the event that it is not enough, the linker could create multiple GOTs. Because 32-bit offsets are used, loads of global data do not require loading the offset into a displacement register; the base plus immediate displacement addressing form can be used.

⁵This document is currently available via <http://www.akkadia.org/drepper/tls.pdf>

The relocations R_X86_64_GOTPC32_TLSDESC, R_X86_64_TLSDESC_C_LL and R_X86_64_TLSDESC are also used for Thread-Local Storage, but are not documented there as of this writing. A description can be found in the document “Thread-Local Storage Descriptors for I 32 and MD64/EM64T”⁶.

In order to make this document self-contained, a description of the TLS relocations follows.

The %fs segment register is used to implement the thread pointer. The linear address of the thread pointer is stored at offset 0 relative to the %fs segment register. The following code loads the thread pointer in the %rax register:

```
movq    %fs:0, %rax
```

R_X86_64_DTPMOD64 resolves to the index of the dynamic thread vector entry that points to the base address of the TLS block corresponding to the module that defines the referenced symbol. R_X86_64_DTPOFF64 and R_X86_64_DTPOFF32 compute the offset from the pointer in that entry to the referenced symbol. The linker generates such relocations in adjacent entries in the GOT, in response to R_X86_64_TLSD and R_X86_64_TLSLD relocations. If the linker can compute the offset itself, because the referenced symbol binds locally, the relocations R_X86_64_64 and R_X86_64_32 may be used instead. Otherwise, such relocations are always in pairs, such that the R_X86_64_DTPOFF64 relocation applies to the word64 right past the corresponding R_X86_64_DTPMOD64 relocation.

R_X86_64_TPOFF64 and R_X86_64_TPOFF32 resolve to the offset from the thread pointer to a thread-local variable. The former is generated in response to R_X86_64_GOTTPOFF, that resolves to a PC-relative address of a GOT entry containing such a 64-bit offset.

R_X86_64_TLSD and R_X86_64_TLSLD both resolve to PC-relative offsets to a DTPMOD GOT entry. The difference between them is that, for R_X86_64_TLSD, the following GOT entry will contain the offset of the referenced symbol into its TLS block, whereas, for R_X86_64_TLSLD, the following GOT entry will contain the offset for the base address of the TLS block. The idea is that adding this offset to the result of R_X86_64_DTPMOD32 for a symbol ought to yield the same as the result of R_X86_64_DTPMOD64 for the same symbol.

R_X86_64_TLSDESC resolves to a pair of word64s, called TLS Descriptor, the first of which is a pointer to a function, followed by an argument. The function is passed a pointer to the this pair of entries in %rax and, using the argument in the second entry, it must compute and return in %rax the offset from the thread pointer to the symbol

⁶This document is currently available via
<http://www.fsfila.org/~lxoliva/writeups/TLS/RFC-TLSDESC-x86.txt>

referenced in the relocation, without modifying any registers other than processor flags. `R_X86_64_GOTPC32_TLSDESC` resolves to the PC-relative address of a TLS descriptor corresponding to the named symbol. `R_X86_64_TLSDESC_C_LL` must annotate the instruction used to call the TLS Descriptor resolver function, so as to enable relaxation of that instruction.

`R_X86_64_IREL_TIVE` is similar to `R_X86_64_REL_TIVE` except that the value used in this relocation is the program address returned by the function, which takes no arguments, at the address of the result of the corresponding `R_X86_64_REL_TIVE` relocation.

One use of the `R_X86_64_IREL_TIVE` relocation is to avoid name lookup for the locally defined `STT_GNU_IFUNC` symbols at load-time. Support for this relocation is optional, but is required for the `STT_GNU_IFUNC` symbols.

4.4.2 Large Models

In order to extend both the PLT and the GOT beyond 2GB, it is necessary to add appropriate relocation types to handle full 64-bit addressing. See figure 4.10.

Table 4.10: Large Model Relocation Types

Name	Value	Field	Calculation
<code>R_X86_64_GOT64</code>	27	<code>word64</code>	<code>G +</code>
<code>R_X86_64_GOTPCREL64</code>	28	<code>word64</code>	<code>G + GOT - P +</code>
<code>R_X86_64_GOTPC64</code>	29	<code>word64</code>	<code>GOT - P +</code>
Deprecated	30		
<code>R_X86_64_PLTOFF64</code>	31	<code>word64</code>	<code>L - GOT +</code>

Chapter 5

Program Loading and Dynamic Linking

5.1 Program Loading

Program loading is a process of mapping file segments to virtual memory segments. For efficient mapping executable and shared object files must have segments whose file offsets and virtual addresses are congruent modulo the page size.

To save space the file page holding the last page of the text segment may also contain the first page of the data segment. The last data page may contain file information not relevant to the running process. Logically, the system enforces the memory permissions as if each segment were complete and separate; segments' addresses are adjusted to ensure each logical page in the address space has a single set of permissions. In the example above, the region of the file holding the end of text and the beginning of data will be mapped twice: at one virtual address for text and at a different virtual address for data.

The end of the data segment requires special handling for uninitialized data, which the system defines to begin with zero values. Thus if a file's last data page includes information not in the logical memory page, the extraneous data must be set to zero, not the unknown contents of the executable file. "Impurities" in the other three pages are not logically part of the process image; whether the system expunges them is unspecified.

One aspect of segment loading differs between executable files and shared objects. Executable file segments typically contain absolute code (see section 3.5 "Coding Examples"). For the process to execute correctly, the segments must reside at the virtual addresses used to build the executable file. Thus the system uses the `p_vaddr` values unchanged as virtual addresses.

On the other hand, shared object segments typically contain position-independent code. This lets a segments virtual address change from one process to another, without invalidating execution behavior. Though the system chooses virtual addresses for individ-

ual processes, it maintains the segments' relative positions. Because position-independent code uses relative addressing between segments, the difference between virtual addresses in memory must match the difference between virtual addresses in the file.

5.1.1 Program header

The following MD64 program header types are defined:

Table 5.1: Program Header Types

Name	Value
PT_GNU_EH_FR ME	0x6474e550
PT_SUNW_EH_FR ME	0x6474e550
PT_SUNW_UNWIND	0x6464e550

PT_GNU_EH_FR ME, PT_SUNW_EH_FR ME and PT_SUNW_UNWIND The segment contains the stack unwind tables. See Section 4.2.4 of this document.¹

5.2 Dynamic Linking

Dynamic Section

Dynamic section entries give information to the dynamic linker. Some of this information is processor-specific, including the interpretation of some entries in the dynamic structure.

Global Offset Table (GOT)

Position-independent code cannot, in general, contain absolute virtual addresses. Global offset tables hold absolute addresses in private data, thus making the addresses available without compromising the position-independence and shareability of a program's text.

program references its global offset table using position-independent addressing and extracts absolute values, thus redirecting position-independent references to absolute locations.

¹The value for these program headers have been placed in the PT_LOOS and PT_HIOS (os specific range) in order to adapt to the existing GNU implementation. New OS's wanting to agree on these program header should also add it into their OS specific range.

If a program requires direct access to the absolute address of a symbol, that symbol will have a global offset table entry. Because the executable file and shared objects have separate global offset tables, a symbol's address may appear in several tables. The dynamic linker processes all the global offset table relocations before giving control to any code in the process image, thus ensuring the absolute addresses are available during execution.

The tables first entry (number zero) is reserved to hold the address of the dynamic structure, referenced with the symbol `_DYN_MIC`. This allows a program, such as the dynamic linker, to find its own dynamic structure without having yet processed its relocation entries. This is especially important for the dynamic linker, because it must initialize itself without relying on other programs to relocate its memory image. On the MD64 architecture, entries one and two in the global offset table also are reserved.

The global offset table contains 64-bit addresses.

For the large models the GOT is allowed to be up to 16EB in size.

Figure 5.1: Global Offset Table

```
extern Elf64_ ddr _GLOB_L_OFFSET_T BLE_ [];
```

The symbol `_GLOB_L_OFFSET_T BLE_` may reside in the middle of the `.got` section, allowing both negative and non-negative offsets into the array of addresses.

Function ddresses

References to the address of a function from an executable file and the shared objects associated with it might not resolve to the same value. References from within shared objects will normally be resolved by the dynamic linker to the virtual address of the function itself. References from within the executable file to a function defined in a shared object will normally be resolved by the link editor to the address of the procedure linkage table entry for that function within the executable file.

To allow comparisons of function addresses to work as expected, if an executable file references a function defined in a shared object, the link editor will place the address of the procedure linkage table entry for that function in its associated symbol table entry. This will result in symbol table entries with section index of `SHN_UNDEF` but a type of `STT_FUNC` and a non-zero `st_value`. A reference to the address of a function from within a shared library will be satisfied by such a definition in the executable.

Some relocations are associated with procedure linkage table entries. These entries are used for direct function calls rather than for references to function addresses. These relocations do not use the special symbol value described above. Otherwise a very tight endless loop would be created.

Procedure Linkage Table

Much as the global offset table redirects position-independent address calculations to absolute locations, the procedure linkage table redirects position-independent function calls to absolute locations. The link editor cannot resolve execution transfers (such as function calls) from one executable or shared object to another. Consequently, the link editor arranges to have the program transfer control to entries in the procedure linkage table. On the

MD64 architecture, procedure linkage tables reside in shared text, but they use addresses in the private global offset table. The dynamic linker determines the destinations' absolute addresses and modifies the global offset table's memory image accordingly. The dynamic linker thus can redirect the entries without compromising the position-independence and shareability of the program's text. Executable files and shared object files have separate procedure linkage tables. Unlike Intel386 BI, this BI uses the same procedure linkage table for both programs and shared objects (see figure 5.2).

Figure 5.2: Procedure Linkage Table (small and medium models)

```
.PLT0: pushq    GOT+8(%rip)          # GOT[1]
        jmp     *GOT+16(%rip)         # GOT[2]
        nopl    0x0(%rax)
.PLT1: jmp     *name1@GOTPCREL(%rip)   # 16 bytes from .PLT0
        pushq    $index1
        jmp     .PLT0
.PLT2: jmp     *name2@GOTPCREL(%rip)   # 16 bytes from .PLT1
        pushq    $index2
        jmp     .PLT0
.PLT3: ...
```

Following the steps below, the dynamic linker and the program “cooperate” to resolve symbolic references through the procedure linkage table and the global offset table.

1. When first creating the memory image of the program, the dynamic linker sets the second and the third entries in the global offset table to special values. Steps below explain more about these values.

2. Each shared object file in the process image has its own procedure linkage table, and control transfers to a procedure linkage table entry only from within the same object file.
3. For illustration, assume the program calls `name1`, which transfers control to the label `.PLT1`.
4. The first instruction jumps to the address in the global offset table entry for `name1`. Initially the global offset table holds the address of the following `pushq` instruction, not the real address of `name1`.
5. Now the program pushes a relocation index (*index*) on the stack. The relocation index is a 32-bit, non-negative index into the relocation table addressed by the `DT_JMPREL` dynamic section entry. The designated relocation entry will have type `R_X86_64_JUMP_SLOT`, and its offset will specify the global offset table entry used in the previous `jmp` instruction. The relocation entry contains a symbol table index that will reference the appropriate symbol, `name1` in the example.
6. After pushing the relocation index, the program then jumps to `.PLT0`, the first entry in the procedure linkage table. The `pushq` instruction places the value of the second global offset table entry (GOT+8) on the stack, thus giving the dynamic linker one word of identifying information. The program then jumps to the address in the third global offset table entry (GOT+16), which transfers control to the dynamic linker.
7. When the dynamic linker receives control, it unwinds the stack, looks at the designated relocation entry, finds the symbol's value, stores the “real” address for `name1` in its global offset table entry, and transfers control to the desired destination.
8. Subsequent executions of the procedure linkage table entry will transfer directly to `name1`, without calling the dynamic linker a second time. That is, the `jmp` instruction at `.PLT1` will transfer to `name1`, instead of “falling through” to the `pushq` instruction.

The `LD_BIND_NOW` environment variable can change the dynamic linking behavior. If its value is non-null, the dynamic linker evaluates procedure linkage table entries before transferring control to the program. That is, the dynamic linker processes relocation entries of type `R_X86_64_JUMP_SLOT` during process initialization. Otherwise, the dynamic linker evaluates procedure linkage table entries lazily, delaying symbol resolution and relocation until the first execution of a table entry.

Relocation entries of type `R_X86_64_TLSDESC` may also be subject to lazy relocation, using a single entry in the procedure linkage table and in the global offset table, at locations given by `DT_TLSDESC_PLT` and `DT_TLSDESC_GOT`, respectively, as described in “Thread-Local Storage Descriptors for I 32 and MD64/EM64T”².

For self-containment, `DT_TLSDESC_GOT` specifies a GOT entry in which the dynamic loader should store the address of its internal TLS Descriptor resolver function, whereas `DT_TLSDESC_PLT` specifies the address of a PLT entry to be used as the TLS descriptor resolver function for lazy resolution from within this module. The PLT entry must push the linkmap of the module onto the stack and tail-call the internal TLS Descriptor resolver function.

Large Models

In the small and medium code models the size of both the PLT and the GOT is limited by the maximum 32-bit displacement size. Consequently, the base of the PLT and the top of the GOT can be at most 2GB apart.

Therefore, in order to support the available addressing space of 16EB, it is necessary to extend both the PLT and the GOT. Moreover, the PLT needs to support the GOT being over 2GB away and the GOT can be over 2GB in size.³

The PLT is extended as shown in figure 5.3 with the assumption that the GOT address is in `%r15`⁴.

²This document is currently available via
<http://www.fsfla.org/~lxoliva/writeups/TLS/RFC-TLSDESC-x86.txt>

³If it is determined that the base of the PLT is within 2GB of the top of the GOT, it is also allowed to use the same PLT layout for a large code model object as that of the small and medium code models.

⁴See Function Prologue.

Figure 5.3: Final Large Code Model PLT

```
.PLT0: pushq  8(%r15)          # GOT[1]
       jmpq   *16(%r15)         # GOT[2]
       nopl   0x0(%rax,%rax,1)
.PLT1: movabs $name1@GOT,%r11 # 16 bytes from .PLT0
       jmp    *(%r11,%r15)
.PLT1a: pushq  $index1        # "call" dynamic linker
       jmp    .PLT0
.PLT2: ...
.PLTx: movabs $namex@GOT,%r11 # 102261125th entry
       jmp    *(%r11,%r15)
.PLTx: pushq  $indexx
       pushq  8(%r15)          # repeat .PLT0 code
       jmpq   *16(%r15)
.PLTy: ...                   # 27 bytes from .PLTx
```

This way, for the first 102261125 entries, each PLT entry besides .PLT0 uses only 21 bytes. Afterwards, the PLT entry code changes by repeating that of .PLT0, when each PLT entry is 27 bytes long. Notice that any alignment consideration is dropped in order to keep the PLT size down.

Each extended PLT entry is thus 5 to 11 bytes larger than the small and medium code model PLT entries.

The functionality of entry .PLT0 remains unchanged from the small and medium code models.

Note that the symbol index is still limited to 32 bits, which would allow for up to 4G global and external functions.

Typically, UNIX compilers support two types of PLT, generally through the options `-fPIC` and `-fPIC`. When building position-independent objects using the large code model, only `-fPIC` is allowed. Using the option `-fPIC` with the large code model remains reserved for future use.

5.2.1 Program Interpreter

The valid program interpreter for programs conforming to the MD64 BI is listed in Table 5.4, which also contains the program interpreter used by Linux.

Figure 5.4: MD64 Program Interpreter

Data Model	Path	Linux Path
LP64	/lib/ld64.so.1	/lib64/ld-linux-x86-64.so.2
ILP32	/lib/ldx32.so.1	/libx32/ld-linux-x32.so.2

5.2.2 Initialization and Termination Functions

The implementation is responsible for executing the initialization functions specified by `DT_INIT`, `DT_INIT_ RR Y`, and `DT_PREINIT_ RR Y` entries in the executable file and shared object files for a process, and the termination (or finalization) functions specified by `DT_FINI` and `DT_FINI_ RR Y`, as specified by the *System V BI*. The user program plays no further part in executing the initialization and termination functions specified by these dynamic tags.

Chapter 6

Libraries

further review of the Intel386 BI is needed.

6.1 C Library

6.1.1 Global Data Symbols

The symbols `_fp_hw`, `__flt_rounds` and `__huge_val` are not provided by the MD64 BI.

6.1.2 Floating Point Environment Functions

ISO C 99 defines the floating point environment functions from `<fenv.h>`. Since MD64 has two floating point units with separate control words, the programming environment has to keep the control values in sync. On the other hand this means that routines accessing the control words only need to access one unit, and the SSE unit is the unit that should be accessed in these cases. The function `fegetround` therefore only needs to report the rounding value of the SSE unit and can ignore the x87 unit.

6.2 Unwind Library Interface

This section defines the Unwind Library interface¹, expected to be provided by any MD64 ps BI-compliant system. This is the interface on which the C++ BI exception-handling facilities are built. We assume as a basis the Call Frame Information tables described in the DW RF Debugging Information Format document.

¹The overall structure and the external interface is derived from the I -64 UNIX System V BI

This section is meant to specify a language-independent interface that can be used to provide higher level exception-handling facilities such as those defined by C++.

The unwind library interface consists of at least the following routines:

```
_Unwind_RaiseException ,  
_Unwind_Resume ,  
_Unwind_DeleteException ,  
_Unwind_GetGR ,  
_Unwind_SetGR ,  
_Unwind_GetIP ,  
_Unwind_GetIPInfo ,  
_Unwind_SetIP ,  
_Unwind_GetRegionStart ,  
_Unwind_GetLanguageSpecificData ,  
_Unwind_ForcedUnwind ,  
_Unwind_GetCF
```

In addition, two data types are defined (`_Unwind_Context` and `_Unwind_Exception`) to interface a calling runtime (such as the C++ runtime) and the above routine. All routines and interfaces behave as if defined `extern "C"`. In particular, the names are not mangled. All names defined as part of this interface have a "`_Unwind_`" prefix.

Lastly, a language and vendor specific personality routine will be stored by the compiler in the unwind descriptor for the stack frames requiring exception processing. The personality routine is called by the unwinder to handle language-specific tasks such as identifying the frame handling a particular exception.

6.2.1 Exception Handler Framework

Reasons for Unwinding

There are two major reasons for unwinding the stack:

- exceptions, as defined by languages that support them (such as C++)
- “forced” unwinding (such as caused by `longjmp` or thread termination)

The interface described here tries to keep both similar. There is a major difference, however.

- In the case where an exception is thrown, the stack is unwound while the exception propagates, but it is expected that the personality routine for each stack frame knows

whether it wants to catch the exception or pass it through. This choice is thus delegated to the personality routine, which is expected to act properly for any type of exception, whether “native” or “foreign”. Some guidelines for “acting properly” are given below.

- During “forced unwinding”, on the other hand, an external agent is driving the unwinding. For instance, this can be the `longjmp` routine. This external agent, not each personality routine, knows when to stop unwinding. The fact that a personality routine is not given a choice about whether unwinding will proceed is indicated by the `_U _FORCE_UNWIND` flag.

To accommodate these differences, two different routines are proposed. `_Unwind_RaiseException` performs exception-style unwinding, under control of the personality routines. `_Unwind_ForcedUnwind`, on the other hand, performs unwinding, but gives an external agent the opportunity to intercept calls to the personality routine. This is done using a proxy personality routine, that intercepts calls to the personality routine, letting the external agent override the defaults of the stack frame’s personality routine.

As a consequence, it is not necessary for each personality routine to know about any of the possible external agents that may cause an unwind. For instance, the C++ personality routine need deal only with C++ exceptions (and possibly disguising foreign exceptions), but it does not need to know anything specific about unwinding done on behalf of `longjmp` or `pthreads cancellation`.

The Unwind Process

The standard BI exception handling/unwind process begins with the raising of an exception, in one of the forms mentioned above. This call specifies an exception object and an exception class.

The runtime framework then starts a two-phase process:

- In the *search* phase, the framework repeatedly calls the personality routine, with the `_U _SE RCH_PH SE` flag as described below, first for the current `%rip` and register state, and then unwinding a frame to a new `%rip` at each step, until the personality routine reports either success (a handler found in the queried frame) or failure (no handler) in all frames. It does not actually restore the unwound state, and the personality routine must access the state through the PI.
- If the search phase reports a failure, e.g. because no handler was found, it will call `terminate()` rather than commence phase 2.

If the search phase reports success, the framework restarts in the *cleanup* phase. Again, it repeatedly calls the personality routine, with the _U _CLE NUP_PH SE flag as described below, first for the current %rip and register state, and then unwinding a frame to a new %rip at each step, until it gets to the frame with an identified handler. At that point, it restores the register state, and control is transferred to the user landing pad code.

Each of these two phases uses both the unwind library and the personality routines, since the validity of a given handler and the mechanism for transferring control to it are language-dependent, but the method of locating and restoring previous stack frames is language-independent.

The two-phase exception-handling model is not strictly necessary to implement C++ language semantics, but it does provide some benefits. For example, the first phase allows an exception-handling mechanism to *dismiss* an exception before stack unwinding begins, which allows *presumptive* exception handling (correcting the exceptional condition and resuming execution at the point where it was raised). While C++ does not support presumptive exception handling, other languages do, and the two-phase model allows C++ to coexist with those languages on the stack.

Note that even with a two-phase model, we may execute each of the two phases more than once for a single exception, as if the exception was being thrown more than once. For instance, since it is not possible to determine if a given catch clause will re-throw or not without executing it, the exception propagation effectively stops at each catch clause, and if it needs to restart, restarts at phase 1. This process is not needed for destructors (cleanup code), so the phase 1 can safely process all destructor-only frames at once and stop at the next enclosing catch clause.

For example, if the first two frames unwound contain only cleanup code, and the third frame contains a C++ catch clause, the personality routine in phase 1, does not indicate that it found a handler for the first two frames. It must do so for the third frame, because it is unknown how the exception will propagate out of this third frame, e.g. by re-throwing the exception or throwing a new one in C++.

The PI specified by the MD64 ps BI for implementing this framework is described in the following sections.

6.2.2 Data Structures

Reason Codes

The unwind interface uses reason codes in several contexts to identify the reasons for failures or other actions, defined as follows:

```

typedef enum {
    _URC_NO_RE SON = 0,
    _URC_FOREIGN_EXCEPTION_C UGHT = 1,
    _URC_F T L_PH SE2_ERROR = 2,
    _URC_F T L_PH SE1_ERROR = 3,
    _URC_NORM L_STOP = 4,
    _URC_END_OF_ST CK = 5,
    _URC_H NDLER_FOUND = 6,
    _URC_INST LL_CONTEXT = 7,
    _URC_CONTINUE_UNWIND = 8
} _Unwind_Reason_Code;

```

The interpretations of these codes are described below.

Exception Header

The unwind interface uses a pointer to an exception header object as its representation of an exception being thrown. In general, the full representation of an exception object is language- and implementation-specific, but is prefixed by a header understood by the unwind interface, defined as follows:

```

typedef void (*_Unwind_Exception_Cleanup_Fn)
    (_Unwind_Reason_Code reason,
     struct _Unwind_Exception *exc);
struct _Unwind_Exception {
    uint64 exception_class;           exception_class;
    _Unwind_Exception_Cleanup_Fn exception_cleanup;   exception_cleanup;
    uint64 private_1;                 private_1;
    uint64 private_2;                 private_2;
};

```

A `_Unwind_Exception` object must be eightbyte aligned. The first two fields are set by user code prior to raising the exception, and the latter two should never be touched except by the runtime.

The `exception_class` field is a language- and implementation-specific identifier of the kind of exception. It allows a personality routine to distinguish between native and foreign exceptions, for example. By convention, the high 4 bytes indicate the vendor (for instance MD\0), and the low 4 bytes indicate the language. For the C++ BI described in this document, the low four bytes are C++\0.

The `exception_cleanup` routine is called whenever an exception object needs to be destroyed by a different runtime than the runtime which created the exception object, for instance if a Java exception is caught by a C++ catch handler. In such a case, a reason code (see above) indicates why the exception object needs to be deleted:

`_URC_FOREIGN_EXCEPTION_C UGHT = 1` This indicates that a different runtime caught this exception. Nested foreign exceptions, or re-throwing a foreign exception, result in undefined behavior.

_URC_F T L_PH SE1_ERROR = 3 The personality routine encountered an error during phase 1, other than the specific error codes defined.

_URC_F T L_PH SE2_ERROR = 2 The personality routine encountered an error during phase 2, for instance a stack corruption.

Normally, all errors should be reported during phase 1 by returning from `_Unwind_RaiseException`. However, landing pad code could cause stack corruption between phase 1 and phase 2. For a C++ exception, the runtime should call `terminate()` in that case.

The private unwinder state (`private_1` and `private_2`) in an exception object should be neither read by nor written to by personality routines or other parts of the language-specific runtime. It is used by the specific implementation of the unwinder on the host to store internal information, for instance to remember the final handler frame between unwinding phases.

In addition to the above information, a typical runtime such as the C++ runtime will add language-specific information used to process the exception. This is expected to be a contiguous area of memory after the `_Unwind_Exception` object, but this is not required as long as the matching personality routines know how to deal with it, and the `exception_cleanup` routine de-allocates it properly.

Unwind Context

The `_Unwind_Context` type is an opaque type used to refer to a system-specific data structure used by the system unwinder. This context is created and destroyed by the system, and passed to the personality routine during unwinding.

```
struct _Unwind_Context
```

6.2.3 Throwing an Exception

`_Unwind_RaiseException`

```
_Unwind_Reason_Code _Unwind_RaiseException  
( struct _Unwind_Exception *exception_object );
```

Raise an exception, passing along the given exception object, which should have its `exception_class` and `exception_cleanup` fields set. The exception object has been allocated by the language-specific runtime, and has a language-specific format, except that it must contain an `_Unwind_Exception` struct (see Exception Header above). `_Unwind_RaiseException` does not return, unless an error condition is found (such as no

handler for the exception, bad stack format, etc.). In such a case, an `_Unwind_Reason_Code` value is returned.

Possibilities are:

`_URC_END_OF_ST CK` The unwinder encountered the end of the stack during phase 1, without finding a handler. The unwind runtime will not have modified the stack. The C++ runtime will normally call `uncaught_exception()` in this case.

`_URC_FT_LPH_SE1_ERROR` The unwinder encountered an unexpected error during phase 1, e.g. stack corruption. The unwind runtime will not have modified the stack. The C++ runtime will normally call `terminate()` in this case.

If the unwinder encounters an unexpected error during phase 2, it should return `_URC_FT_LPH_SE2_ERROR` to its caller. In C++, this will usually be `__cxa_throw`, which will call `terminate()`.

The unwind runtime will likely have modified the stack (e.g. popped frames from it) or register context, or landing pad code may have corrupted them. As a result, the caller of `_Unwind_RaiseException` can make no assumptions about the state of its stack or registers.

`_Unwind_ForcedUnwind`

```
typedef _Unwind_Reason_Code (*_Unwind_Stop_Fn)
    (int version,
     _Unwind_ction actions,
     uint64 exceptionClass,
     struct _Unwind_Exception *exceptionObject,
     struct _Unwind_Context *context,
     void *stop_parameter );
_Unwind_Reason_Code_Unwind_ForcedUnwind
    ( struct _Unwind_Exception *exception_object,
      _Unwind_Stop_Fn stop,
      void *stop_parameter );
```

Raise an exception for forced unwinding, passing along the given exception object, which should have its `exception_class` and `exception_cleanup` fields set. The exception object has been allocated by the language-specific runtime, and has a language-specific format, except that it must contain an `_Unwind_Exception` struct (see Exception Header above).

Forced unwinding is a single-phase process (phase 2 of the normal exception-handling process). The `stop` and `stop_parameter` parameters control the termination of the unwind process, instead of the usual personality routine query. The `stop` function parameter is called for each unwind frame, with the parameters described for the usual personality routine below, plus an additional `stop_parameter`.

When the `stop` function identifies the destination frame, it transfers control (according to its own, unspecified, conventions) to the user code as appropriate without returning, normally after calling `_Unwind_DeleteException`. If not, it should return an `_Unwind_Reason_Code` value as follows:

`_URC_NO_RE SON` This is not the destination frame. The unwind runtime will call the frame's personality routine with the `_U_FORCE_UNWIND` and `_U_CLE_NUP_PH_SE` flags set in actions, and then unwind to the next frame and call the stop function again.

`_URC_END_OF_ST CK` In order to allow `_Unwind_ForcedUnwind` to perform special processing when it reaches the end of the stack, the unwind runtime will call it after the last frame is rejected, with a `NULL` stack pointer in the context, and the stop function must catch this condition (i.e. by noticing the `NULL` stack pointer). It may return this reason code if it cannot handle end-of-stack.

`_URC_F_T_L_PH_SE2_ERROR` The stop function may return this code for other fatal conditions, e.g. stack corruption.

If the stop function returns any reason code other than `_URC_NO_RE SON`, the stack state is indeterminate from the point of view of the caller of `_Unwind_ForcedUnwind`. Rather than attempt to return, therefore, the unwind library should return `_URC_F_T_L_PH_SE2_ERROR` to its caller.

Example: `longjmp_unwind()`

The expected implementation of `longjmp_unwind()` is as follows. The `setjmp()` routine will have saved the state to be restored in its customary place, including the frame pointer. The `longjmp_unwind()` routine will call `_Unwind_ForcedUnwind` with a stop function that compares the frame pointer in the context record with the saved frame pointer. If equal, it will restore the `setjmp()` state as customary, and otherwise it will return `_URC_NO_RE SON` or `_URC_END_OF_ST CK`.

If a future requirement for two-phase forced unwinding were identified, an alternate routine could be defined to request it, and an actions parameter flag defined to support it.

`_Unwind_Resume`

```
void _Unwind_Resume  
    (struct _Unwind_Exception *exception_object);
```

Resume propagation of an existing exception e.g. after executing cleanup code in a partially unwound stack. A call to this routine is inserted at the end of a landing pad that

performed cleanup, but did not resume normal execution. It causes unwinding to proceed further.

`_Unwind_Resume` should not be used to implement re-throwing. To the unwinding runtime, the catch code that re-throws was a handler, and the previous unwinding session was terminated before entering it. Re-throwing is implemented by calling `_Unwind_RaiseException` again with the same exception object.

This is the only routine in the unwind library which is expected to be called directly by generated code: it will be called at the end of a landing pad in a "landing-pad" model.

6.2.4 Exception Object Management

`_Unwind_DeleteException`

```
void _Unwind_DeleteException
    (struct _Unwind_Exception *exception_object);
```

Deletes the given exception object. If a given runtime resumes normal execution after catching a foreign exception, it will not know how to delete that exception. Such an exception will be deleted by calling `_Unwind_DeleteException`. This is a convenience function that calls the function pointed to by the `exception_cleanup` field of the exception header.

6.2.5 Context Management

These functions are used for communicating information about the unwind context (i.e. the unwind descriptors and the user register state) between the unwind library and the personality routine and landing pad. They include routines to read or set the context record images of registers in the stack frame corresponding to a given unwind context, and to identify the location of the current unwind descriptors and unwind frame.

`_Unwind_GetGR`

```
uint64 _Unwind_GetGR
    (struct _Unwind_Context *context, int index);
```

This function returns the 64-bit value of the given general register. The register is identified by its index as given in figure 3.36.

During the two phases of unwinding, no registers have a guaranteed value.

`_Unwind_SetGR`

```
void _Unwind_SetGR
    (struct _Unwind_Context *context,
     int index,
     uint64 new_value);
```

This function sets the 64-bit value of the given register, identified by its index as for `_Unwind_GetGR`.

The behavior is guaranteed only if the function is called during phase 2 of unwinding, and applied to an unwind context representing a handler frame, for which the personality routine will return `_URC_INST_LL_CONTEXT`. In that case, only registers `%rdi`, `%rsi`, `%rdx`, `%rcx` should be used. These scratch registers are reserved for passing arguments between the personality routine and the landing pads.

`_Unwind_GetIP`

```
uint64 _Unwind_GetIP
    (struct _Unwind_Context *context);
```

This function returns the 64-bit value of the instruction pointer (IP).

During unwinding, the value is guaranteed to be the address of the instruction immediately following the call site in the function identified by the unwind context. This value may be outside of the procedure fragment for a function call that is known to not return (such as `_Unwind_Resume`).

Applications which unwind through asynchronous signals and other non-call locations should use `_Unwind_GetIPInfo` below, and the additional flag that function provides.

`_Unwind_GetIPInfo`

```
uint64 _Unwind_GetIPInfo
    (struct _Unwind_Context *context, int *ip_before_insn);
```

This function returns the same value as `_Unwind_GetIP`. In addition, the argument `ip_before_insn` must not be null, and `*ip_before_insn` is updated with a flag which indicates whether the returned pointer is at or after the first not yet fully executed instruction.

If `*ip_before_insn` is false, the application calling `_Unwind_GetIPInfo` should assume that the instruction pointer provided points after a call instruction which has not yet returned. In general, this means that the application should use the preceding call instruction as the instruction pointer location of the unwind context. Typically, this can be approximated by subtracting one from the returned instruction pointer.

If `*ip_before_insn` is true, then the instruction pointer does not refer to an active call site. Usually, this means that the instruction pointer refers to the point at which an asynchronous signal arrived. In this case, the application should use the instruction pointer returned from `_Unwind_GetIPInfo` as the instruction pointer location of the unwind context, without adjustment.

`_Unwind_SetIP`

```
void _Unwind_SetIP  
    (struct _Unwind_Context *context,  
     uint64 new_value);
```

This function sets the value of the instruction pointer (IP) for the routine identified by the unwind context.

The behavior is guaranteed only when this function is called for an unwind context representing a handler frame, for which the personality routine will return `_URC_INST_LL_CONTEXT`. In this case, control will be transferred to the given address, which should be the address of a landing pad.

`_Unwind_GetLanguageSpecificData`

```
uint64 _Unwind_GetLanguageSpecificData  
    (struct _Unwind_Context *context);
```

This routine returns the address of the language-specific data area for the current stack frame.

This routine is not strictly required: it could be accessed through `_Unwind_SetIP` using the documented format of the DW_RF Call Frame Information Tables, but since this work has been done for finding the personality routine in the first place, it makes sense to cache the result in the context. We could also pass it as an argument to the personality routine.

`_Unwind_GetRegionStart`

```
uint64 _Unwind_GetRegionStart  
    (struct _Unwind_Context *context);
```

This routine returns the address of the beginning of the procedure or code fragment described by the current unwind descriptor block.

This information is required to access any data stored relative to the beginning of the procedure fragment. For instance, a call site table might be stored relative to the beginning of the procedure fragment that contains the calls. During unwinding, the function returns the start of the procedure fragment containing the call site in the current stack frame.

`_Unwind_GetCF`

```
uint64 _Unwind_GetCF  
    (struct _Unwind_Context *context);
```

This function returns the 64-bit Canonical Frame address which is defined as the value of `%rsp` at the call site in the previous frame. This value is guaranteed to be correct any time the context has been passed to a personality routine or a stop function.

6.2.6 Personality Routine

```
_Unwind_Reason_Code (*__personality_routine)
(int version,
 _Unwind_ ction actions,
 uint64 exceptionClass,
 struct _Unwind_Exception *exceptionObject,
 struct _Unwind_Context *context);
```

The personality routine is the function in the C++ (or other language) runtime library which serves as an interface between the system unwind library and language-specific exception handling semantics. It is specific to the code fragment described by an unwind info block, and it is always referenced via the pointer in the unwind info block, and hence it has no ps BI-specified name.

Parameters

The personality routine parameters are as follows:

version Version number of the unwinding runtime, used to detect a mis-match between the unwinder conventions and the personality routine, or to provide backward compatibility. For the conventions described in this document, version will be 1.

actions Indicates what processing the personality routine is expected to perform, as a bit mask. The possible actions are described below.

exceptionClass An 8-byte identifier specifying the type of the thrown exception. By convention, the high 4 bytes indicate the vendor (for instance MD\0), and the low 4 bytes indicate the language. For the C++ BI described in this document, the low four bytes are C++\0. This is not a null-terminated string. Some implementations may use no null bytes.

exceptionObject The pointer to a memory location recording the necessary information for processing the exception according to the semantics of a given language (see the Exception Header section above).

context Unwinder state information for use by the personality routine. This is an opaque handle used by the personality routine in particular to access the frame's registers (see the Unwind Context section above).

return value The return value from the personality routine indicates how further unwind should happen, as well as possible error conditions. See the following section.

Personality Routine Actions

The actions argument to the personality routine is a bitwise OR of one or more of the following constants:

```
typedef int _Unwind_ction;
const _Unwind_ction _U_SE_RCH_PH_SE = 1;
const _Unwind_ction _U_CLE_NUP_PH_SE = 2;
const _Unwind_ction _U_H_NDLER_FR_ME = 4;
const _Unwind_ction _U_FORCE_UNWIND = 8;
```

_U_SE_RCH_PH_SE Indicates that the personality routine should check if the current frame contains a handler, and if so return `_URC_HANDLER_FOUND`, or otherwise return `_URC_CONTINUE_UNWIND`. `_U_SE_RCH_PH_SE` cannot be set at the same time as `_U_CLE_NUP_PH_SE`.

_U_CLE_NUP_PH_SE Indicates that the personality routine should perform cleanup for the current frame. The personality routine can perform this cleanup itself, by calling nested procedures, and return `_URC_CONTINUE_UNWIND`. Alternatively, it can setup the registers (including the IP) for transferring control to a "landing pad", and return `_URC_INST_LL_CONTEXT`.

_U_H_NDLER_FR_ME During phase 2, indicates to the personality routine that the current frame is the one which was flagged as the handler frame during phase 1. The personality routine is not allowed to change its mind between phase 1 and phase 2, i.e. it must handle the exception in this frame in phase 2.

_U_FORCE_UNWIND During phase 2, indicates that no language is allowed to "catch" the exception. This flag is set while unwinding the stack for `longjmp` or during thread cancellation. User-defined code in a catch clause may still be executed, but the catch clause must resume unwinding with a call to `_Unwind_Resume` when finished.

Transferring Control to a Landing Pad

If the personality routine determines that it should transfer control to a landing pad (in phase 2), it may set up registers (including IP) with suitable values for entering the landing pad (e.g. with landing pad parameters), by calling the context management routines above. It then returns `_URC_INST_LL_CONTEXT`.

Prior to executing code in the landing pad, the unwind library restores registers not altered by the personality routine, using the context record, to their state in that frame

before the call that threw the exception, as follows. All registers specified as callee-saved by the base BI are restored, as well as scratch registers `%rdi`, `%rsi`, `%rdx`, `%rcx` (see below). Except for those exceptions, scratch (or caller-saved) registers are not preserved, and their contents are undefined on transfer.

The landing pad can either resume normal execution (as, for instance, at the end of a C++ catch), or resume unwinding by calling `_Unwind_Resume` and passing it the `exceptionObject` argument received by the personality routine. `_Unwind_Resume` will never return.

`_Unwind_Resume` should be called if and only if the personality routine did not return `_Unwind_H_NDLER_FOUND` during phase 1. As a result, the unwinder can allocate resources (for instance memory) and keep track of them in the exception object reserved words. It should then free these resources before transferring control to the last (handler) landing pad. It does not need to free the resources before entering non-handler landing-pads, since `_Unwind_Resume` will ultimately be called.

The landing pad may receive arguments from the runtime, typically passed in registers set using `_Unwind_SetGR` by the personality routine. For a landing pad that can call to `_Unwind_Resume`, one argument must be the `exceptionObject` pointer, which must be preserved to be passed to `_Unwind_Resume`.

The landing pad may receive other arguments, for instance a switch value indicating the type of the exception. Four scratch registers are reserved for this use (`%rdi`, `%rsi`, `%rdx`, `%rcx`)².

Rules for Correct Inter-Language Operation

The following rules must be observed for correct operation between languages and/or runtimes from different vendors:

An exception which has an unknown class must not be altered by the personality routine. The semantics of foreign exception processing depend on the language of the stack frame being unwound. This covers in particular how exceptions from a foreign language are mapped to the native language in that frame.

If a runtime resumes normal execution, and the caught exception was created by another runtime, it should call `_Unwind_DeleteException`. This is true even if it understands the exception object format (such as would be the case between different C++ runtimes).

A runtime is not allowed to catch an exception if the `_U_FORCE_UNWIND` flag was passed to the personality routine.

²GCC uses `%rax` and `%rdx` instead.

Example: Foreign Exceptions in C++. In C++, foreign exceptions can be caught by a `catch(...)` statement. They can also be caught as if they were of a `_foreign_exception` class, defined in `<exception>`. The `_foreign_exception` may have subclasses, such as `_java_exception` and `_ada_exception`, if the runtime is capable of identifying some of the foreign languages.

The behavior is undefined in the following cases:

- `_foreign_exception` catch argument is accessed in any way (including taking its address).
- `_foreign_exception` is active at the same time as another exception (either there is a nested exception while catching the foreign exception, or the foreign exception was itself nested).
- `uncaught_exception()`, `set_terminate()`, `set_unexpected()`, `terminate()`, or `unexpected()` is called at a time a foreign exception exists (for example, calling `set_terminate()` during unwinding of a foreign exception).

In these cases might involve accessing C++ specific content of the thrown exception, for instance to chain active exceptions.

Otherwise, a catch block catching a foreign exception is allowed:

- to resume normal execution, thereby stopping propagation of the foreign exception and deleting it, or
- to re-throw the foreign exception. In that case, the original exception object must be unaltered by the C++ runtime.

A catch-all block may be executed during forced unwinding. For instance, a `longjmp` may execute code in a `catch(...)` during stack unwinding. However, if this happens, unwinding will proceed at the end of the catch-all block, whether or not there is an explicit re-throw.

Setting the low 4 bytes of exception class to C++\0 is reserved for use by C++ runtimes compatible with the common C++ BI.

6.3 Unwinding Through ssembler Code

For successful unwinding on MD64 every function must provide a valid debug information in the DW RF Debugging Information Format. In high level languages (e.g. C/C++),

Fortran, da, ...) this information is generated by the compiler itself. However for hand-written assembly routines the debug info must be provided by the author of the code. To ease this task some new assembler directives are added:

.cfi_startproc is used at the beginning of each function that should have an entry in **.eh_frame**. It initializes some internal data structures and emits architecture dependent initial CFI instructions. Each **.cfi_startproc** directive has to be closed by **.cfi_endproc**.

.cfi_endproc is used at the end of a function where it closes its unwind entry previously opened by **.cfi_startproc** and emits it to **.eh_frame**.

.cfi_def_cfa REGISTER, OFFSET defines a rule for computing CF as: take address from REGISTER and add OFFSET to it.

.cfi_def_cfa_register REGISTER modifies a rule for computing CF . From now on REGISTER will be used instead of the old one. The offset remains the same.

.cfi_def_cfa_offset OFFSET modifies a rule for computing CF . The register remains the same, but OFFSET is new. Note that this is the absolute offset that will be added to a defined register to compute the CF address.

.cfi_adjust_cfa_offset OFFSET is similar to **.cfi_def_cfa_offset** but OFFSET is a relative value that is added or subtracted from the previous offset.

.cfi_offset REGISTER, OFFSET saves the previous value of REGISTER at offset OFFSET from CF .

.cfi_rel_offset REGISTER, OFFSET saves the previous value of REGISTER at offset OFFSET from the current CF register. This is transformed to **.cfi_offset** using the known displacement of the CF register from the CF . This is often easier to use, because the number will match the code it is annotating.

.cfi_escape EXPRESSION[, ...] allows the user to add arbitrary bytes to the unwind info. One might use this to add OS-specific CFI opcodes, or generic CFI opcodes that the assembler does not support.

Figure 6.1: Examples for Unwinding in ssembler

```
# - function with local variable allocated on the stack
.type    func_locvars,@function
func_locvars:
.cfi_startproc
# allocate space for local vars
sub      $0x1234, %rsp
.cfi_adjust_cfa_offset 0x1234
# body
...
# release space of local vars and return
add      $0x1234, %rsp
.cfi_adjust_cfa_offset -0x1234
ret
.cfi_endproc

# - function that moves frame pointer to another register
# and then allocates space for local variables
.type    func_otherreg,@function
func_otherreg:
.cfi_startproc
# save frame pointer to r12
movq    %rsp, %r12
.cfi_def_cfa_register r12
# allocate space for local vars
# (no .cfi_{def,adjust}_cfa_offset needed here,
# because CF is computed from r12!)
sub      $100,%rsp
# body
...
# restore frame pointer from r12
movq    %r12, %rsp
.cfi_def_cfa_register rsp
ret
.cfi_endproc
```

Chapter 7

Development Environment

During compilation of C or C++ code at least the symbols in table 7.1 are defined by the pre-processor ¹.

Table 7.1: Predefined Pre-Processor Symbols

<code>__amd64</code>	Defined for both LP64 and ILP32 programming models.
<code>__amd64__</code>	Defined for both LP64 and ILP32 programming models.
<code>__x86_64</code>	Defined for both LP64 and ILP32 programming models.
<code>__x86_64__</code>	Defined for both LP64 and ILP32 programming models.
<code>_LP64</code>	Defined for LP64 programming model.
<code>__LP64__</code>	Defined for LP64 programming model.
<code>_ILP32</code>	Defined for ILP32 programming model.
<code>__ILP32__</code>	Defined for ILP32 programming model.

¹ `__LP64` and `__LP64__` were added to GCC 3.3 in March, 2003.

Chapter 8

Execution Environment

Not done yet.

Chapter 9

Conventions

1

¹This chapter is used to document some features special to the MD64 BI. The different sections might be moved to another place or removed completely.

9.1 C++

For the C++ BI we will use the I -64 C++ BI and instantiate it appropriately. The current draft of that BI is available at:

<http://mentorembedded.github.io/cxx-abi/>

9.2 Fortran

formal Fortran BI does not exist. Most Fortran compilers are designed for very specific high performance computing applications, so Fortran compilers use different passing conventions and memory layouts optimized for their specific purpose. For example, Fortran applications that must run on distributed memory machines need a different data representation for array descriptors (also known as dope vectors, or fat pointers) than applications running on symmetric multiprocessor shared memory machines.

normative BI for Fortran is therefore not desirable. However, for interoperability of different Fortran compilers, as well as for interoperability with other languages, this section provides some guidelines for data types representation, and argument passing. The guidelines in this section are derived from the GNU Fortran 77 (G77) compiler, and are also followed by the GNU Fortran 95 (`gfortran`) compiler (restricted to Fortran 77 features). Other Fortran compilers already available for MD64 at the time of this writing may use different conventions, so compatibility is not guaranteed.

When this text uses the term *Fortran procedure*, the text applies to both Fortran `FUNCTION` and `SUBROUTINE` subprograms as well as for alternate `ENTRY` points, unless specifically stated otherwise.

Everything not explicitly defined in this BI is left to the implementation.

9.2.1 Names

External names in Fortran are names of entities visible to all subprograms at link time. This includes names of `COMMON` blocks and Fortran procedures. To avoid name space conflicts with linked-in libraries, all external names have to be mangled. And to avoid name space conflicts of mangled external names with local names, all local names must also be mangled. The mangling scheme is straightforward as follows:

- all names that do not have any underscores in it should have *one* underscore appended
- all external names containing one or more underscores in it (wherever) should have *two* underscores appended².
- all external names should be mapped to lower case, following the traditional UNIX model for Fortran compilers

For examples see figure 9.1:

²Historically, this is to be compatible with f2c.

Figure 9.1: Example mapping of names

Fortran external name	Linker name
FOO	foo_
foo	foo_
Foo	foo_
foo_	foo_
f_oo	f_oo_

The entry point of the main program unit is called `M_IN__`. The symbol name for the blank common block is `_BLNK__`. the external name of the unnamed `BLOCK_D_T` routine is `_BLOCK_D_T__`.

9.2.2 Representation of Fortran Types

For historical reasons, GNU Fortran 77 maps Fortran programs to the C BI, so the data representation can be explained best by providing the mapping of Fortran types to C types used by G77 on MD64³ as in figure 9.2. The “`TYPE*N`” notation specifies that variables or aggregate members of type `TYPE` shall occupy `N` bytes of storage.

Figure 9.2: Mapping of Fortran to C types

Fortran	Data kind	Equivalent C type
<code>INTEGER*4</code>	Default integer	<code>signed int</code>
<code>INTEGER*8</code>	Double precision integer	<code>signed long</code>
<code>RE_L*4</code>	Single precision FP number	<code>float</code>
<code>RE_L*8</code>	Double precision FP number	<code>double</code>
<code>COMPLEX*4</code>	Single precision complex FP number	<code>complex float</code>
<code>COMPLEX*8</code>	Double precision complex FP number	<code>complex double</code>
<code>LOGIC L</code>	Boolean logical type	<code>signed int</code>
<code>CH_R_CTER</code>	Text string	<code>char[] + length</code>

The values for type `LOGIC L` are `.TRUE.` implemented as 1 and `.F_LSE.` implemented as 0.

³G77 provides a header `g2c.h` with the equivalent C type definitions for all supported Fortran scalar types.

Data objects with a `CH R CTER` type⁴ are represented as an array of characters of the C `char` type (not guaranteed to be “\0” terminated) with a separate length counter to distinguish between `CH R CTER` data objects with a length parameter, and aggregate types of `CH R CTER` data objects, possibly also with a length parameter.

Layout of other aggregate types is implementation defined. GNU Fortran puts all arrays in contiguous memory in column-major order. GNU Fortran 95 builds an equivalent C struct for derived types without reordering the type fields. Other compilers may use other representations as needed. The representation and use of Fortran 90/95 array descriptors is implementation defined. Note that array indices start at 1 by default.

Fortran 90/95 allow different kinds of each basic type using the `kind` type parameter of a type. Kind type parameter values are implementation defined.

Layout of the commonly used Cray pointers is implementation defined.

9.2.3 Argument Passing

For each given Fortran 77 procedure, an equivalent C prototype can be derived. Once this equivalent C prototype is known, the C BI conventions should be applied to determine how arguments are passed to the Fortran procedure.

G77 passes all (user defined) formal arguments of a procedure by reference. Specifically, pointers to the location in memory of a variable, array, array element, a temporary location that holds the result of evaluating an expression or a temporary or permanent location that holds the value of a constant (xf. g77 manual) are passed as actual arguments.

Artificial compiler generated arguments may be passed by value or by reference as they are inherently compiler and hence implementation specific.

Data objects with a `CH R CTER` type are passed as a pointer to the character string and its length, so that each `CH R CTER` formal argument in a Fortran procedure results in two actual arguments in the equivalent C prototype. The first argument occupies the position in the formal argument list of the Fortran procedure. This argument is a pointer to the array of characters that make up the string, passed by the caller. The second argument is appended to the end of the user-specified formal argument list. This argument is of the default integer type and its value is the length of the array of characters, that is the length, passed as the first argument. This length is passed by value. When more than one `CH R CTER` argument is present in an argument list, the length arguments are appended in the order the original arguments appear. The above discussion also applies to sub-strings.

This BI does not define the passing of optional arguments. They are allowed only in Fortran 90/95 and their passing is implementation defined.

⁴This includes sub-strings.

This BI does not define array functions (function returning arrays). They are allowed only in Fortran 90/95 and requires the definition of array descriptors.

Note that Fortran 90/95 procedure arguments with the `INTENT(IN)` attribute should also be passed by reference if the procedure is to be linked with code written in Fortran 77. Fortran 77 does not and can not support the `INTENT` attribute because it has no concept of explicit interfaces. It is therefore not possible to declare the callee's arguments as `INTENT(IN)`. Fortran 77 compiler must assume that all procedure arguments are `INTENT(INOUT)` in the Fortran 90/95 sense.

9.2.4 Functions

The calling of statement functions is implementation defined (as they are defined only locally, the compiler has the freedom to apply any calling convention it likes).

Subroutines with alternate returns (e.g. "SUBROUTINE X(*,*)" called as "C LL X(*10,*20)") are implemented as functions returning an `INTEGER` of the default kind. The value of this returned integer is whatever integer is specified in the "RETURN" statement for the subroutine⁵, or 0 for a `RETURN` statement without an argument. It is up to the caller to jump to the corresponding alternate return label. The actual alternate-return arguments are omitted from the calling sequence.

In example:

```
SUBROUTINE SHOW_ LTERN TE_RETURN (N)
    INTEGER N
    C LL  LTERN TE_RETURN_EX MPLE (N, *10, *20, *30)
    WRITE (*,*) 'OK - Normal Return'
    RETURN
10    WRITE (*,*) '1st alternate return'
    RETURN
20    WRITE (*,*) '2nd alternate return'
    RETURN
30    WRITE (*,*) '2nd alternate return'
    RETURN
END

SUBROUTINE LTERN TE_RETURN_EX MPLE (N, *, *, *)
    INTEGER N
    IF (N .EQ. 0 ) RETURN      ! Implicit "RETURN 0"
    IF ( N .EQ. 1 ) RETURN 1
    IF ( N .EQ. 2 ) RETURN 2
    RETURN 3
END
```

⁵This integer indicates the position of an alternate return from the subroutine in the formal argument list

Here the `SUBROUTINE LTERN TE_RETURN_EX MPLE` is implemented as a function returning an `INTEGER*4` with value 0 if N is 0, 1 if N is 1, 2 if N is 2 and 3 for all other values of N. This return value is used by the caller as if the actual call were replaced by this sequence:

```
INTEGER X
X = C LL  LTERN TE_RETURN_EX MPLE (N)
GOTO (10, 20, 30), X
```

In all the effect is that the index of the returned to label (starting from 1) will be contained in `%rax` after the call.

Alternate `ENTRY` points of a `SUBROUTINE` or `FUNCTION` should be treated as separate sub-programs, as mandated by the Fortran standard. I.e. arguments passed to an alternate `ENTRY` should be passed as if the alternate `ENTRY` is a separate `SUBROUTINE` or `FUNCTION`. If a `FUNCTION` has alternate `ENTRY` points, the result of each of the alternate `ENTRY` points must be returned as if the alternate `ENTRY` is a separate `FUNCTION` with the result type of the alternate `ENTRY`. The external naming of alternate `ENTRY` points follows section 9.2.1.

9.2.5 COMMON blocks

In absence of any `EQUIV LENCE` declaration involving variables in `COMMON` blocks the layout of a `COMMON` block is exactly the same as the layout of the equivalent C structure (with types of variables substituted according to section 9.2.2), including the alignment requirements.

This BI defines the layout under presence of `EQUIV LENCE` statements only in some cases:

- the layout of the `COMMON` block must not change if one ignores the `EQUIV LENCE`, which amongst other things means:
 - If two arrays are equivalenced, the larger array must be named in the `COMMON` block, and there must be complete inclusion, in particular the other array may not extend the size of the equivalenced segment. It may also not change the alignment requirement.
 - If an array element and a scalar are equivalenced, the array must be named in the `COMMON` block and it must not be smaller than the scalar. The type of the scalar must not require bigger alignment than the array.
 - if two scalars are equivalenced they must have the same size and alignment requirements.

Other cases are implementation defined.

Because the Fortran standard allows the blank `COMMON` block to have different sizes in different subprograms, it may be impossible to determine if it is small enough to fit in the `.bss` section. When compiling for the medium or large code models the blank `COMMON` block should therefore always be put in the `.lbss` section.

9.2.6 Intrinsics

This sections lists the set of intrinsics which has to be supported at minimum by a conforming compiler. They are separated by origin. They follow regular calling and naming conventions.

The signature of intrinsics uses the syntax `return type argtype1,argtype2,...`, where the individual types can be the following characters: **V** (as in `void`) designates a `SUBROUTINE`, **L** a `LOGIC L`, **I** an `INTEGER`, **R** a `RE L`, and **C** a `CH R CTER`. Hence `I(R,L)` designates a `FUNCTION` returning an `INTEGER` and taking a `RE L` and a `LOGIC L`. If an argument is an array, this is indicated using a trailing number, e.g. **I13** is an `INTEGER` array with 13 elements. If a `CH R CTER` argument or return value has a fixed length, this is indicated using an asterisk and a trailing number, for example **C*16** is a `CH R CTER(len=16)`. If a `CH R CTER` argument of arbitrary length must be passed, the trailing number is replaced with **N**, for example **C*N**.

Table 9.1: Mil intrinsics

Name	Signature	Meaning
BTest	L(I,I)	Test bit
I_nd	I(I,I)	Boolean ND
IOr	I(I,I)	Boolean OR
IEOr	I(I,I)	Boolean XOR
Not	I(I)	Boolean NOT
IBClr	I(I,I)	Clear a bit
IBits	I(I,I,I)	Extract a bit subfield of a variable
IBSet	I(I,I)	Set a bit
IShft	I(I,I)	Logical bit shift
IShftC	I(I,I,I)	Circular bit shift
MvBits	V(I,I,I,I,I)	Move a bit field

BTest (I, Pos) Returns .TRUE. if bit `Pos` in `I` is set, returns .F LSE. otherwise.

I_and (I, J) Returns value resulting from a boolean AND on each pair of bits in `I` and `J`.

IOr (I, J) Returns value resulting from a boolean OR on each pair of bits in `I` and `J`.

IEOr (I, J) Returns value resulting from a boolean XOR on each pair of bits in `I` and `J`.

Not (I) Returns value resulting from a boolean NOT on each bit in `I`.

IBClr (I, Pos) Returns the value of `I` with bit `Pos` cleared (set to zero).

IBits (I, Pos, Len) Extracts a subfield starting from bit position `Pos` and with a length (towards the most significant bit) of `Len` bits from `I`. The result is right-justified and the remaining bits are zeroed.

IBSet (I, Pos) Returns the value of `I` with the bit in position `Pos` set to one.

IShft (I, Shift) `I` bits of `I` are shifted `Shift` places. `Shift.GT.0` indicates a left shift, `Shift.EQ.0` indicates no shift, and `Shift.LT.0` indicates a right shift. Bits shifted out from the least (when shifting right) or most (when shifting left) significant position are lost. Bits shifted in at the opposite end are not set (i.e. zero).

IShftC (I, Shift, Size) The rightmost `Size` bits of the argument `I` are shifted circularly `Shift` places. The unshifted bits of the result are the same as the unshifted bits of `I`.

MvBits (From, FromPos, Len, To, ToPos) Move `Len` bits of `From` from bit positions `FromPos` through `FromPos+Len-1` to bit positions `ToPos` through `ToPos+Len-1` of `To`. The bit portions of `To` that are not affected by the movement of bits are unchanged.

Table 9.2: F77 intrinsics

Name	Meaning
bs	bsolute value
Cos	rc cosine
Int	Truncate to whole number
NInt	Round to nearest whole number
Sin	rc sine
Tan	rc Tangent
Tan2	rc Tangent
Char	Character from code
Cmplx	Construct <code>COMPLEX(KIND=1)</code> value
Conjg	Complex conjugate
Cos	Cosine
CosH	Hyperbolic cosine
Dble	Convert to double precision
DiM	Difference magnitude (non-negative subtract)
DProd	Double-precision product
Exp	Exponential
IChar	Code for character
Index	Locate a <code>CH R CTER</code> substring
Int	Convert to <code>INTEGER</code> value truncated to whole number
Len	Length of character entity
LGe	Lexically greater than or equal
LGt	Lexically greater than
LLe	Lexically less than or equal
LLt	Lexically less than
Log	Natural logarithm
Log10	Common logarithm
Max	Maximum value
Min	Minimum value
Mod	Remainder
NInt	Convert to <code>INTEGER</code> value rounded to nearest whole number
Real	Convert value to type <code>RE L(KIND=1)</code>
Sin	Sine
SinH	Hyperbolic sine
SqRt	Square root
Tan	Tangent
TanH	Hyperbolic tangent

Refer to the Fortran 77 language standard for signature and definition of the F77 intrinsics listed in table 9.2. These intrinsics can have a prefix as per the standard hence the table is not exhaustive.

Table 9.3: F90 intrinsics

Name	Meaning
Char	SCII character from code
Bit_Size	Number of bits in arguments type
CPU_Time	Get current CPU time
I_Char	SCII code for character
Len_Trim	Get last non-blank character in string
System_Clock	Get current system clock value

Refer to the Fortran 90 language standard for signature and definition of the F90 intrinsics listed in table 9.3.

Table 9.4: Math intrinsics

Name	Signature	Meaning
BesJ0	R(R)	Bessel function
BesJ1	R(R)	Bessel function
BesJN	R(I,R)	Bessel function
BesY0	R(R)	Bessel function
BesY1	R(R)	Bessel function
BesYN	R(I,R)	Bessel function
ErF	R(R)	Error function
ErFC	R(R)	Complementary error function
IRand	I(I)	Random number
Rand	R(I)	Random number
SRand	V(I)	Random seed

BesJ0 (X) Calculates the Bessel function of the first kind of order 0 of X. Returns a RE L of the same kind as x.

BesJ1 (X) Calculates the Bessel function of the first kind of order 1 of X. Returns a RE L of the same kind as x.

BesJN (N, X) Calculates the Bessel function of the first kind of order N of X. Returns a RE L of the same kind as x.

BesY0 (X) Calculates the Bessel function of the second kind of order 0 of X. Returns a RE L of the same kind as x.

BesY1 (X) Calculates the Bessel function of the second kind of order 1 of X. Returns a RE L of the same kind as x.

BesYN (N, X) Calculates the Bessel function of the second kind of order N of X. Returns a RE L of the same kind as x.

ErF (X) Calculates the error function of X. Returns a RE L of the same kind as x.

ErFC (X) Calculates the complementary error function of X, i.e. $1 - \text{ERF}(X)$. Returns a RE L of the same kind as x.

IRand (Flag) Flag is optional. Returns a uniform quasi-random number up to a system-dependent limit. If `Flag .EQ. 0` or `Flag` is not passed, the next number in sequence is returned. If `Flag .EQ. 1`, the generator is restarted. If `Flag` has any other value, the generator is restarted with the value of `Flag` as the new seed.

Rand (Flag) Flag is optional. Returns a uniform quasi-random number between 0 and 1. If `Flag .EQ. 0` or `Flag` is not passed, the next number in sequence is returned. If `Flag .EQ. 1`, the generator is restarted. If `Flag` has any other value, the generator is restarted with the value of `Flag` as the new seed.

SRand (Seed) Reinitializes the random number generator for `IRand` and `Rand` with the seed in `Seed`.

Table 9.5: Unix intrinsics

Name	Signature	Meaning
bort	V()	bort the program
ccess	I(C,C)	Check file accessibility
DTime	V(R2,R)	Get elapsed time since last call
ETime	V(R2,R)	Get elapsed time for process
Flush	V(I)	Flush buffered output
FNum	I(I)	Get file descriptor from Fortran unit number
FStat	V(I,I13,I)	Get file information
GError	V(C*N)	Get error message for last error
Get rg	V(I,C*N)	Obtain command-line argument
GetCWD	V(C*N,I)	Get current working directory
GetEnv	V(C*N,C*N)	Get environment variable
GetGId	I()	Get process group ID
GetPId	I()	Get process ID
GetUID	I()	Get process user ID
GetLog	V(C*N)	Get login name
HostNm	V(C*N,I)	Get host name
I rgC	I()	Obtain count of command-line arguments
IDate	V(I3)	Get local date info
IErrNo	I()	Get error number for last error
ITime	V(I3)	Get local time of day
LStat	V(C*N,I13,I)	Get file information
PError	V(C*N)	Print error message for last error
Rename	V(C*N,C*N,I)	Rename file
Sleep	V(I)	Sleep for a specified time
System	V(C*N,I)	Invoke shell (system) command

bort () Prints a message and potentially causes a core dump.

ccess (Name, Mode) Checks file `Name` for accessibility in the mode specified by `Mode`. Returns 0 if the file is accessible in that mode, otherwise an error code. `Name` must be a NULL-terminated string of CH R CTER (i.e. a C-style string). Trailing blanks in `Name` are ignored. `Mode` must be a concatenation of any of the following characters: **r** meaning test for read permission, **w** meaning test for write permission, **x** meaning test for execute/search permission, or a space meaning test for existence of the file.

DTime (*T rray, Result*) When called for the first time, returns the number of seconds of runtime since the start of the program in *Result*, the user component of this runtime in *T rray(1)*, and the system time in *T rray(2)*. Subsequent invocations values based on accumulations since the previous invocation.

ETime (*T rray, Result*) Returns the number of seconds of runtime since the start of the program in *Result*, the user component of this runtime in *T rray(1)*, and the system time in *T rray(2)*. Subsequent invocations values based on accumulations since the previous invocation.

Flush (*Unit*) Flushes the Fortran I/O unit with ID *Unit*. The unit must be open for output. If the optional *Unit* argument is omitted, all open units are flushed.

FNum (*Unit*) Returns the UNIX(tm) file descriptor number corresponding to the Fortran I/O unit *Unit*. The unit must be open.

FStat (*Unit, S rray, Status*) Obtains data about the file open on Fortran I/O unit *Unit* and places it in the array *S rray*. The values in this array are as follows:

1. Device ID
2. Inode number
3. File mode
4. Number of links
5. Owner's UID
6. Owner's GID
7. ID of device containing directory entry for file
8. File size (bytes)
9. Last access time
10. Last modification time
11. Last file status change time
12. Preferred I/O block size (-1 if not available)
13. Number of blocks allocated (-1 if not available)

If an element is not available, or not relevant on the host system, it is returned as 0 except when indicated otherwise in the above list. If the optional *Status* argument is supplied, it contains 0 on success or a nonzero error code upon return.

`Gerror (Message)` Returns the system error message corresponding to the last system error (`errno` in C). The message is returned in `Message`. If `Message` is longer than the error message, it is padded with blanks after the message. If `Message` is not long enough to hold the error message, the error message is truncated to the length of `Message`.

`Get rg (Pos, Value)` Returns in `Value` the command-line argument in position `Pos`. If there are fewer than `Pos` command-line arguments, `Value` is filled with blanks. If `Pos` is 0, the name of the program is returned. If `Value` is longer than the command-line argument, it is padded with blanks after the argument. If `Value` is not long enough to hold the command-line argument, the argument is truncated to the length of `Value`.

`GetCWD (Name, Status)` Returns in `Name` the current working directory. If the optional `Status` argument is supplied, it contains 0 on success or a nonzero error code upon return.

`GetEnv (Name, Value)` Returns in `Value` the environment variable identified with `Name`. If `Name` has not been set, `Value` is filled with blanks. `null` character marks the end of the name in `Name`. Trailing blanks in `Name` are ignored. If `Value` is longer than the environment variable, it is padded with blanks after the variable. If `Value` is not long enough to hold the environment variable, the variable is truncated to the length of `Value`.

`GetGId ()` Returns the group ID for the current process.

`GetPId ()` Returns the process ID for the current process.

`GetUID ()` Returns the user ID for the current process.

`GetLog (Login)` Returns the login name for the process in `Login`, or a blank string if the host system does not support `getlogin(3)`. If `Login` is longer than the login name, it is padded with blanks after the login name. If `Login` is not long enough to hold the login name, the login name is truncated to the length of `Login`.

`HostNm (Name, Status)` Returns in `Name` system's host name. If the optional `Status` argument is supplied, it contains 0 on success or a nonzero error code upon return. If `Name` is longer than the host name, it is padded with blanks after the host name. If `Name` is not long enough to hold the host name, the host name is truncated to the length of `Name`.

`I rgC ()` Returns the number of command-line arguments. The program name itself is not included in this number.

IDate (*T rray*) Returns the current local date day, month, year in elements 1, 2, and 3 of *Tarray*, respectively. The year has four significant digits.

IErrno () Returns the last system error number (*errno* in C).

ITime (*T rray*) Returns the current local time hour, minutes, and seconds in elements 1, 2, and 3 of *T rray*, respectively.

LStat (*File*, *S rray*, *Status*) Obtains data about a file named *File* and places it in the array *S rray*. The values in this array are as follows:

1. Device ID
2. Inode number
3. File mode
4. Number of links
5. Owner's UID
6. Owner's GID
7. ID of device containing directory entry for file
8. File size (bytes)
9. Last access time
10. Last modification time
11. Last file status change time
12. Preferred I/O block size (-1 if not available)
13. Number of blocks allocated (-1 if not available)

If an element is not available, or not relevant on the host system, it is returned as 0 except when indicated otherwise in the above list. If the optional *Status* argument is supplied, it contains 0 on success or a nonzero error code upon return.

PError (*MsgPrefix*) Prints a newline-terminated error message corresponding to the last system error. This is prefixed by the string *MsgPrefix*, a colon and a space. The error message is printed on the C *stderr* stream.

Rename (*Path1*, *Path2*, *Status*) Renames the file named *Path1* to *Path2*. null character marks the end of the names. Trailing blanks are ignored. If the optional *Status* argument is supplied, it contains 0 on success or a nonzero error code upon return.

`Sleep (Seconds)` Causes the program to pause for `Seconds` seconds.

`System (Command, Status)` Passes the string in `Command` to a shell though `system(3)`. If the optional argument `Status` is present, it contains the value returned by `system(3)`.

Chapter 10

ILP32 Programming Model

"x32" is commonly used to refer to MD64 ILP32 programming model.

10.1 Parameter Passing

When a value of pointer type is returned or passed in a register, bits 32 to 63 shall be zero.

10.2 Address Space

ILP32 binaries reside in the lower 32 bits of the 64-bit virtual address space and all addresses are 32 bits in size. They should conform to small code model or small position independent code model (PIC) described in Section 3.5.1.

10.3 Thread-Local Storage Support

ILP32 Thread-Local Storage (TLS) support is based on LP64 TLS implementation with some modifications.

10.3.1 Global Thread-Local Variable

For a global thread-local variable x:

```
extern __thread int x;
```

General Dynamic Model Load address of x into $\%rax$

Table 10.1: General Dynamic Model Code Sequence

LP64			ILP32		
0x00	.byte	0x66	0x00	leaq	$x@tlsqd(\%rip), \%rdi$
0x01	leaq	$x@tlsqd(\%rip), \%rdi$	0x07	.word	0x6666
0x08	.word	0x6666	0x09	rex64	
0x0a	rex64		0x0a	call	$__tls_get_addr@plt$
0x0b	call	$__tls_get_addr@plt$			

In TLSDESC code sequence, `leal` instruction must be encoded with `rex` prefix even if it isn't required by destination register. If the `leal` encoding has a variable length, linker can't tell where it starts and can't safely perform GDesc -> IE/LE optimization.

Table 10.2: General Dynamic Model Code Sequence with TLSDESC

LP64			ILP32		
0x00	leaq	$x@tlsdesc(\%rip), \%rax$	0x00	rex leal	$x@tlsdesc(\%rip), \%eax$
0x07	call	$*x@tlsdesc(\%rax)$	0x07	call	$*x@tlsdesc(\%eax)$
0x08	add	$%fs:0x0, \%eax$	0x09	add	$%fs:0x0, \%eax$

Initial Exec Model Load address of x into $\%rax$. Instruction `addl` must be encoded with `rex` prefix even if it isn't required by destination register. Otherwise linker can't safely perform IE -> LE optimization.

Table 10.3: Initial Exec Model Code Sequence

LP64			ILP32		
0x00	movq	$%fs:0, \%rax$	0x00	movl	$%fs:0, \%eax$
0x09	addq	$x@gottpoff(\%rip), \%rax$	0x08	rex addl	$x@gottpoff(\%rip), \%eax$

Initial Exec Model, II Load value of x into $\%edi$. $\%fs:(\%eax)$ memory operand can't be used for ILP32 since its effective address is the base address of $\%fs$ + value of $\%eax$ zero-extended to a 64-bit result, which is incorrect with negative value in $\%eax$.

Table 10.4: Initial Exec Model Code Sequence, II

LP64				ILP32			
0x00	movq	x@gottpoff(%rip),%rax		0x00	movq	x@gottpoff(%rip),%rax	
0x07	movl	%fs:(%rax),%edi		0x07	movl	%fs:(%rax),%edi	

10.3.2 Static Thread-Local Variable

For a static thread-local variable x:

```
static __thread int x;
```

Local Dynamic Model Load address of x into %rax

Table 10.5: Local Dynamic Model Code Sequence With Lea

LP64				ILP32			
0x00	leaq	x@tlsld(%rip),%rdi		0x00	leaq	x@tlsld(%rip),%rdi	
0x07	call	__tls_get_addr@plt		0x07	call	__tls_get_addr@plt	
0x0c	leaq	x@dtppoff(%rax),%rax		0x0c	leal	x@dtppoff(%rax),%eax	

or

Table 10.6: Local Dynamic Model Code Sequence With dd

LP64				ILP32			
0x00	leaq	x@tlsld(%rip),%rdi		0x00	leaq	x@tlsld(%rip),%rdi	
0x07	call	__tls_get_addr@plt		0x07	call	__tls_get_addr@plt	
0x0c	addq	\$x@dtppoff,%rax		0x0c	addl	\$x@dtppoff,%eax	

For code sequence with TLSDESC, local dynamic model is similar to general dynamic model. The same encoding requirement for leal instruction also applies.

Table 10.7: General Dynamic Model Code Sequence with TLSDESC

LP64				ILP32			
0x00	leaq	x@tlsdesc(%rip),%rax		0x00	rex leal	x@tlsdesc(%rip),%eax	
0x07	call	*x@tlsdesc(%rax)		0x07	call	*x@tlsdesc(%eax)	
0x08	add	%fs:0x0,%eax		0x09	add	%fs:0x0,%eax	

Local Dynamic Model, II Load value of x into $\%edi$

Table 10.8: Local Dynamic Model Code Sequence, II

LP64				ILP32			
0x00	movq	%fs:0,%rax		0x00	movl	%fs:0,%eax	
0x09	movl	x@dtpoff(%rax),%edi		0x08	movl	x@dtpoff(%rax),%edi	

Local Exec Model Load address of x into $\%rax$

Table 10.9: Local Exec Model Code Sequence With Lea

LP64				ILP32			
0x00	movq	%fs:0,%rax		0x00	movl	%fs:0,%eax	
0x09	leaq	x@tpoff(%rax),%rax		0x08	leal	x@tpoff(%rax),%eax	

or

Table 10.10: Local Exec Model Code Sequence With dd

LP64				ILP32			
0x00	movq	%fs:0,%rax		0x00	movl	%fs:0,%eax	
0x09	addq	\$x@tpoff,%rax		0x08	addl	\$x@tpoff,%eax	

Local Exec Model, II Load value of x into $\%edi$

Table 10.11: Local Exec Model Code Sequence, II

LP64				ILP32			
0x00	movq	%fs:0,%rax		0x00	movl	%fs:0,%eax	
0x09	movl	x@tpoff(%rax),%edi		0x08	movl	x@tpoff(%rax),%edi	

Local Exec Model, III Load value of x into $\%edi$

Table 10.12: Local Exec Model Code Sequence, III

LP64				ILP32			
0x00	movl	%fs:x@tpoff,%edi		0x00	movl	%fs:x@tpoff,%edi	

10.3.3 TLS Linker Optimization

General Dynamic To Initial Exec Load address of x into $\%rax$

Table 10.13: GD -> IE Code Transition

GD				IE			
0x00	leaq	x@tlsqd(%rip),%rdi		0x00	movl	%fs:0,%eax	
0x07	.word	0x6666		0x08	addq	x@gottpoff(%rip),%rax	
0x09	rex64						
0x0a	call	__tls_get_addr@plt					

Table 10.14: GDesc -> IE Code Transition

GDesc				LE			
0x00	rex leal	x@tlsdes(%rip),%eax		0x00	rex movl	x@gottpoff(%rip),%eax	
0x07	call	*foo@TLSC LL(%eax)		0x07	nopl	(%rax)	

Table 10.15: GD -> LE Code Transition

	GD			LE		
0x00	leaq	x@tlsqd(%rip),%rdi		0x00	movl	%fs:0,%eax
0x07	.word	0x6666		0x08	leal	x@tpoff(%rax),%eax
0x09	rex64					
0x0a	call	__tls_get_addr@plt				

Table 10.16: GDesc -> LE Code Transition

	GDesc			LE		
0x00	rex	leal	x@tlsdes(%rip),%eax	0x00	rex	movl \$x@tpoff,%eax
0x07	call	*	foo@TLSC LL(%eax)	0x07	nopl	(%rax)

Initial Exec To Local Exec Load address of `x` into `%rax`

Table 10.17: IE -> LE Code Transition With Lea

	IE			LE		
0x00	movl	%fs:0,%eax		0x00	movl	%fs:0,%eax
0x08	rex	addl	x@gottpoff(%rip),%eax	0x08	rex	leal x@tpoff(%rax),%eax

or

Table 10.18: IE -> LE Code Transition With dd

	IE			LE		
0x00	movl	%fs:0,%r12d		0x00	movl	%fs:0,%r12d
0x09	addl	x@gottpoff(%rip),%r12d		0x09	addl	\$x@tpoff,%r12d

Initial Exec To Local Exec, II Load value of `x` into `%edi`.

Table 10.19: IE -> LE Code Transition, II

IE			LE		
0x00	movq	x@gottpoff(%rip),%rax	0x00	movq	x@tpoff,%rax
0x07	movl	%fs:(%rax),%edi	0x07	movl	%fs:(%rax),%edi

Local Dynamic to Local Exec Load address of `x` into `%rax`

Table 10.20: LD -> LE Code Transition With Lea

LD			LE		
0x00	leaq	x@tlsld(%rip),%rdi	0x00	nopl	0x0(%rax)
0x07	call	__tls_get_addr@plt	0x04	movl	%fs:0,%eax
0x0c	leal	x@dtppoff(%rax),%eax	0x0c	leal	x@tpoff(%rax),%eax

or

Table 10.21: LD -> LE Code Transition With dd

LD			LE		
0x00	leaq	x@tlsld(%rip),%rdi	0x00	nopl	0x0(%rax)
0x07	call	__tls_get_addr@plt	0x04	movl	%fs:0,%eax
0x0c	addq	\$x@dtppoff,%rax	0x0c	addl	\$x@tpoff,%eax

Local Dynamic To Local Exec, II Load value of `x` into `%edi`.

Table 10.22: LD -> LE Code Transition, II

LD			LE		
0x00	leaq	x@tlsld(%rip),%rdi	0x00	nopl	0x0(%rax)
0x07	call	__tls_get_addr@plt	0x04	movl	%fs:0,%eax
0x0c	movl	x@dtppoff(%rax),%eax	0x0c	movl	x@tpoff(%rax),%eax

10.4 Kernel Support

Kernel should limit stack and addresses returned from system calls bewteen 0x00000000 to 0xffffffff.

10.5 Coding Examples

lthough ILP32 binaries run in the 64-bit mode, not all 64-bit instructions are supported. This section discusses example code sequences for fundamental operations which are different from the 64-bit mode.

10.5.1 Indirect Branch

Since indirect branch via memory loads a 64-bit address at the memory location, it is not supported in ILP32. Indirect branch via register should be used instead. The 32-bit address from memory is loaded into the lower 32 bits of a register, which will automatically zero-extend the upper 32 bits of the register. Then the indirect call can be performed via the 64-bit register.

Table 10.23: Indirect Branch

LP64	ILP32
call *%rax	call *%rax
call *func_p(%rip)	movl func_p(%rip), %eax call *%rax
call *func_p	movl func_p, %eax call *%rax
jmp *%rax	jmp *%rax
jmp *func_p(%rip)	movl func_p(%rip), %eax jmp *%rax
jmp *func_p	movl func_p, %eax jmp *%rax

Chapter 11

Iternate Code Sequences For Security

11.1 Code Sequences without PLT

Procedure Linkage Table (PLT), see Section 5.2 for detail, is used to access external functions defined in shared object and support

Lazy symbol resolution The function address is resolved only when it is called the first time at run-time.

Canonical function address The PLT entry of the external function is used as its address, aka function pointer.

The first instruction in the PLT entry is an indirect branch via the Global Offset Table (GOT), see Section 5.2 for detail, entry of the external function, which is set up in such a way that it will be updated to the address of the function body the first time when the function is called. Since the GOT entry is writable, any address may be written to it at run-time, which is a potential security risk.

11.1.1 Indirect Call via the GOT Slot

For small and medium models, different code sequence is used to avoid PLT:

Figure 11.1: Function Call without PLT (Small and Medium Models)

extern void func (void); func (void);	.globl func call *func@GOTPCREL(%rip)
--	--

The direct branch is replaced by an indirect branch via the GOT slot, which is similar to the first instruction in the PLT slot.

Figure 11.2: Function address without PLT (Small and Medium Models)

extern void func (void); void* ptr (void) { return func; }	.globl func func: movq func@GOTPCREL(%rip), %rax ret
--	---

Instead using the PLT slot as function address, the function address is retrieved from the GOT slot.

If linker determines the function is defined locally, it converts indirect branch via the GOT slot to direct branch with a `nop` prefix and converts load via the GOT slot to load immediate or `lea`, see Section B.2 for details.

After dynamic linker resolved all symbols by updating GOT entries with symbol addresses, GOT can be made read-only and overwriting GOT becomes a hard error immediately. Since PLT is no longer used to call external function, lazy symbol resolution is disabled and a function can only be interposed during symbol resolution at startup. Tools and features which depend on lazy symbol resolution will not work properly. However, there are also a few side benefits:

No extra direct branch to PLT entry Since indirect branch is 6 byte long and direct branch is 5 byte long, when indirect branch via the GOT slot is used to call a local function, code size will be increased by one byte for each call. Since one PLT slot has 16 bytes, there will be code size increase when indirect branch via the GOT slot is used to call an external function more than 16 times.

Custom calling convention Since external function is called directly via the GOT slot, instead of invoking dynamic linker to lookup function symbol when called the first time, parameters can be passed differently from what is specified in this document.

11.1.2 Thread-Local Storage without PLT

TLS code sequences for general and local dynamic models can be updated to replace direct call to `__tls_get_addr` via the PLT entry, with indirect call to `__tls_get_addr` via the GOT slot, see Figure 11.3. Since direct `call` instruction is 4-byte long and indirect `call` instruction is 5-byte long, the extra one byte must be handled properly.

Figure 11.3: `__tls_get_addr` Call

Direct via PLT	Indirect via GOT
<code>call __tls_get_addr@PLT</code>	<code>call *__tls_get_addr@GOTPCREL(%rip)</code>

General Dynamic Model for Global Variable

For general dynamic model, one `0x66` prefix before `call` instruction is removed to make room for indirect call:

```
extern __thread int x;
```

the following alternate code sequence loads address of `x` into `%rax` without PLT:

Table 11.1: General Dynamic Model Code Sequence (LP64)

With PLT			Without PLT		
0x00	.byte	0x66	0x00	.byte	0x66
0x01	leaq	<code>x@tlsqd(%rip),%rdi</code>	0x01	leaq	<code>x@tlsqd(%rip),%rdi</code>
0x08	.word	0x6666	0x08	.byte	0x66
0x0a	rex64		0x09	rex64	
0x0b	call	<code>__tls_get_addr@PLT</code>	0xa	call	<code>*__tls_get_addr@GOTPCREL(%rip)</code>

Table 11.2: General Dynamic Model Code Sequence (ILP32)

With PLT			Without PLT		
0x00	leaq	<code>x@tlsqd(%rip),%rdi</code>	0x00	leaq	<code>x@tlsqd(%rip),%rdi</code>
0x07	.word	0x6666	0x07	.byte	0x66
0x09	rex64		0x08	rex64	
0xa	call	<code>__tls_get_addr@PLT</code>	0x09	call	<code>*__tls_get_addr@GOTPCREL(%rip)</code>

Static Thread-Local Variable

For local dynamic model, indirect call is used instead of direct call:

```
static __thread int x;
```

the following alternate code sequence loads the address of the TLS block of the module, which contains variable `x`, into `%rax` without PLT:

Table 11.3: Local Dynamic Model Code Sequence (LP64)

With PLT				Without PLT			
0x00	leaq	x@tlsld(%rip),%rdi		0x00	leaq	x@tlsld(%rip),%rdi	
0x07	call	__tls_get_addr@PLT		0x07	call	*__tls_get_addr@GOTPCREL(%rip)	

Table 11.4: Local Dynamic Model Code Sequence (ILP32)

With PLT				Without PLT			
0x00	leaq	x@tlsld(%rip),%rdi		0x00	leaq	x@tlsld(%rip),%rdi	
0x07	call	__tls_get_addr@PLT		0x07	call	*__tls_get_addr@GOTPCREL(%rip)	

TLS Linker Optimization

Since the code sequence with indirect call for general dynamic model has the same length as the one with direct call, linker just needs to recognize new instruction pattern to convert general dynamic access to initial exec or local exec accesses.

General Dynamic to Initial Exec To load address of `x` into `%rax`:

Table 11.5: GD -> IE Code Transition (LP64)

GD				IE			
0x00	.byte	0x66		0x00	movq	%fs:0,%rax	
0x01	leaq	x@tlsqd(%rip),%rdi		0x09	addq	x@gottpoff(%rip),%rax	
0x09	.byte	0x66					
0x0a	rex64						
0x0b	call	*__tls_get_addr@GOTPCREL(%rip)					

Table 11.6: GD -> IE Code Transition (ILP32)

GD			IE		
0x00	leaq	x@tlsqd(%rip),%rdi	0x00	movl	%fs:0, %eax
0x07	.byte	0x66	0x08	addq	x@gottpoff(%rip),%rax
0x08	rex64				
0x0a	call	*__tls_get_addr@GOTPCREL(%rip)			

General Dynamic to Local Exec To load address of `x` into `%rax`:

Table 11.7: GD -> LE Code Transition (LP64)

GD			LE		
0x00	.byte	0x66	0x00	movq	%fs:0, %rax
0x01	leaq	x@tlsqd(%rip),%rdi	0x09	leaq	x@tpoff(%rax),%rax
0x08	.byte	0x66			
0x09	rex64				
0x0a	call	*__tls_get_addr@GOTPCREL(%rip)			

Table 11.8: GD -> LE Code Transition (ILP32)

GD			LE		
0x00	leaq	x@tlsqd(%rip),%rdi	0x00	movl	%fs:0, %eax
0x07	.byte	0x66	0x08	leaq	x@tpoff(%rax),%rax
0x08	rex64				
0x09	call	*__tls_get_addr@GOTPCREL(%rip)			

Local Dynamic to Local Exec For local dynamic model to local exec model transition, linker generates 4 `0x66` prefixes, instead of 3, before `mov` instruction for LP64 and generate a 5-byte `nop`, instead of 4-byte, before `mov` instruction for ILP32. To load the address of the TLS block of the module, which contains variable `x`, into `%rax` without PLT:

Table 11.9: LD -> LE Code Transition (LP64)

LD				LE			
0x00	leaq	x@tlsld(%rip),%rdi		0x00	.long	0x66666666	
0x07	call	*__tls_get_addr@GOTPCREL(%rip)		0x04	movq	%fs:0,%rax	

Table 11.10: LD -> LE Code Transition (ILP32)

LD				LE			
0x00	leaq	x@tlsld(%rip),%rdi		0x00	nopw	0x0(%rax)	
0x07	call	*__tls_get_addr@GOTPCREL(%rip)		0x05	movl	%fs:0,%eax	

Chapter 12

Intel MPX Extension

Intel MPX (Memory Protection Extensions) provides 4 128-bit wide bound registers (%bnd0 - %bnd3). For purposes of parameter passing and function return, the lower 64 bits of %bndN specify lower bound of the corresponding parameter, and the upper 64 bits specify upper bound of the parameter. The upper bound is represented in one's complement form.¹

12.1 Parameter Passing and Returning of Values

POINTER class is added for passing and returning pointer types and **INTEGER** class is updated as follows:

POINTER This class consists of pointer types.

INTEGER This class consists of integral types (except pointer types) that fit into one of the general purpose registers.

12.1.1 Classification

Pointers and integers are classified as:

- Pointers are in the **POINTER** class.
- Arguments of types (signed and unsigned) `_Bool`, `char`, `short`, `int`, `long` and `long long` are in the **INTEGER** class.

¹MPX is not supported under ILP32 since MPX requires 64-bit pointers.

Figure 12.1: Bound Register Usage

Register	Usage	Preserved across function calls
%bnd0–%bnd3	used to pass/return bounds of pointer arguments/return values	No

The classification of aggregate (structures and arrays) and union types is updated as follows:

1. When a C++ object is passed by invisible reference, the object is replaced in the parameter list by a pointer that has class POINTER.
2. If one of the classes is POINTER, the result is the POINTER.

12.1.2 Passing

For parameter passing, if the class is INTEGER or POINTER, the next available register of the sequence %rdi, %rsi, %rdx, %rcx, %r8 and %r9 is used.

Bounds passing

Intel MPX provides IS extensions that allow passing bounds for a pointer argument that specify memory area that may be legally accessed by dereferencing the pointer. This paragraph describes how the bounds are passed to the callee.

Several functions used in the description below are defined as follows:

BOUND_M_P_STORE(bnd, addr, ptr) This function executes Intel MPX `bndstx` instruction. `ptr` argument is used to initialize index field of the memory operand of the `bndstx` instruction, `addr` is encoded in base and/or displacement fields of the memory operand, `bnd` is encoded in the register operand.

BOUND_M_P_LO_D(addr, ptr) This function executes Intel MPX `bndlidx` instruction. `ptr` argument is used to initialize index field of the memory operand of the `bndlidx` instruction, `addr` is encoded in base and/or displacement fields of the memory operand.

The following algorithm is used to decide how bounds are passed for each eightbyte:

1. If the class is INTEGER, the eightbyte is passed in a general purpose register and the function being called uses varargs or stdarg, then the class is converted to POINTER. Artificial bounds that allow accessing all memory are created for pointers contained in the eightbyte.
2. If the class is POINTER and the eightbyte is passed in a general purpose register, then the bounds associated with the pointers contained in the eightbyte are passed in the next available registers from the sequence `%bnd0`, `%bnd1`, `%bnd2` and `%bnd3`. If there are no bound registers available for the bounds of an argument, then the bounds are passed in a CPU defined manner by executing `BOUND_M P_STORE(bnd, addr, ptr)` function, where `bnd` is the current bounds of the argument, `addr` is the address of stack location beyond the location of the callee's return address (that will be put on stack by the corresponding call instruction), `ptr` is the actual value of the pointer argument. For each call, there may be up to two such pointer arguments, the first one has its bounds associated with (*<return address stack location>-8*) address, and the second one - with (*<return address stack location>-16*).
3. If the class is POINTER and the eightbyte is passed on the stack, or the class is MEMORY and the argument contains pointer members, then the bounds associated with each pointer contained in the eightbyte are passed in a CPU defined manner by executing `BOUND_M P_STORE(bnd, addr, ptr)` function, where `bnd` is the current bounds of the pointer argument, `addr` is the address of the pointer argument's stack location, `ptr` is the actual value of the pointer argument. If the eightbyte may contain parts of partially overlapping pointers, then bounds associated with the pointers are ignored and special bounds that allow accessing all memory are passed for such pointers.

The callee uses the same algorithm to classify the incoming parameters. If a parameter is passed to the callee using `BOUND_M P_STORE`, then the callee fetches the passed bounds using `BOUND_M P_LOAD(addr, ptr)`, where `addr` is the same address passed to the corresponding `BOUND_M P_STORE` in the caller, and `ptr` is the actual value of the pointer parameter fetched by the callee either from a general purpose register or from a stack location.

When passing arguments with bounds to functions, function prototypes must be provided. Otherwise, the run-time behavior is undefined.

12.1.3 Returning of Values

The returning of values is updated:

- If the class is INTEGER or POINTER, the next available register of the sequence %rax, %rdx is used.

Returning of Bounds

The returning of bounds is done according to the following algorithm:

1. Classify the return type with the classification algorithm.
2. If the type has class MEMORY, on return %bnd0 must contain bounds of the “hidden” first argument that has been passed in by the caller in %rdi.
3. If the class is POINTER, the next available register of the sequence %bnd0, %bnd1, %bnd2, %bnd3 is used to return bounds of the pointers contained in the eightbyte.

As an example of bound passing conventions, consider the declaration and the function call is show in Figure 12.2. The corresponding bound registers allocation is given in Figure 12.3, the stack frame offset given shows the frame before calling the function.

Figure 12.2: Bounds Passing Example

```
extern void func (int *p1, int *p2, int *p3,
                  int *p4, int *p5, int x,
                  int *p6);

func(p1, p2, p3, p4, p5, x, p6);
```

Figure 12.3: Bounds Allocation Example

Bound Registers	Stack Frame Offset for BOUND_M_P_STORE	General Purpose Registers	Stack Frame Offset
%bnd0: p1	-16: p5 ²	%rdi: p1	0: p6
%bnd1: p2	-24: p6	%rsi: p2	
%bnd2: p3		%rdx: p3	
%bnd3: p4		%rcx: p4	
		%r8: p5	
		%r9: x	

12.1.4 Variable Argument Lists

Passing variable arguments is updated as follows:

The Register Save rea

If a function taking a variable argument list is compiled for Intel MPX, then the bounds passed for the argument registers saved to the register save area are saved for each argument register in the prolog by executing `BOUND_M P_STORE(bnd, addr, ptr)` function (12.1.2), where `bnd` is the current bounds of the pointer argument, `addr` is the address of the argument register's location in register save area and `ptr` is the actual value of the argument register.

The va_arg Macro

The `va_arg(l, type)` implementation is updated as follows:

1. Fetch `type` from `l->reg_save_area` with an offset of `l->gp_offset` and/or `l->fp_offset`. This may require copying to a temporary location in case the parameter is passed in different register classes or requires an alignment greater than 8 for general purpose registers and 16 for XMM registers. If `type` specifies a pointer, then the bounds of the argument being fetched are loaded by executing `BOUND_M P_LO D(l->reg_save_area + l->gp_offset, ptr)` (12.1.2), where `ptr` is the actual value fetched from `l->reg_save_area` with an offset of `l->gp_offset`.

12.2 Program Loading and Dynamic Linking

To preserve bound registers for symbol lookup in branch instructions with the BND (0xf2) prefix, linker should generate the BND procedure linkage table (see figure 12.4) together with an additional procedure linkage table (see figure 12.5).

²Before the call to `func()` the return address of the call is not yet put on the stack, thus offset -16 accounts for the push of the return address that will be made by the call instruction.

Figure 12.4: BND Procedure Linkage Table (small and medium models)

```
.PLT0: pushq    GOT+8(%rip)          # GOT[1]
       bnd jmp   *GOT+16(%rip)        # GOT[2]
       nopl    (%rax)
.PLT1: pushq    $index1              # 16 bytes from .PLT0
       bnd jmp   PLT0
       nopl    0(%rax,%rax,1)
.PLT2: pushq    $index2              # 16 bytes from .PLT1
       bnd jmp   .PLT0
       nopl    0(%rax,%rax,1)
.PLT3: ...
```

Figure 12.5: dditional Procedure Linkage Table (small and medium models)

```
. PLT1: bnd jmp  name1@GOTPCREL(%rip)
         nop
. PLT2: bnd jmp  name2@GOTPCREL(%rip)      # 8 bytes from . PLT1
         nop
. PLT3: ...                                # 8 bytes from . PLT2
```

To support indirect branches with the BND (0xf2) prefix (see figure 12.6), branches in all BND procedure linkage table entries must have the BND (0xf2) prefix.

Figure 12.6: Indirect branch

```
movq      fp@GOTPCREL(%rip), %rax
bnd jmp  *(%rax)
...
.globl  fp
.section .data.rel,"aw",@progbits
.align   8
.type    fp, @object
.size    fp, 8
fp:
.quad    memcpy
```

When the BND procedure linkage table is used, the initial value of the global offset table entry for the external function is the address of the corresponding entry of the additional procedure linkage table.

ppendix

Linux Conventions

This chapter describes some details that are only relevant to GNU/Linux systems and the Linux kernel.

.1 Execution of 32-bit Programs

The MD64 processors are able to execute 64-bit MD64 and also 32-bit ia32 programs. Libraries conforming to the Intel386 BI will live in the normal places like /lib, /usr/lib and /usr/bin. Libraries following the MD64, will use lib64 subdirectories for the libraries, e.g /lib64 and /usr/lib64. Programs conforming to Intel386 BI and to the MD64 BI will share directories like /usr/bin. In particular, there will be no /bin64 directory.

.2 MD64 Linux Kernel Conventions

The section is informative only.

.2.1 Calling Conventions

The Linux MD64 kernel uses internally the same calling conventions as user-level applications (see section 3.2.3 for details). User-level applications that like to call system calls should use the functions from the C library. The interface between the C library and the Linux kernel is the same as for the user-level applications with the following differences:

1. User-level applications use as integer registers for passing the sequence `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8` and `%r9`. The kernel interface uses `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8` and `%r9`.
2. system-call is done via the `syscall` instruction. The kernel destroys registers `%rcx` and `%r11`.
3. The number of the syscall has to be passed in register `%rax`.
4. System-calls are limited to six arguments, no argument is passed directly on the stack.
5. Returning from the `syscall`, register `%rax` contains the result of the system-call. value in the range between -4095 and -1 indicates an error, it is `-errno`.
6. Only values of class INTEGER or class MEMORY are passed to the kernel.

2.2 Stack Layout

The Linux kernel may align the end of the input argument area to a 8, instead of 16, byte boundary. It does not honor the red zone (see section [3.2.2](#)) and therefore this area is not allowed to be used by kernel code. Kernel code should be compiled by GCC with the option `-mno-red-zone`.

2.3 Miscellaneous Remarks

Linux Kernel code is not allowed to change the x87 and SSE units. If those are changed by kernel code, they have to be restored properly before sleeping or leaving the kernel. On preemptive kernels also more precautions may be needed.

ppendix B

Linker Optimization

This chapter describes optimizations which may be performed by linker.

B.1 Combine GOTPLT and GOT Slots

In the small and medium models, when there are both PLT and GOT references to the same function symbol, normally linker creates a GOTPLT slot for PLT entry and a GOT slot for GOT reference. run-time JUMP_SLOT relocation is created to update the GOT-PLT slot and a run-time GLOB_D_T relocation is created to update the GOT slot. Both JUMP_SLOT and GLOB_D_T relocations apply the same symbol value to GOTPLT and GOT slots, respectively, at run-time.

As an optimization, linker may combine GOTPLT and GOT slots into a single GOT slot and remove the run-time JUMP_SLOT relocation. It replaces the regular PLT entry:

Figure B.1: Procedure Linkage Table Entry Via GOTPLT Slot

```
.PLT: jmp      [GOTPLT slot]
      pushq   relocation index
      jmp     .PLT0
```

with an GOT PLT entry with an indirect jump via the GOT slot:

Figure B.2: Procedure Linkage Table Entry Via GOT Slot

```
.PLT: jmp      [GOT slot]
      nop
```

and resolves the PLT reference to the GOT PLT entry. Indirect `jmp` is an 5-byte instruction. `nop` can be encoded as a 3-byte instruction or a 11-byte instruction for 8-byte or 16-byte PLT slot. separate PLT with 8-byte slots may be used for this optimization.

This optimization isn't applicable to the `STT_GNU_IFUNC` symbols since their GOT-PLT slots are resolved to the selected implementation and their GOT slots are resolved to their PLT entries.

This optimization must be avoided if pointer equality is needed since the symbol value won't be cleared in this case and the dynamic linker won't update the GOT slot. Otherwise, the resulting binary will get into an infinite loop at run-time.

B.2 Optimize GOTPCRELX Relocations

The MD64 instruction encoding supports converting certain instructions on memory operand with `R_X86_64_GOTPCRELX` or `R_X86_64_REX_GOTPCRELX` relocations against symbol, `foo`, into a different form on immediate operand if `foo` is defined locally.

Convert call and jmp Convert memory operand of `call` and `jmp` into immediate operand.

Table B.1: Call and Jmp Conversion

Memory Operand	Immediate Operand
<code>call *foo@GOTPCREL(%rip)</code>	<code>nop call foo</code>
<code>call *foo@GOTPCREL(%rip)</code>	<code>call foo nop</code>
<code>jmp *foo@GOTPCREL(%rip)</code>	<code>jmp foo nop</code>

Convert mov Convert memory operand of `mov` into immediate operand. When position-independent code is disabled and `foo` is defined locally in the lower 32-bit address space, memory operand in `mov` can be converted into immediate operand. Otherwise, `mov` must be changed to `lea`.

Table B.2: Mov Conversion

Memory Operand	Immediate Operand
mov foo@GOTPCREL(%rip), %reg	mov \$foo, %reg
mov foo@GOTPCREL(%rip), %reg	lea foo(%rip), %reg

Convert Test and Binop Convert memory operand of `test` and `binop` into immediate operand, where `binop` is one of `adc`, `add`, `and`, `cmp`, `or`, `sbb`, `sub`, `xor` instructions, when position-independent code is disabled.

Table B.3: Test and Binop Conversion

Memory Operand	Immediate Operand
test %reg, foo@GOTPCREL(%rip)	test \$foo, %reg
binop foo@GOTPCREL(%rip), %reg	binop \$foo, %reg

Index

.cfi_adjust_cfa_offset, 98
.cfi_def_cfa, 98
.cfi_def_cfa_offset, 98
.cfi_def_cfa_register, 98
.cfi_endproc, 98
.cfi_escape, 98
.cfi_offset, 98
.cfi_rel_offset, 98
.cfi_startproc, 98
.eh_frame, 98
%rax, 53
_U _CLE NUP_PH SE, 86
_U _FORCE_UNWIND, 85
_U _SE RCH_PH SE, 85
_Unwind_Context, 84
_Unwind_DeleteException, 84
_Unwind_Exception, 84
_Unwind_ForcedUnwind, 84, 85
_Unwind_GetCF , 84
_Unwind_GetGR, 84
_Unwind_GetIP, 84
_Unwind_GetIPInfo, 84
_Unwind_GetLanguageSpecificData, 84
_Unwind_GetRegionStart, 84
_Unwind_RaiseException, 84, 85
_Unwind_Resume, 84
_Unwind_SetGR, 84
_Unwind_SetIP, 84
__float128, 15

auxiliary vector, 34
boolean, 17
byte, 15
C++, 103
Call Frame Information tables, 83
code models, 37
Convert call and jmp, 143
Convert mov, 143
Convert Test and Binop, 144

double quadword, 15
doubleword, 15
DT_FINI, 82
DT_FINI_ RR Y, 82
DT_INIT, 82
DT_INIT_ RR Y, 82
DT_JMPREL, 79
DT_PREINIT_ RR Y, 82
DW RF Debugging Information Format, 83, 97

eightbyte, 15
exceptions, 28
exec, 32

fegetround, 83
fourbyte, 15

General Dynamic Model, 120
General Dynamic To Initial Exec, 123
global offset table, 77

halfword, 15

Initial Exec Model, 120
Initial Exec Model, II, 120
Initial Exec To Local Exec, 124
Initial Exec To Local Exec, II, 124

Kernel code model, 38

Large code model, 38
Large position independent code model, 39
Local Dynamic Model, 121
Local Dynamic Model, II, 122
Local Dynamic to Local Exec, 125
Local Dynamic To Local Exec, II, 125
Local Exec Model, 122
Local Exec Model, II, 122
Local Exec Model, III, 123
`longjmp`, 84

Medium code model, 38
Medium position independent code model, 39
micro-architecture levels, 13

PIC, 39, 119
Procedure Linkage Table, 70
procedure linkage table, 78–79
program interpreter, 81

quadword, 15

`R_X86_64_JUMP_SLOT`, 79
`R_X86_64_TLSDESC`, 80
red zone, 21, 141
register save area, 54

`signal`, 28
sixteenbyte, 15
`size_t`, 17
Small code model, 38
small code model, 119
Small position independent code model, 39
small position independent code model, 119

`terminate()`, 85
Thread-Local Storage, 72
twobyte, 15

Unwind Library interface, 83

`va_arg`, 56, 137

`va_list`, 55

`va_start`, 55

`word`, 15